

# A Formally Verified NAT

# Background

- Middleboxes are important and everywhere.
- But middleboxes are buggy.
- The network becomes vulnerable due to middlebox crashes.

# What does the paper do?

- Formally verifying the functional correctness of a network function (middlebox).
- Verify a NAT (Network Address Translation) middlebox following RFC 3022 standard.
- Propose a framework, called Vigor, for combining symbolic execution with theorem prover.

# Choose Verification Technique

- Symbolic execution
  - KLEE: <http://klee.github.io/>, from an OSDI 2008 paper
- Theorem proving
  - Coq: <https://coq.inria.fr/>
  - Agda, Z3, F-star

# Symbolic Execution

- Human efforts: fully automated, low.
- Machine efforts: exhaustive search, super high.
- An assertion:
  - It is impossible for symbolic execution to verify complicated programs, as it takes forever.


# Theorem proving

- Human effort: high, I can show you my personal experience.
  - 19 lines of implementation.
  - 304 lines of proof.
- Machine effort: I would say moderate, a complicated proof usually requires hours to run.
  - Better than never stop.

# Find a balance

- Combine the best of symbolic execution and theorem proving.

## Vigor



# Vigor

- Use symbolic execution to check stateless part.
- Use theorem proving to check stateful data structures.
- Combine them together into a full-program verification.



# Vigor

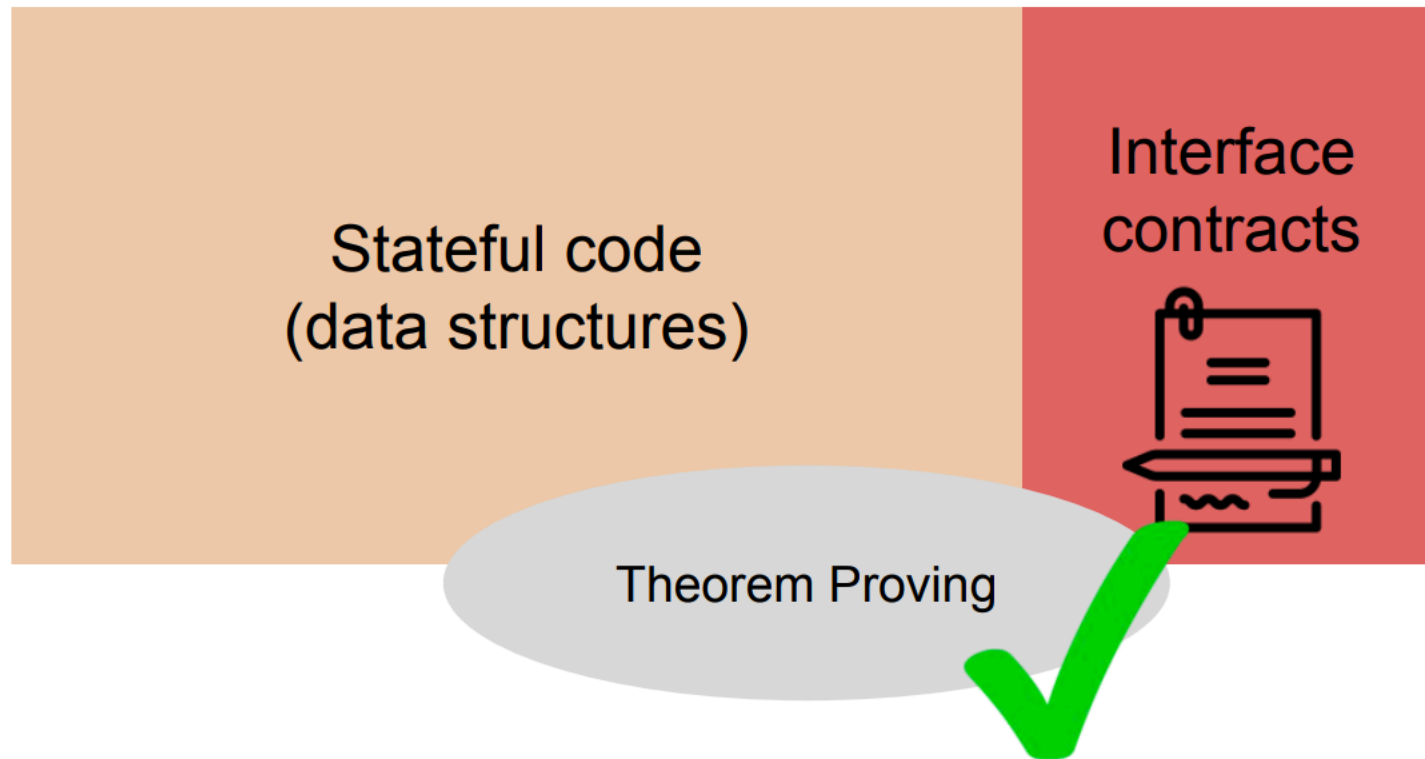
Stateful code  
(data structures)

Interface  
contracts



Stateless code  
(application logic)


# Theorem proving



# Theorem proving

- They use Z3 theorem prover.
- C code with annotations.
- Annotations:
  - Pre-condition: what you expect from the function input.
  - Post-condition: what you expect when the function finishes.
  - Function body annotations: for auxiliary


# Symbolic execution



# Symbolic Execution

- They use KLEE symbolic execution engine.
- Must avoid state explosion.
  - Replace stateful code with symbolic modules.
- Add another source to verify.

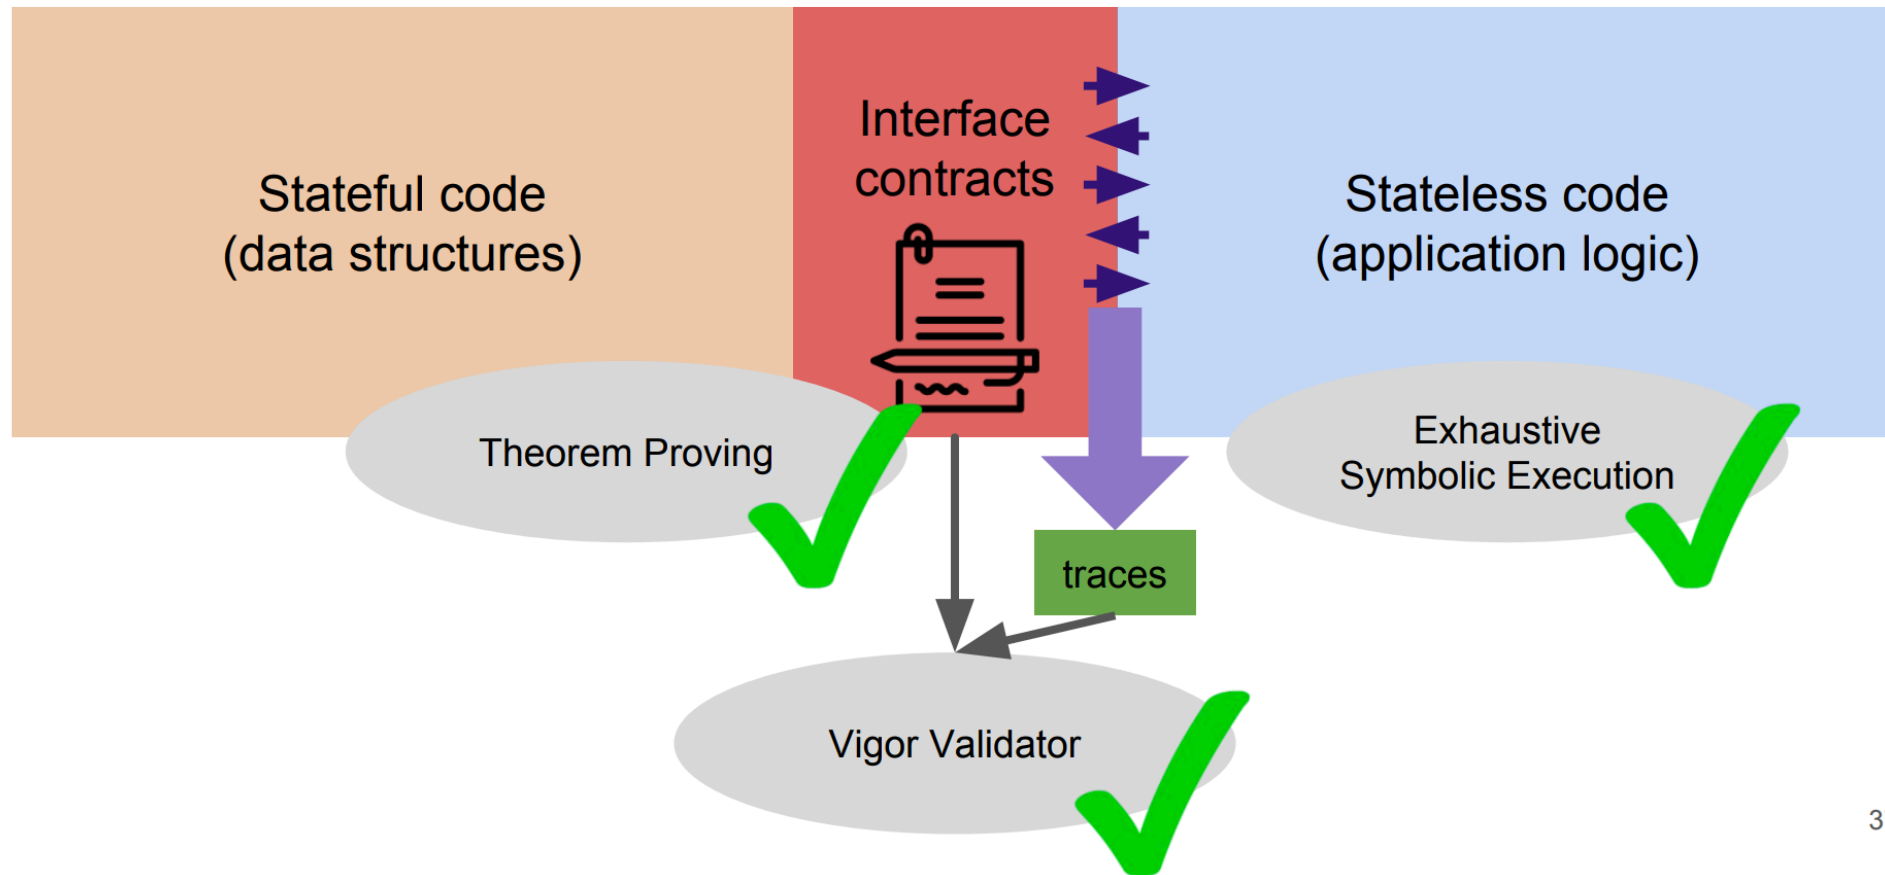
# Verify the Correctness of Symbolic Modules



# Verify the Correctness of Symbolic Modules


- They design a validator written in OCaml.
- Verify the traces collected during symbolic execution, against the interface contract they designed when verifying the stateful code.

# Vigor



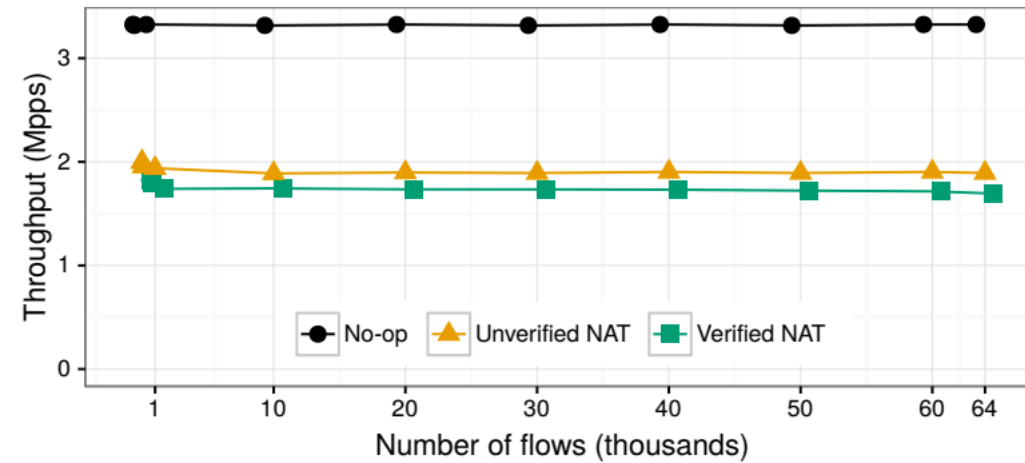


# Evaluation



**Figure 12: Average latency for probe flows. Confidence intervals are approximately 20 nanosec, not visible at this scale.**

# Evaluation



**Figure 14: Maximum throughput with a maximum loss rate of 0.1%.**

# Discussion

- Vigor's approach is new, but:
  - Ad-hoc.
  - Broke code modularity.
  - Can not scale to more complicated NFs, like Snort IDS.
- I would prefer a systematic programming framework, where you can do everything within this framework, not stitch things together.

# Discussion

- Theorem provers are more powerful than before.
- **“machine-checked mathematical proofs are the next big enabler of practical program structuring techniques”** – Adam Chlipala
- Proof automation:
  - You can provide some hints, the prover figures out the rest of the proof.
  - SMT solver

# F-star Language

- Developed by Microsoft and INRIA.
- Support most of the modern theorem proving features.
- A subset of F-star can be compiled to C for efficiency.

# Full Stack NFV Verification in F-star

- User space packet I/O framework (like netmap/dpdk).
- User space TCP/IP stack (like mTCP).
- Verifying complicated NF (like Snort IDS).

# Full Stack NFV Verification in F-star

- Huge correctness guarantee.
- Still good performance:
  - Verify in F-star.
  - Extract to efficient C code.
  - Compile with modern C compilers for speed.
- Good research opportunities.