

Optimal Online Multi-Instance Acquisition in IaaS Clouds

Wei Wang, Baochun Li, Ben Liang

USENIX ICAC'13

TPDS'14

Tradeoffs in Cloud Pricing Options

- On-demand Instances

- No commitment

- Pay-as-you-go

- Reserved Instances

- Reservation fee + discounted price


- Suitable for long-term usage commitment

Instance Type	Pricing Option	Upfront	Hourly
Standard Small	On-Demand	\$0	\$0.08
	1-Year Reserved	\$69	\$0.039
Standard Medium	On-Demand	\$0	\$0.16
	1-Year Reserved	\$138	\$0.078



Multi-Instance Acquisition Problem

- Workload (demand) is time-varying



- When should I reserve an instance?
- How many instances should I reserve?

Predict Future?

- Existing work relies on prediction of future demand
- However...
 - Prediction is needed for long-term future
 - Instance reservation period is typically months to years
 - Precise prediction is impossible
 - Demand prediction is limited
 - E.g., start-up companies, new services

How well can we make reservation decisions online, without any a *prior* information of the future demand?

Main Contributions

- Propose two online algorithms that offer the best provable cost guarantees
 - Deterministic: $(2-\alpha)$ -competitive
 - Randomized: $(e/(e-1+\alpha))$ -competitive
 - α is the hourly discount due to reservation ($0 \leq \alpha \leq 1$)
 - They only consider one instance type

Main Contributions

- Propose two online algorithms that offer the best provable cost guarantees
 - Deterministic: $(2-\alpha)$ -competitive
 - Randomized: $(e/(e-1+\alpha))$ -competitive
 - α is the hourly discount due to reservation ($0 \leq \alpha \leq 1$)
 - They only consider one instance type

Instance Type	Pricing Option	Upfront	Hourly
Standard Small	On-Demand	\$0	\$0.08
	1-Year Reserved	\$69	\$0.039
Standard Medium	On-Demand	\$0	\$0.16
	1-Year Reserved	\$138	\$0.078



Main Contributions

- Propose two online algorithms that offer the best provable cost guarantees
 - Deterministic: $(2-\alpha)$ -competitive
 - Randomized: $(e/(e-1+\alpha))$ -competitive
 - α is the hourly discount due to reservation ($0 \leq \alpha \leq 1$)
 - They only consider one instance type
- Study practical performance gains using Google workload traces

Problem Model

Pricing of On-demand Instances and Reserved Instances

- On-demand Instances
 - Fixed hourly price: p
 - Cost of running for h hours: ph
- Reserved Instances
 - Upfront reservation fee + discounted hourly price
 - Reservation fee is normalized to 1
 - Reservation period: τ
 - Cost of running for h hours: $1 + \alpha ph$
 - α is the hourly discount due to reservation ($0 \leq \alpha \leq 1$)

User demand and reservation

At t time slot, the user

- Has demand for d_t instances (time-varying)
- Newly reserves r_t instances
 - Available reserved instances: $\sum_{i=t-\tau+1}^t r_i$
- Purchases o_t on-demand instances
 - Total # of instances the user could launch at t :

$$o_t + \sum_{i=t-\tau+1}^t r_i \geq d_t$$

Offline Problem Formulation

Cost of on-demand instances

Cost of reserved instances

$$\min_{\{r_t, o_t\}} C = \sum_{t=1}^T (o_t p + r_t + \alpha p (d_t - o_t)),$$

s.t. $o_t + \sum_{i=t-\tau+1}^t r_i \geq d_t,$

$o_t, r_t \in \{0, 1, 2, \dots\}, t = 1, \dots, T.$

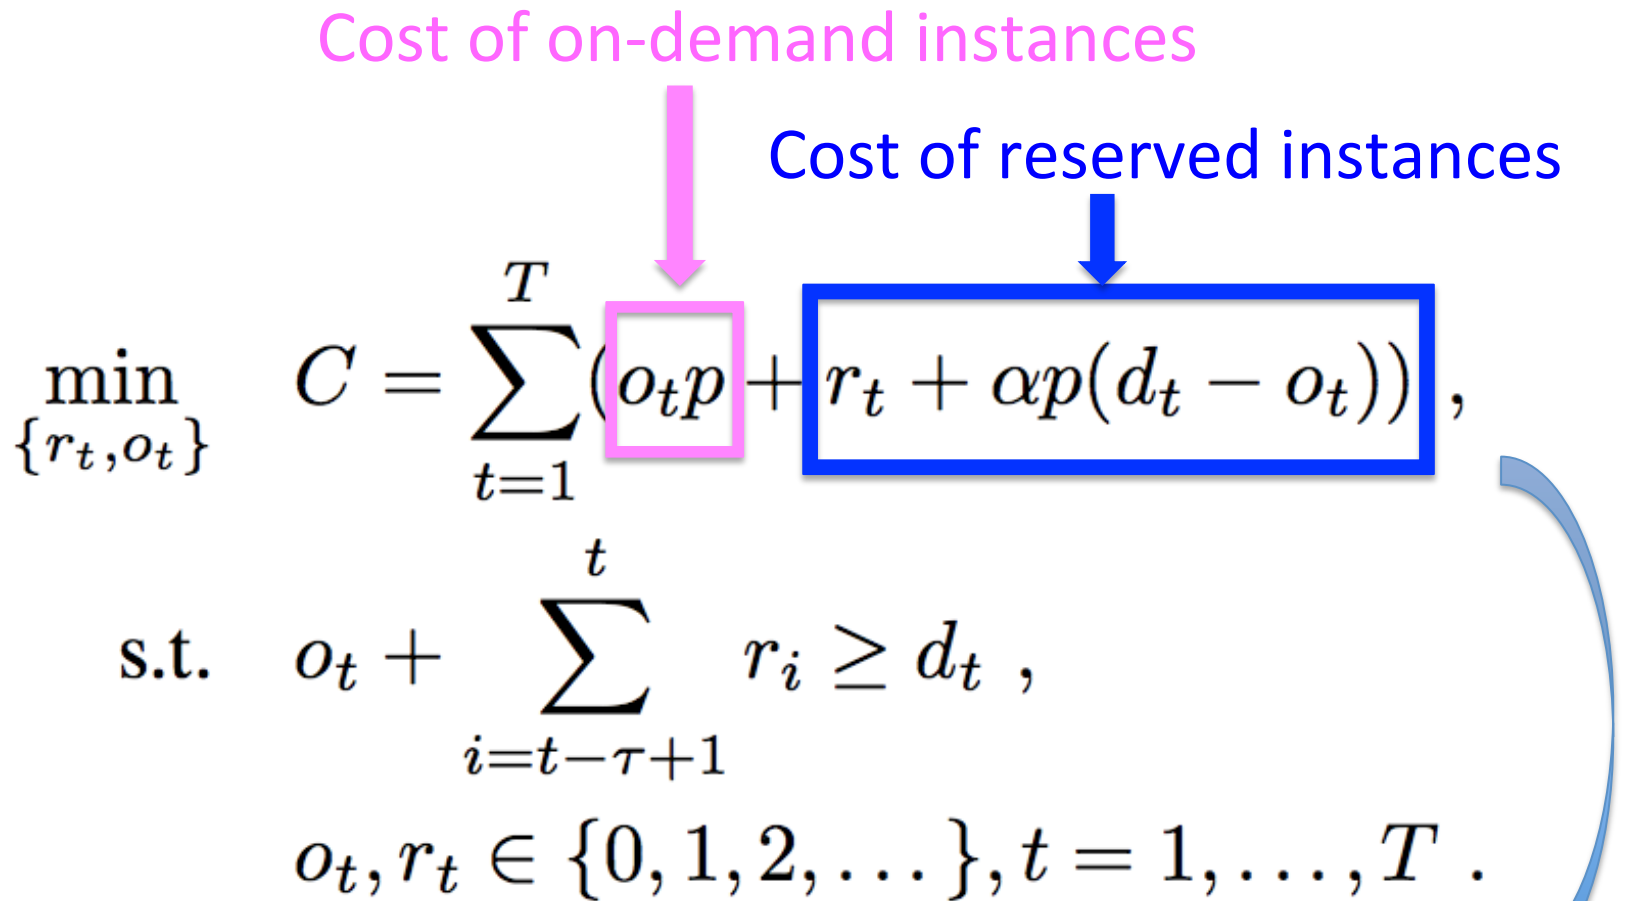
Offline Problem Formulation

Cost of on-demand instances

Cost of reserved instances

$$\min_{\{r_t, o_t\}} C = \sum_{t=1}^T (o_t p + r_t + \alpha p (d_t - o_t)),$$

s.t. $o_t + \sum_{i=t-\tau+1}^t r_i \geq d_t,$

$$o_t, r_t \in \{0, 1, 2, \dots\}, t = 1, \dots, T.$$


Hourly fee of reserved instances is based on usage

$$t = 1: \quad d_t = 7, r_t = 5, o_t = 2, d_t - o_t = 5, \tau = 2$$

$$\begin{aligned} \min_{\{r_t, o_t\}} \quad & C = \sum_{t=1}^T (o_t p + r_t + \alpha p (d_t - o_t)), \\ \text{s.t.} \quad & o_t + \sum_{i=t-\tau+1}^t r_i \geq d_t, \\ & o_t, r_t \in \{0, 1, 2, \dots\}, t = 1, \dots, T. \end{aligned}$$

Hourly fee of reserved instances are based on usage

$$t = 1: \quad d_t = 7, r_t = 5, o_t = 2, d_t - o_t = 5, \tau = 2$$

$$t = 2: \quad d_t = 4, r_1 = 5, r_2 = 0, o_t = 0, d_t - o_t = 4$$

$$\begin{aligned} \min_{\{r_t, o_t\}} \quad & C = \sum_{t=1}^T (o_t p + r_t + \alpha p (d_t - o_t)), \\ \text{s.t.} \quad & o_t + \sum_{i=t-\tau+1}^t r_i \geq d_t, \\ & o_t, r_t \in \{0, 1, 2, \dots\}, t = 1, \dots, T. \end{aligned}$$

Hourly fee of reserved instances are based on usage

$$t = 1: \quad d_t = 7, r_t = 5, o_t = 2, d_t - o_t = 5, \tau=2$$

$$t = 2: \quad d_t = 4, r_1 = 5, r_2 = 0, o_t = 0, d_t - o_t = 4$$

$$\begin{aligned} \min_{\{r_t, o_t\}} \quad & C = \sum_{t=1}^T (o_t p + r_t + \alpha p (d_t - o_t)), \\ \text{s.t.} \quad & o_t + \sum_{i=t-\tau+1}^t r_i \geq d_t, \\ & o_t, r_t \in \{0, 1, 2, \dots\}, t = 1, \dots, T. \end{aligned}$$

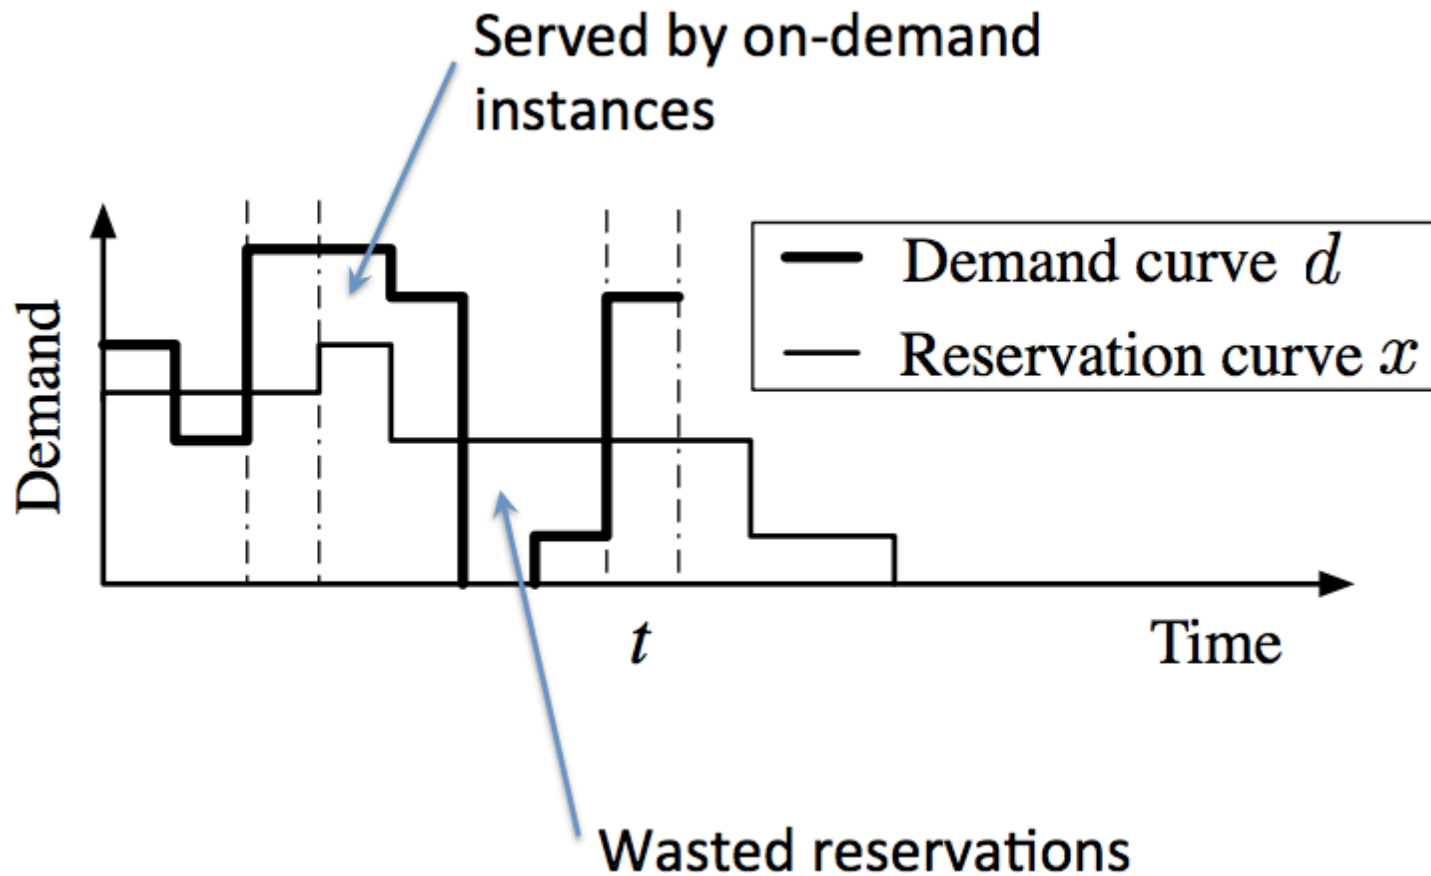
Hourly fee of reserved instances are based on usage

Lower-bound of competitive ratio

- Lemma 1: The best competitive ratio of this problem is at least $2-\alpha$ for deterministic online algorithms, and is at least $e/(e-1+\alpha)$ for randomized algorithms.
- R. Rleischer, “On the Bahncard Problem”, 2001

Optimal deterministic online algorithm

Demand and Reservation Curves

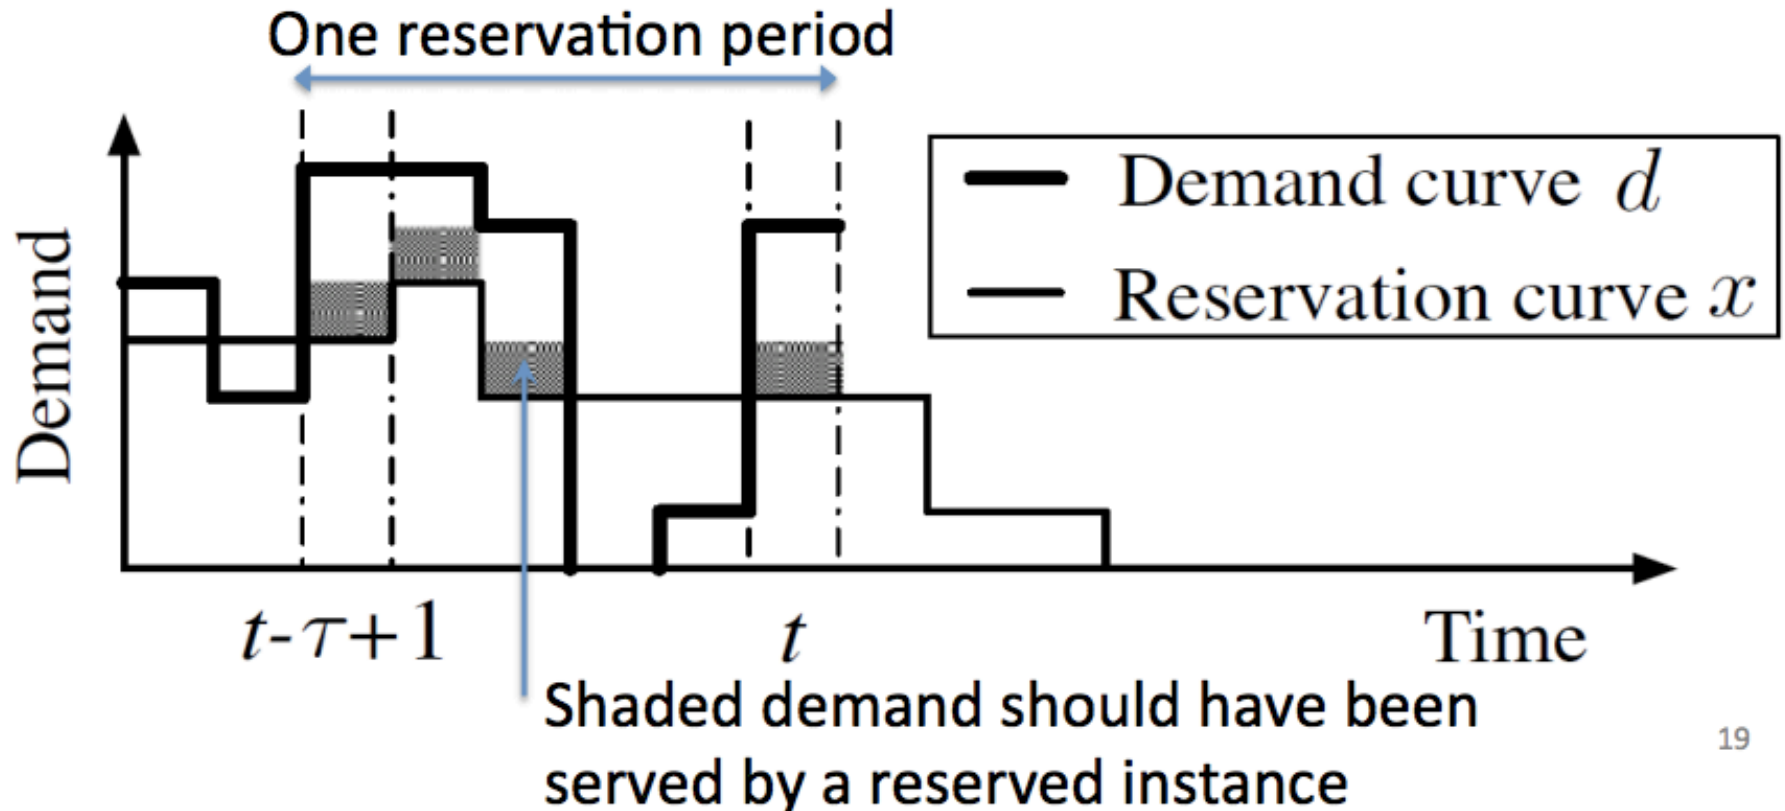


Break-Even Point

- Let c be the cost of running an on-demand instance that within a reservation period
- Using a reserved instance instead, the cost is $1 + \alpha c$
- Break-even point: $\beta = 1 + \alpha\beta \Rightarrow \beta = 1/(1-\alpha)$
 - If $c = \beta$, $c = 1 + \alpha c$, break even
 - If $c < \beta$, $c < 1 + \alpha c$, use an on-demand instance
 - If $c > \beta$, $c > 1 + \alpha c$, use a reserved instance


Regret and Compensation

- At time t , look back for a reservation period
 - If the incurred on-demand cost $> \beta$, reserve a new instance $r_t := r_t + 1$
- $\beta/p = 4$




Update reservation curve

If a new instance is reserved, lift the reservation curve (mark the compensated demand) by 1, both backward and forward for τ time slots



Repeat until no regret

- Repeat to reserve more instances, until the incurred on-demand instances cost $< \beta$.



- Proposition 1: This deterministic algorithm is $(2-\alpha)$ -competitive, hence is the optimal among all the deterministic algorithms of this problem

Optimal Randomized Online Algorithm

Basic Idea

- Strike balance between reserving too aggressively or too conservatively
- Randomly pick a threshold z instead of the break-even point β according to the density function:

$$f(z) = \begin{cases} (1 - \alpha)e^{(1-\alpha)z} / (e - 1 + \alpha), & z \in [0, \beta), \\ \delta(z - \beta) \cdot \alpha / (e - 1 + \alpha), & \text{o.w.}, \end{cases}$$

Basic Idea

- Strike balance between reserving too aggressively or too conservatively
- Randomly pick a threshold z instead of the break-even point β according to the density function:

$$f(z) = \begin{cases} (1 - \alpha)e^{(1-\alpha)z} / (e - 1 + \alpha), & z \in [0, \beta), \\ \delta(z - \beta) \cdot \alpha / (e - 1 + \alpha), & \text{o.w.}, \end{cases}$$


- Pick $z = \beta$ with probability of 1

- Proposition 2: This randomized algorithm is $(e/(e-1+))$ -competitive, and hence is an optimal among all the randomized algorithms of this problem


Trace Driven Simulation

Dataset and Preprocessing

- Google Cluster Traces
 - 900+ user's usage traces in a month
 - Users' computing demand data converted to IaaS instance demands




Users are classified into 3 groups based on demand fluctuation level



Five Algorithms


- All-On-Demand
- All-Reserved
- Deterministic
- Randomized
- Separate (At each demand level, run a Bahncard Algorithm)

CDF of Cost Normalized to the All-On-Demand




(b) Cost CDF (high fluctuation)

CDF of Cost Normalized to the All-On-Demand



(c) Cost CDF (medium fluctuation)

CDF of Cost Normalized to the All-On-Demand



(d) Cost CDF (stable demands)

Thank you

Q & A