

DeepXplore: Automated Whitebox Testing of Deep Learning Systems



Kexin Pei¹, Yinzhi Cao², Junfeng Yang¹, Suman Jana¹

¹Columbia University, ²Lehigh University




Deep learning (DL) has matched human performance!

- Image recognition, speech recognition, machine translation, intrusion detection...
- Wide deployment in real-world systems



Amazon Echo



Iphone X



Google Earphone




Deep learning is increasingly used in safety-critical systems

- Deep learning correctness and security is crucial



Self-driving car



Medical diagnosis



Malware detection




Unreliable deep learning: fooling self-driving cars



DRV_C1:right




DRV_C2:right




DRV_C3:right




Unreliable deep learning: fooling face recognition




Real
glasses



Lujo Bauer



=



John Malkovich

100% success





Existing DL testing methods are seriously limited

- Common practice: measure accuracy on a test input set of randomly chosen data
- Problem 1: how good is the coverage of the test set?
 - DL decision logic is incredibly complex
 - More fundamentally, what is testing coverage metric for DL?
- Problem 2: it requires expensive labeling effort
 - Data in test set must be manually labelled
 - To enlarge the test set, we need to manually label more data



Existing DL testing methods are seriously limited (cont.)

- Adversarial testing (Szegedy et al. ICLR'14): find corner-case inputs imperceptible to human but induce errors
 - Problem 1: how good is the coverage of the test set?
 - Problem 2: it requires expensive labeling effort
 - Problem 3: Not realistic. (Theoretical, assumes a very powerful adversary. [Sharif et al. CCS'16])



School bus



Carefully crafted noise

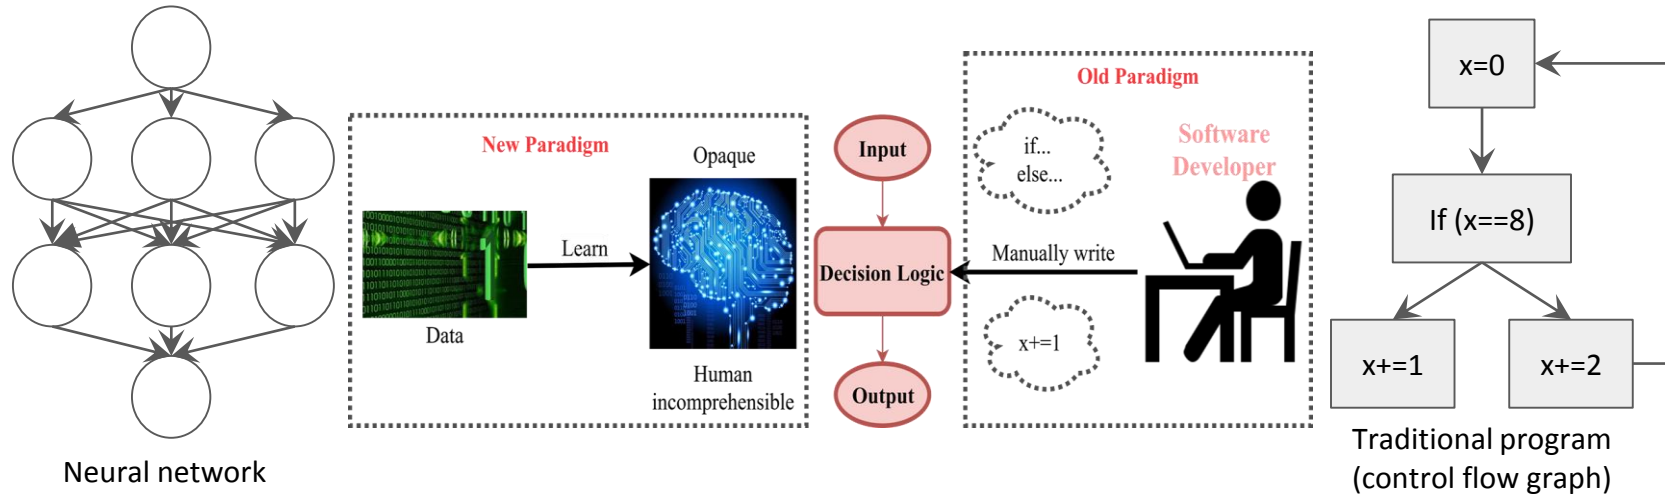


Ostrich



Many traditional software testing techniques don't apply to DL

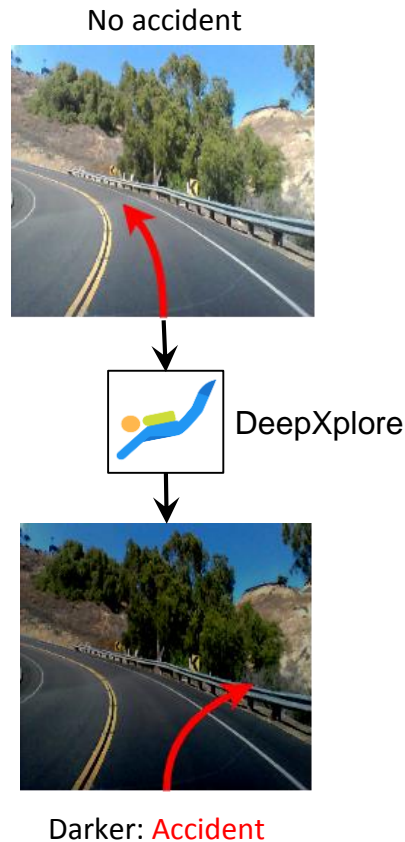
- DL decision logic is embedded in neurons and layers, not in code





Quick Summary of DeepXplore

- The first step towards systematic testing of Deep Neural Nets (DNNs)
- Neuron coverage: first testing coverage metric for deep neural net
- Automated: cross-check multiple DNNs
- Realistic: physically realizable transformations
- Effective:
 - 15 State-of-the-art DNNs on 5 large datasets (ImageNet, Self-driving cars, PDF/Android malware)
 - Numerous corner-case errors
 - 50% more neuron coverage than existing testing






Outline

- Quick deep learning primer
- Workflow of DeepXplore
 - Design
 - Detail of Neuron coverage
- Implementation
- Evaluation setup and results summary



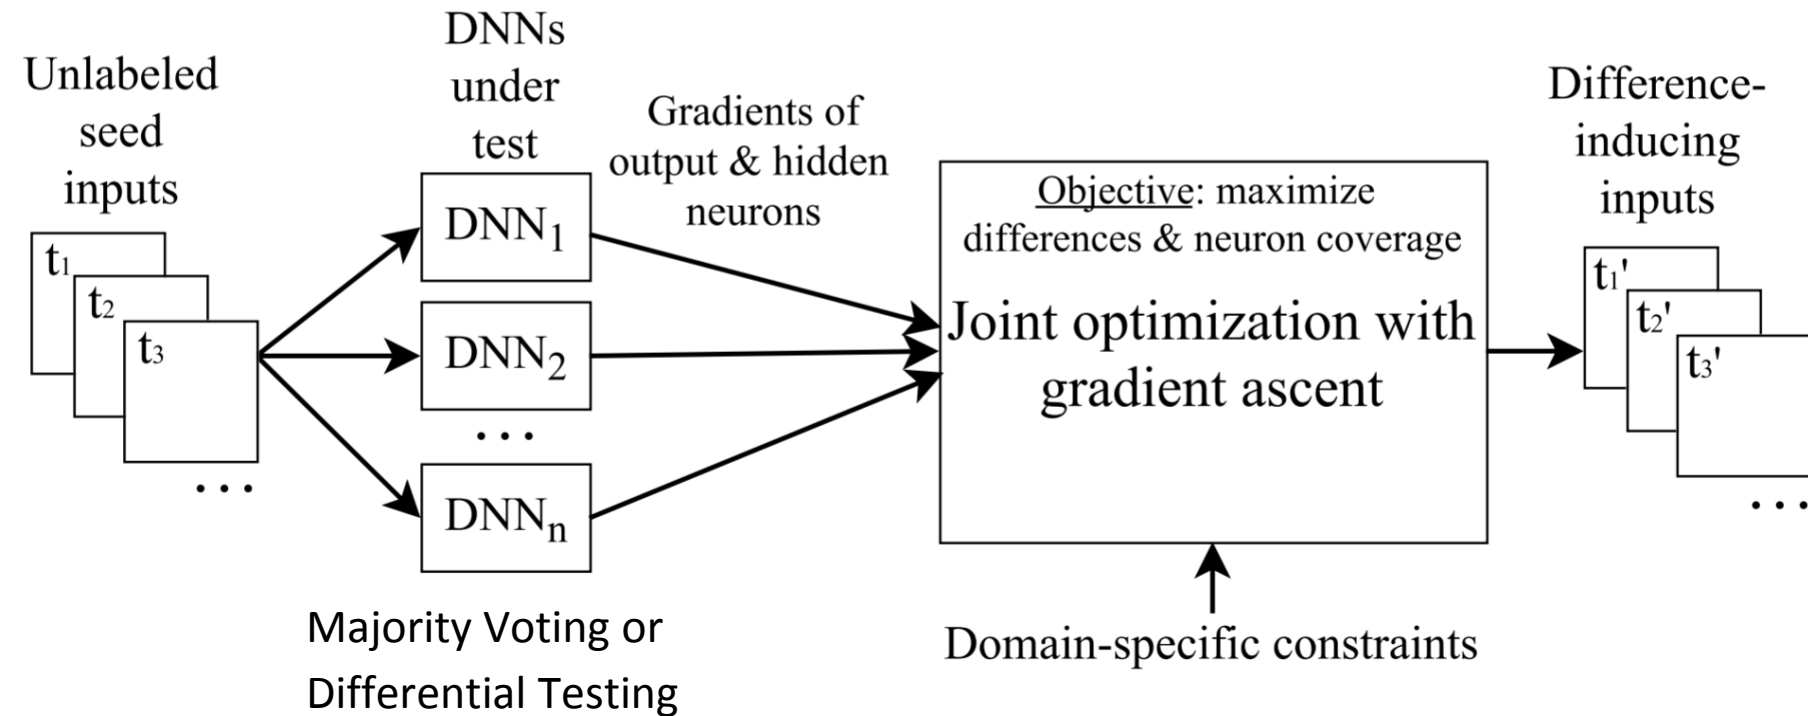
Deep learning primer

- A neural network is a function $f(X) \rightarrow Y$
 - Trainable parameters (W_i) on each edge and nonlinear activation function at each neuron
 - DNN learns the weights during training
- Training: Given training set (X, Y) , adjust W to minimize the prediction error (slow)
- Inference: Simply propagates X through layers (fast)







How DeepXplore Works?





How To Maximize Differences?





How To Maximize Differences?

- MAXIMIZE

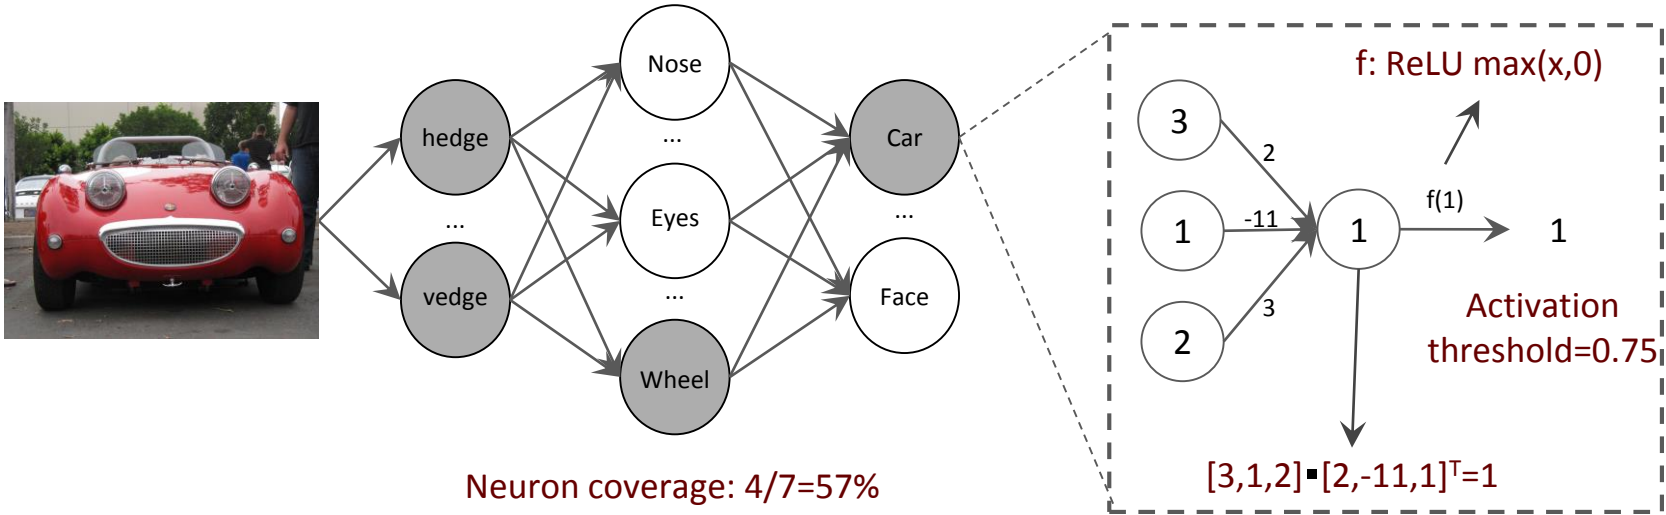
{sum of classification probabilities of other networks - classification probability of the debugging network}

$$obj_1(\mathbf{x}) = \sum_{k \neq j} F_k(\mathbf{x})[c] - \lambda_1 \cdot F_j(\mathbf{x})[c]$$



How To Maximize Neuron Coverage?

- Neuron coverage = # neurons activated at least once / # total neurons
- A neuron is activated if its output is larger than a threshold, e.g., 0.





How To Maximize Neuron Coverage?

- Randomly select an inactivated neuron and maximize its output.

$$obj_2(\mathbf{x}) = f_n(\mathbf{x})$$



Joint Optimization

- Weighted sum of the above two objectives.

$$obj_{joint} = (\sum_{i \neq j} F_i(\mathbf{x})[c] - \lambda_1 F_j(\mathbf{x})[c]) + \lambda_2 \cdot f_n(\mathbf{x})$$

- Use **Gradient Ascent** to solve the above problem, the **variable is \mathbf{x}** , i.e., the input of the neural network.
- Given an original input of the network \mathbf{x} , we may find a modified input \mathbf{x}' *that causes different behaviors, but may not.*



How to achieve multiple goals simultaneously

- Goal 1: systematically find corner cases
 - Generate inputs that maximize neuron coverage
- Goal 2: find DNN errors without manual labels
 - Differential testing: use multiple DNNs as cross-referencing oracles
- Goal 3: generate realistic test inputs
 - Domain-specific constraints, e.g., a pixel value must be an integer within $[0, 255]$

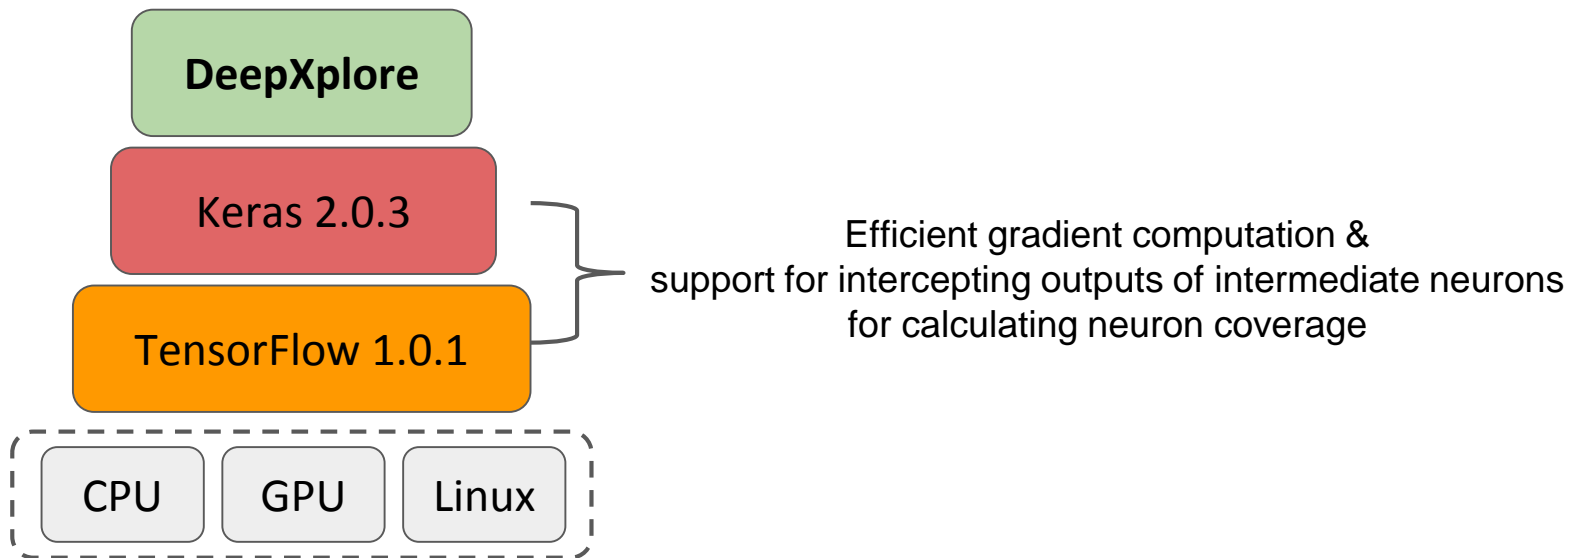


Outline

- Quick deep learning primer
- Workflow of DeepXplore
 - Design
 - Detail of Neuron coverage
- **Implementation**
- Evaluation setup and results summary



Implementation





Outline

- Quick deep learning primer
- Workflow of DeepXplore
 - Design
 - Detail of Neuron coverage
- Implementation
- Evaluation setup and results summary



Evaluation setup and results summary

Dataset	Description	DNNs	Original random testing accuracy	Avg. neuron coverage improvement over random/adversarial	Avg. Violations found by DeepXplore (2000 seeds)
MNIST	Handwritten digits	LeNet variants	98.63%	30.5% → 70%	1,289
ImageNet	General Images in 1000 categories	VGG16, VGG19, ResNet15	93.91%	1% → 69%	1,980
Driving	Udacity self-driving car competition dataset	Nvidia Dave-2 variants	99.94%	3.2% → 59%	1,839
Contagio/ VirusTotal	PDF malware	Fully connected	96.29%	18% → 70%	1,048
Drebin	Android malware	Fully connected	96.03%	18.5% → 40%	2,000



Sample corner-case errors for images

Driving:



Go straight



Turn right

Driving:



Go straight



Turn left

ImageNet:



Light cardigan



Diaper

ImageNet:




bolete




buckeye

MNIST

:




6




2

MNIST

:



5



3



Performance

- For all tested DL models, on average **DeepXplore generated one test input demonstrating incorrect behavior within one second** while running only on a commodity laptop.
- Given a set of DNNs, we can use majority voting to automatically label the exceptional cases generated by DeepXplore. Using these as new labeled training examples improves accuracy by 1-3%.



Conclusions

- Systematically testing DL for realistic corner cases is a hard problem
- DeepXplore is the first step for systematic DL testing
 - Neuron coverage: first testing coverage metric for deep neural net
 - Automated: differential testing by cross-checking multiple DNNs
 - Realistic: physically realizable transformations
 - Effective: find numerous unexpected corner-case errors
- A lot of exciting new research problems!
 - Build analysis tools for testing and verification of ML
 - Build better debugging support for opaque ML



Comments

- I like the ideas:
 - Apply traditional software debugging techniques (code coverage, differential testing) to debug DNNs.
 - Formulate the testing problem as an joint-optimization problem and solve it using gradient-based method.
- This paper focuses on **classification problem** in supervised learning.
- Neuron coverage can not cover all possible inputs.
- The constraints are not flexible enough and the resulting inputs can only be the input used for training, plus some minor modifications.
- DL systems probably can't handle all safety issues alone, we need non-DNN “safety-wrapper”.



Thank you!
Questions?