

Responsive Multipath TCP in SDN-based Datacenters and Dynamic Scaling of Virtual Network Functions

Duan Jingpu

University of Hong Kong

jpduan@cs.hku.hk

Feb. 2, 2015

1 First Section

- Responsive Multipath TCP in SDN-based Datacenters

2 Second Section

- Dynamic Scaling of Virtual Network Functions

Responsive Multipath TCP in SDN-based Datacenters

- Section 1: Responsive Multipath TCP in SDN-based Datacenters.

Background and Motivation

- In datacenter network, content replication, virtual machine migration, and data shuffling in MapReduce tasks generate many large flows.
- These large flows constitute the majority of datacenter traffic.
- Efficient transfer of these large flows is crucial to the performance of a datacenter network.
- But current design of both transport layer protocol and routing scheme is not efficient enough in supporting the high throughput of large flows.

Routing in Datacenter Network

- Datacenter network has redundant paths that connect servers from different racks.
- How to route?

Routing in Datacenter Network

- Rely on ECMP to balance the traffic on different paths.
- But ECMP is not intelligent.
- Two large flows may be routed on the same paths, so none of them can achieve optimal throughput.

Routing in Datacenter Network

- Calculate route using SDN controller, like Hedera.
- More intelligent, can avoid paths collision, constantly re-balance entire traffic.
- But controller needs to constantly polls traffic statistics from all switches. Bad scalability and responsiveness.
- Existing flows may be re-routed to another path. Possible packet reordering and packet loss.

Datcenter Transportation Protocol

- Traditional TCP has been shown to be inefficient in datacenter network.
- People design TCP variants that target the needs of datacenter network, such as DCTCP and D2TCP.
- Their limitation is that they are not multi-path protocol.
- Failed to fully utilize path diversity in datacenter network.

Datacenter Transportation Protocol

- Use multi-path based transportation protocol to exploit available bandwidth.
- Split one TCP flow into multiple subflows.
- So that available bandwidth is more efficiently used.
- How to split?


Datacenter Transportation Protocol

- Split at the switch.
- Reply on advanced switch features.
- Hard to deploy in a large scale.

Datacenter Transportation Protocol


- Split at server, using MPTCP.
- Then flows are still routed using ECMP. We still have all mentioned potential problems.
- Number of subflows used by each flow is fixed. Add unnecessary overhead, can't react to traffic conditions.

Our Work




- Responsive MPTCP system for SDN-based datacenters.
- Controller + Monitor design.

Our Work




- Controller estimates the demand for each MPTCP flow.
- Controller sends updated demand to the monitor.

Our Work




- Monitor records updated demand and current throughput.
- Monitor issues path adding request when necessary.

Our Work



- Controller calculates needed additional subflows.
- Controller sends this information back to monitor.

Our Work



- Dynamically decide the number of subflows to be used by each MPTCP flow and the best subflow path.
- With this system, better throughput, smaller overhead.

Controller

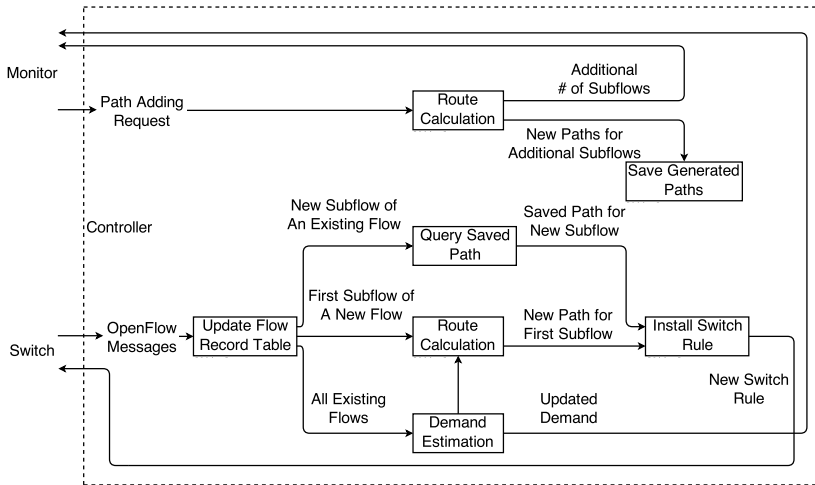


Figure: Architecture of Controller

Controller: Demand Estimation

- Fat-tree datacenter network is a fully non-blocking network.
- Bandwidth bottleneck along a flows path is either senders access link or receivers access link.
- We employ the demand estimation algorithm in Hedera to calculate flow's demand.
 - Proportionally increase the sending rate at sender's access link.
 - Proportionally decrease excessive sending rate at receiver's access link.
 - Stop until all sending rates stabilize.

Controller: Demand Estimation

Algorithm: Demand Estimation and Dispatching

Input: demand record d_t , new flow f_n or expired flow f_e

```
1:  $old\_d\_t \leftarrow d\_t$ ;  
2: update  $d\_t$  to include  $f\_n$  or exclude  $f\_e$ ;  
3: run Hedera demand estimation algorithm with  $d\_t$  as input;  
4: for each flow  $f$  in  $d\_t$  do  
5:   if  $f$  is not in  $old\_d\_t$  then  
6:     dispatch  $d\_t[f]$  to the monitor at sender of  $f$ ;  
7:   else  
8:     if  $|d\_t[f] - old\_d\_t[f]| > \delta_{DDT} * old\_d\_t[f]$  then  
9:       dispatch  $d\_t[f]$  to the monitor at sender of  $f$ ;  
10:    end if  
11:  end if  
12: end for
```


Controller: Route Calculation

- Route calculation algorithm computes a set of new paths for the flow.
- It uses demand gap as input.
- The goal of route calculation algorithm is to cover as much the demand gap as possible.

- Property 1: Aggregate MPTCP Flows Fairly Share Link Bandwidth.
 - When multiple subflows of a MPTCP flow traverse the same link, they are viewed as one aggregate flow.
 - Aggregate MPTCP flows fairly share the link bandwidth, because of MPTCP congestion control protocol.
 - So when n aggregate flows are contending for a link with bandwidth B , the expected throughput of each aggregate flow can be calculated as B/n .


Controller: Route Calculation

- Property 2: Marginal Throughput.
 - Subflows at most Share Links in the Aggregation-ToR Layer, as illustrated in the following figure.



Controller: Route Calculation

- Property 2: Marginal Throughput.
 - Subflows at most Share Links in the Aggregation-ToR Layer.
 - Marginal throughput quantifies the achievable throughput of a subflow on a new path.
 - Especially when new subflow and existing subflows share links on this new path.



Controller: Route Calculation

- We apply the following two rules to calculate the expected throughput of a new subflow sf of flow f on a given path p .
- **Rule 1:** If there is no other subflow of f that uses p 's aggregation-ToR uplink, find out the largest number of aggregate flows sharing the same link on p , n_{max} . The expected throughput of sf on p is B/n_{max} .
- **Rule 2:** If there are m existing subflows of f that use p 's aggregation-ToR uplink, calculate p 's marginal throughput m_t , as well as the upper bound of sf 's expected throughput B/n_{max} .

Controller: Demand Estimation

Algorithm: Route Calculation

Input: flow f , demand gap d_g , number of existing subflows n_sf

```
1: set number of additional subflows  $n\_new\_sf = 0$ ;  
2: while  $n\_new\_sf + n\_sf < M$  do  
3:   find out a subflow path  $p$  with the largest expected throughput  $e\_t$ ;  
4:   if  $f$  is a new flow then  
5:     install switch rules on  $p$ ;  
6:     return;  
7:   end if  
8:   if  $e\_t == 0$  then  
9:     return;  
10:  end if  
11:  save path  $p$ ;  
12:   $n\_new\_sf += 1$ ;  
13:  if  $e\_t > d\_g$  then  
14:    notify the sender monitor of  $f$  of  $n\_new\_sf$ ;  
15:    return;  
16:  else  
17:     $d\_g = d\_g - e\_t$ ;  
18:  end if  
19: end while  
20: notify the sender monitor of  $f$  of  $n\_new\_sf$ ;  
21: return;
```

- Monitor is a daemon program running on each server.
- It keeps track of current throughput of each active flow.
- It constantly receives updated demand from controller.
- It issues path-adding request when it detects a significant gap between current throughput and updated demand.
- Monitor daemon is periodically executed every t_m seconds.

Monitor

Algorithm: Monitor Loop

Input: monitored flows $f[n]$, previous flow rates $p_r[n]$, counters $count[n]$, flow demand $demand[n]$

```
1: for  $i = 1 : n$  do
2:   obtain instant throughput  $i\_r$  for flow  $f[i]$ ;
3:    $p\_r[i] = 0.2 * p\_r[i] + 0.8 * i\_r$ ;
4:   if  $|i\_r - p\_r[i]| < \delta_{RVT} * p\_r[i]$  then
5:      $count[i] = count[i] + 1$ ;
6:     if  $count[i] == R$  then
7:        $count[i] = 0$ ;
8:       if  $p\_r[i] < (1 - \delta_{DGT}) * demand[i]$  then
9:          $gap = demand[i] - p\_r[i]$ ;
10:        issue path adding request for  $f[i]$  with  $gap$ ;
11:        return;
12:      end if
13:    else
14:      return;
15:    end if
16:  else
17:     $count[i] = 0$ ;
18:    return;
19:  end if
20: end for
```

Performance Evaluation

- We evaluate our responsive MPTCP system using a NS3 simulator.
- We simulate a datacenter network with a 8-Ary fat-tree topology which contains 128 servers connected using 1Gbps links.
- Up to 4 subflows per MPTCP connection is allowed.
- Performance under different parameter settings is evaluated.
- Comparison against ECMP and Hedera is conducted.

TABLE I: Summary of Experiment Results



	Random Traffic		Permutation Traffic		Shuffling Traffic	
	NS	TP(Mbps)	NS	TP(Mbps)	JCT(ms)	NS
$t_m = 2ms, \delta_{DGT} = 0.15$	3.72	521	3.88	766	1402	1.92
$t_m = 4ms, \delta_{DGT} = 0.25$	3.40	513	3.82	767	1441	1.28
$t_m = 8ms, \delta_{DGT} = 0.35$	2.47	506	3.75	760	1454	1.03
4 subflows+ECMP	4	530	4	734	1332	4
2 subflows+ECMP	2	493	2	643	1394	2
1 subflow+Hedera	1	438	1	747	1652	1

NS: Average number of subflows per flow.



TP: Average throughput per flow.

JCT: Average shuffle completion time per MapReduce job.



Performance Evaluation



Performance Evaluation



Performance Evaluation



Dynamic Scaling of Virtual Network Functions

- Section 2: Dynamic Scaling of Virtual Network Functions.

- Network middleboxes, i.e. NAT, Firewall, Proxy, are running everywhere.
- Many of them are usually implemented as proprietary hardware, making them notoriously hard to upgrade and scale.
- With cloud and virtualization technology, network middleboxes with a software-implementation can be run on virtual machines.

- Efforts have been made on improving the performance of NFV system:
 - Increase the processing speed of software network middlebox running on virtual machine.
 - Dynamically scale network middleboxes in face of traffic change.

IMS System

- We start our research by studying a specific network middlebox system: Ip Multi-media Subsystem (IMS).
- We use an open-source implementation of IMS system, Project Clearwater.

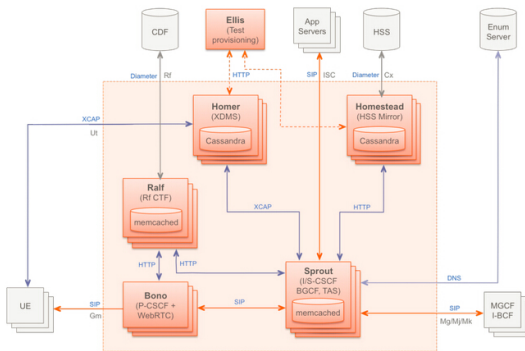



Figure: Architecture of IMS System, from Project Clearwater Website

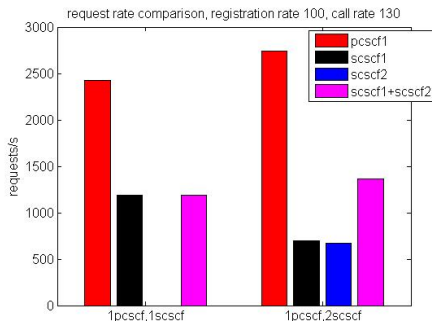
Identifying the Bottleneck

- In order to dynamically scale the system, we need to identify the bottleneck in the system.
- We design the following simple experiment to show the bottleneck of the system.



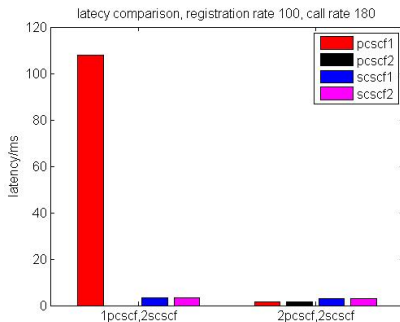
Identifying the Bottleneck

- With registration rate 100 and call rate 130, SCSCF overloads.
- SCSCF performs blocking operations that retrieve user information from database.
- When it overloads, it fails to serve some calls and for some calls to terminate.




Identifying the Bottleneck

- With registration rate 100 and call rate 180, PCSCF overloads.
- PCSCF performs non-blocking operations, acting as a pure proxy.
- When it overloads, its request processing latency increases drastically.




Identifying the Bottleneck

- With registration rate 100 and call rate 320, PCSCF overloads again.
- SCSCF will overload soon, if workload keep increasing.



Identifying the Bottleneck

- With registration rate 100 and call rate 320, PCSCF overloads again.
- SCSCF will overload soon, if workload keep increasing.



What We Learn

- Under different workload, different part of the system becomes bottleneck.
- It's hard to design an accurate mathematical model that takes workload as input and output the number of required resources.

What We Plan to Do

- Monitor the overloading signal. (Latency, CPU usage, Successfully Completed Requests)
- Scale the system in react to the overloading signal.
- After scaling, need to balance the load on each instance.
- Scale up vs scale down:
 - It's easier to scale up than to scale down.
 - Because it's even harder to determine whether you can scale down.

- Future Work:

- Along the direction of datacenter network, design more efficient datacenter transport protocols with control plane assistance and implement them in real system.
- Along the direction of NFV, carry on this existing research and find out more interesting from it.

The End

- Thank you and Q & A!