

# Human-level Control Through Deep Reinforcement Learning

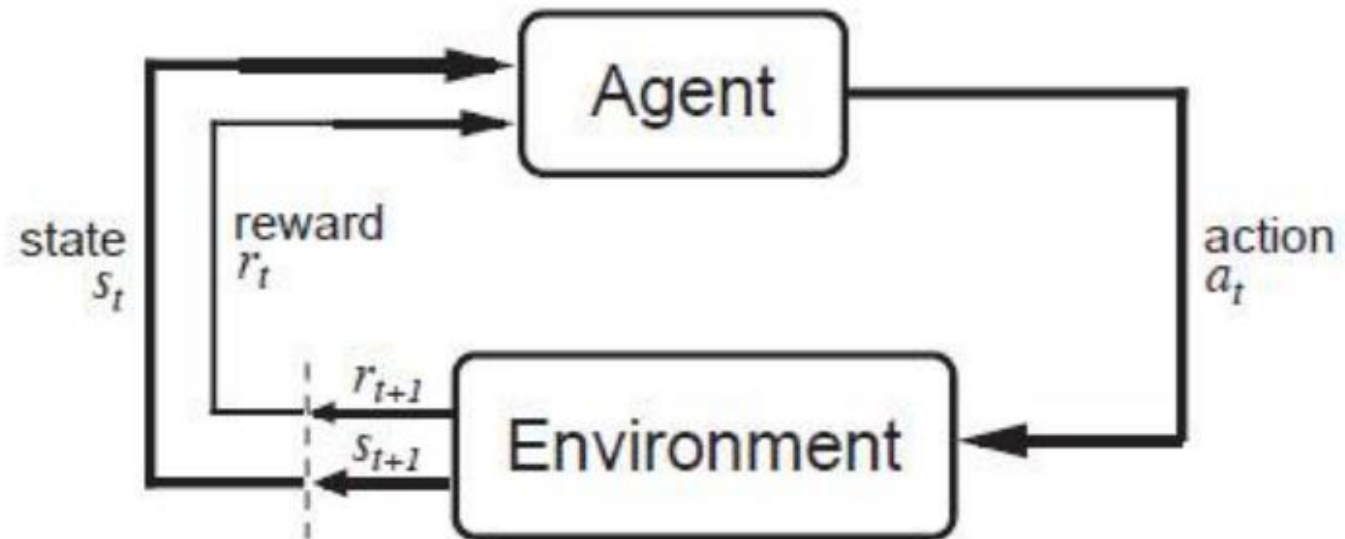
**Google DeepMind**

**Nature 2015**

# Introduction

## Markov Decision Process

- State
- Action
- Reward



# Introduction

- Action value function
  - Discounted future reward (environment is stochastic)

$$\begin{aligned} R_t &= r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \cdots + \gamma^{n-t} r_n \\ &= r_t + \gamma (r_{t+1} + \gamma (r_{t+2} + \cdots)) \\ &= r_t + \gamma R_{t+1} \end{aligned}$$

$$Q^\pi(s, a) = E_\pi\{R_t | s_t = s, a_t = a\} = E_\pi\left\{\sum_{k=0}^{\infty} \gamma^k r_{t+k+1} | s_t = s, a_t = a\right\}$$

# Introduction

- Q-learning
- Bellman Equation

$$Q(s, a) = r + \gamma \max_{a'} Q(s', a')$$

```
initialize  $Q[num\_states, num\_actions]$  arbitrarily  
observe initial state  $s$   
repeat  
    select and carry out an action  $a$   
    observe reward  $r$  and new state  $s'$   
     $Q[s, a] = Q[s, a] + \alpha(r + \gamma \max_{a'} Q[s', a'] - Q[s, a])$   
     $s = s'$   
until terminated
```

-Limitation:

Value Iteration

1. Very limited states/actions
2. Can not generalize to unobserved states

# Introduction

Deep Q-network (DQN)

- Q learning plus
- Function approximator:

Deep neural networks to approximate optimal action-value function

$$Q^*(s,a) = \max_{\pi} \mathbb{E}[r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots | s_t = s, a_t = a, \pi].$$

# Problem

## Stability issues with Deep RL

- Naïve Q-learning **oscillates** or **diverges** with neural nets
  1. Data is sequential
    - Successive samples are correlated, non-i.i.d.
  2. Policy changes rapidly with slight changes to Q-values
    - Policy may oscillate
    - Distribution of data can swing from one extreme to another

# Experience Replay

To remove correlations, build data-set from agent's own experience

- Take **action  $a_t$  according to  $\epsilon$ -greedy policy**  
(Choose “best” action with probability  $1 - \epsilon$ , and selects a random action with probability  $\epsilon$ )
- Store transition  $(s_t, a_t, r_t, s_{t+1})$  in **replay memory  $\mathcal{D}$**  (Huge data base to store historical samples)
- Sample **random mini-batch of transitions  $(s, a, r, s')$**  from  $\mathcal{D}$
- Optimize MSE between Q-network and Q-learning targets, e.g.

$$\mathcal{L}_i(\theta_i) = \mathbb{E}_{s,a,r,s' \sim \mathcal{D}} \left[ \underbrace{\left( r + \gamma \max_{a'} Q(s', a'; \theta_i) \right)}_{\text{target}} - Q(s, a; \theta_i) \right]^2$$

# Fixed target Q-network

To avoid oscillations, **fix parameters used in Q-learning target**

- Compute Q-learning targets w.r.t. **old, fixed parameters  $\theta_i^-$**

$$r + \gamma \max_{a'} Q(s', a'; \theta_i^-)$$

- Optimize MSE between Q-network and Q-learning targets

$$\mathcal{L}_i(\theta_i) = \mathbb{E}_{s,a,r,s' \sim \mathcal{D}} \left[ \left( r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta_i) \right)^2 \right]$$

- **Periodically update fixed parameters  $\theta_i^- \leftarrow \theta_i$**



# Core components of DQN

Deep Q-Network provides a stable solution to deep value-based RL

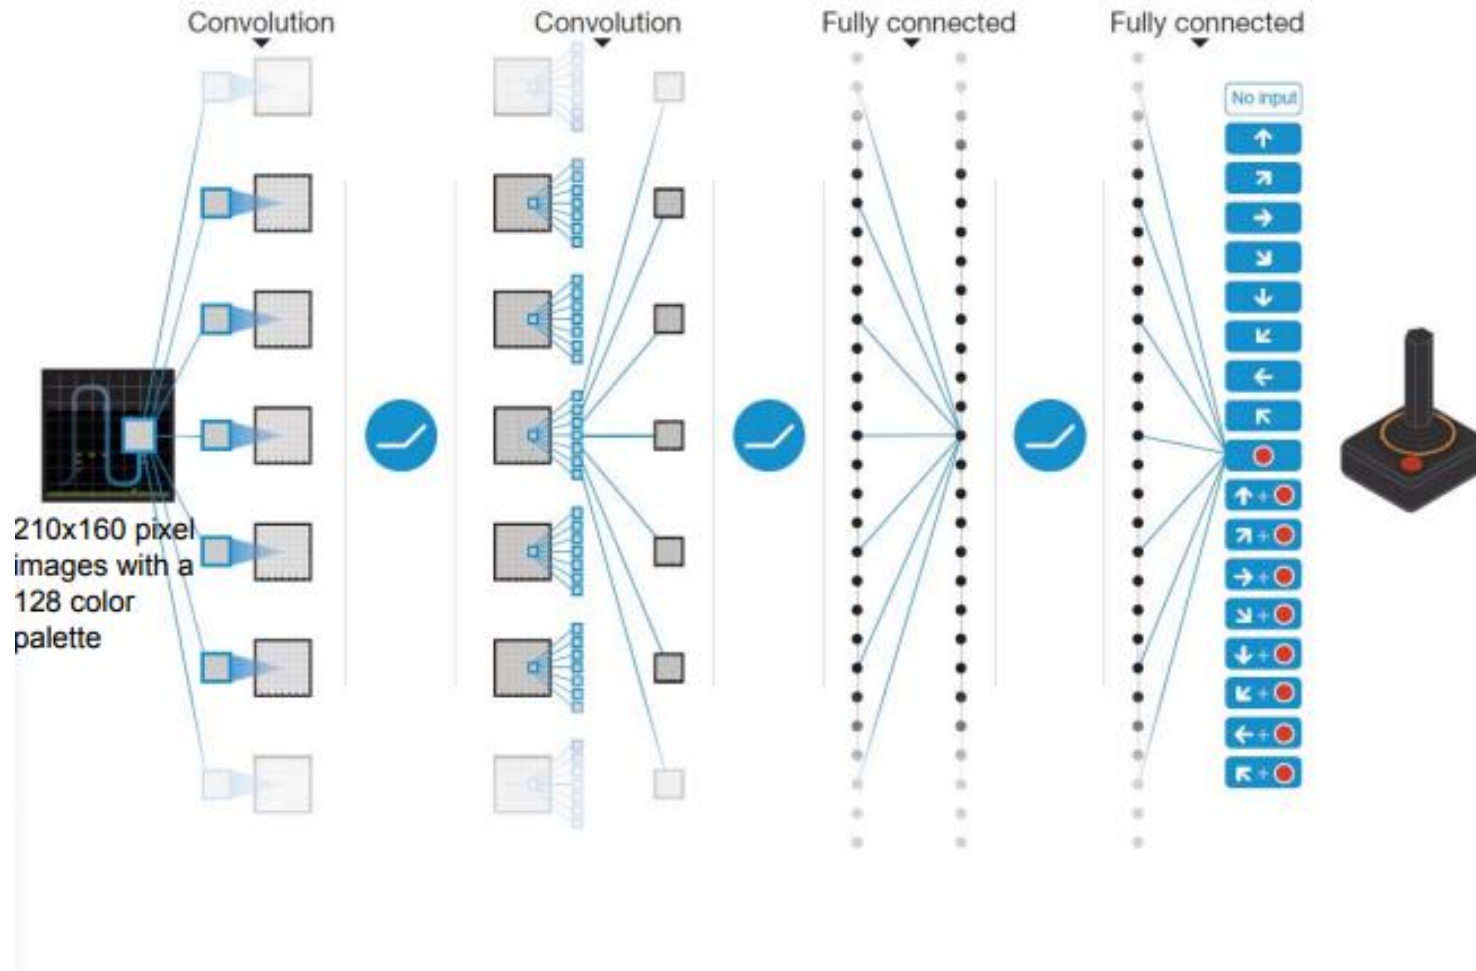
1. Use **experience replay**

- Break correlations in data, bring us back to i.i.d. setting
- Learn from all past policies
- Using off-policy Q-learning

2. Freeze **target Q-network**

- Avoid oscillations
- Break correlations between Q-network and target

# Model: Train this agent on Atari 2600 games



- The **input** to the neural network consists of an **84x84x4 image** produced by the pre-processing map  $\phi$
- Input state is stack of raw pixels from **last 4 frames**

# Model

- State: Sequences of action and observation

$$s_t = x_1, a_1, x_2, \dots, a_{t-1}, x_t,$$

- Action: Legal game action set

$$\mathcal{A} = \{1, \dots, K\}.$$

- Reward: Change in game score

# How to train DQN

Loss function :

$$\mathcal{L}_i(\theta_i) = \mathbb{E}_{s,a,r,s' \sim \mathcal{D}} \left[ \left( r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta_i) \right)^2 \right]$$

Differentiating the loss function w.r.t. the weights we arrive at following gradient :

$$\nabla_{\theta_i} \mathcal{L}_i(\theta_i) = \mathbb{E}_{s,a,r,s' \sim \mathcal{D}} \left[ \left( r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta_i) \right) \nabla_{\theta_i} Q(s, a; \theta_i) \right]$$

Do gradient descent:

$$\theta_{i+1} = \theta_i + \alpha \cdot \nabla_{\theta_i} \mathcal{L}_i(\theta_i)$$

# Algorithm

**Algorithm 1: deep Q-learning with experience replay.**

Initialize replay memory  $D$  to capacity  $N$

Initialize action-value function  $Q$  with random weights  $\theta$

Initialize target action-value function  $\hat{Q}$  with weights  $\theta^- = \theta$

**For** episode = 1,  $M$  **do**

    Initialize sequence  $s_1 = \{x_1\}$  and preprocessed sequence  $\phi_1 = \phi(s_1)$

**For**  $t = 1, T$  **do**

        With probability  $\varepsilon$  select a random action  $a_t$

        otherwise select  $a_t = \operatorname{argmax}_a Q(\phi(s_t), a; \theta)$

        Execute action  $a_t$  in emulator and observe reward  $r_t$  and image  $x_{t+1}$

        Set  $s_{t+1} = s_t, a_t, x_{t+1}$  and preprocess  $\phi_{t+1} = \phi(s_{t+1})$

        Store transition  $(\phi_t, a_t, r_t, \phi_{t+1})$  in  $D$

        Sample random minibatch of transitions  $(\phi_j, a_j, r_j, \phi_{j+1})$  from  $D$

        Set  $y_j = \begin{cases} r_j & \text{if episode terminates at step } j+1 \\ r_j + \gamma \max_{a'} \hat{Q}(\phi_{j+1}, a'; \theta^-) & \text{otherwise} \end{cases}$


        Perform a gradient descent step on  $(y_j - Q(\phi_j, a_j; \theta))^2$  with respect to the network parameters  $\theta$

        Every  $C$  steps reset  $\hat{Q} = Q$

**End For**

**End For**

# Evaluation



# Evaluation

## DQN

Game	With replay, with target Q	With replay, without target Q	Without replay, with target Q	Without replay, without target Q
Breakout	316.8	240.7	10.2	3.2
Enduro	1006.3	831.4	141.9	29.1
River Raid	7446.6	4102.8	2867.7	1453.0
Seaquest	2894.4	822.6	1003.0	275.8
Space Invaders	1088.9	826.3	373.2	302.0

# Conclusion

- A single architecture can successfully learn policy with only minimal prior knowledge
- And the successful integration of RL with deep neural network was mainly dependent on a replay algorithm
- However, games demanding more temporally extended planning strategy still are challenge for all existing agent including DQN
- Experience replay can be better