



Performance Isolation and Fairness for Multi-Tenant Cloud Storage

Zhang Zhizhong Oct 17, 2012

Setting: Shared Storage in the Cloud




Predictable Performance is Hard




Multiple co-located tenants \Rightarrow resource contention

Predictable Performance is Hard



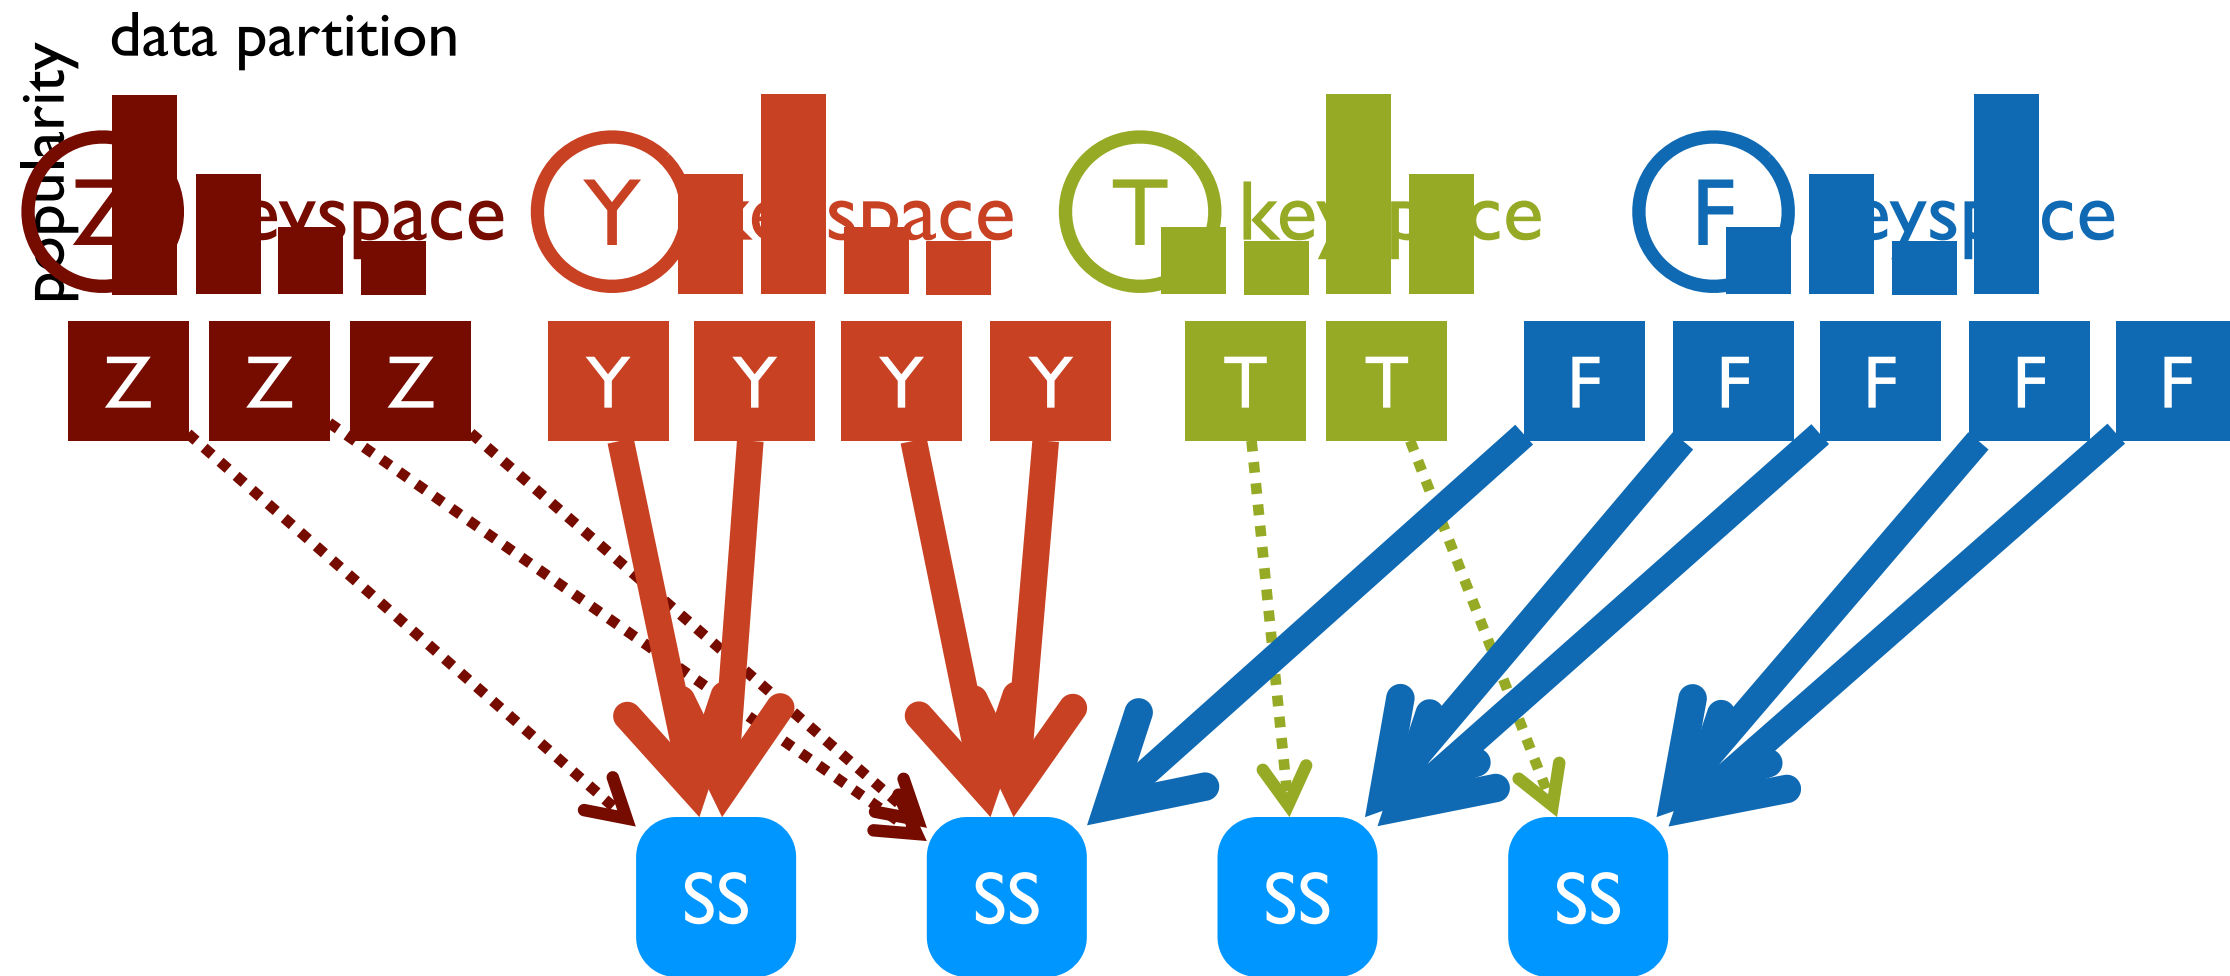
Multiple co-located tenants \Rightarrow resource contention

Predictable Performance is Hard



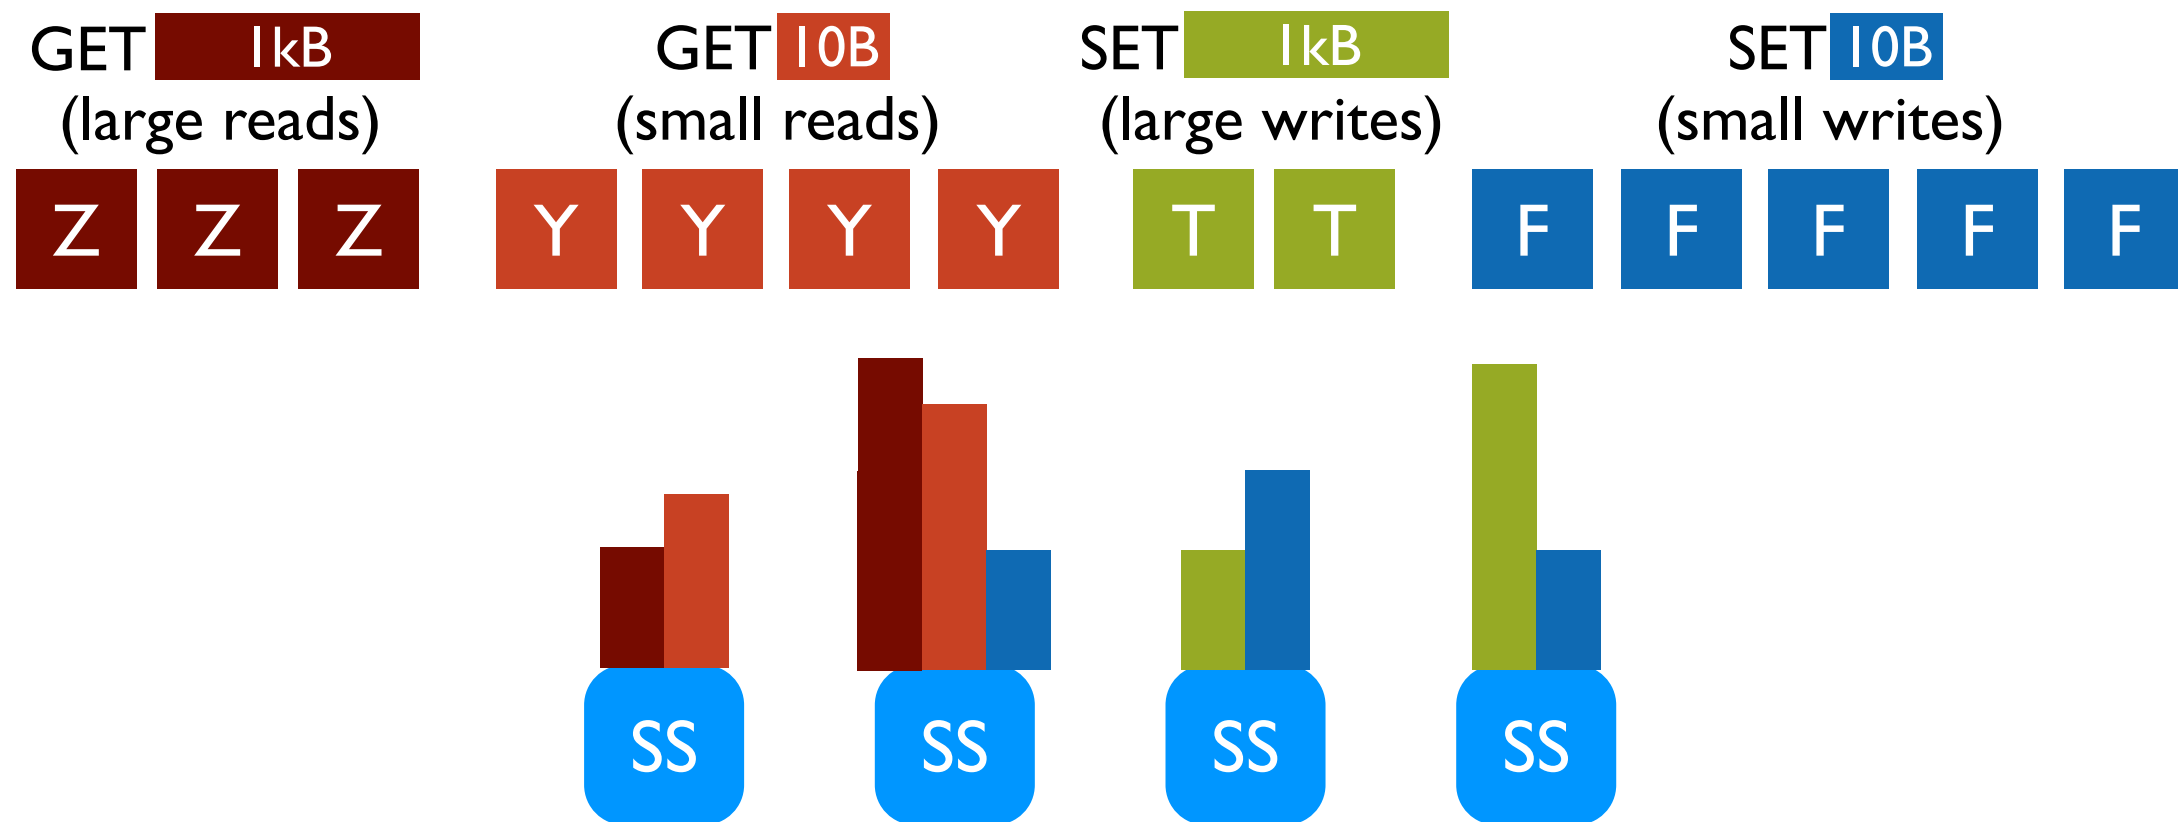
Multiple co-located tenants \Rightarrow resource contention
Distributed system \Rightarrow distributed resource allocation

Predictable Performance is Hard



Multiple co-located tenants \Rightarrow resource contention
Distributed system \Rightarrow distributed resource allocation

Predictable Performance is Hard



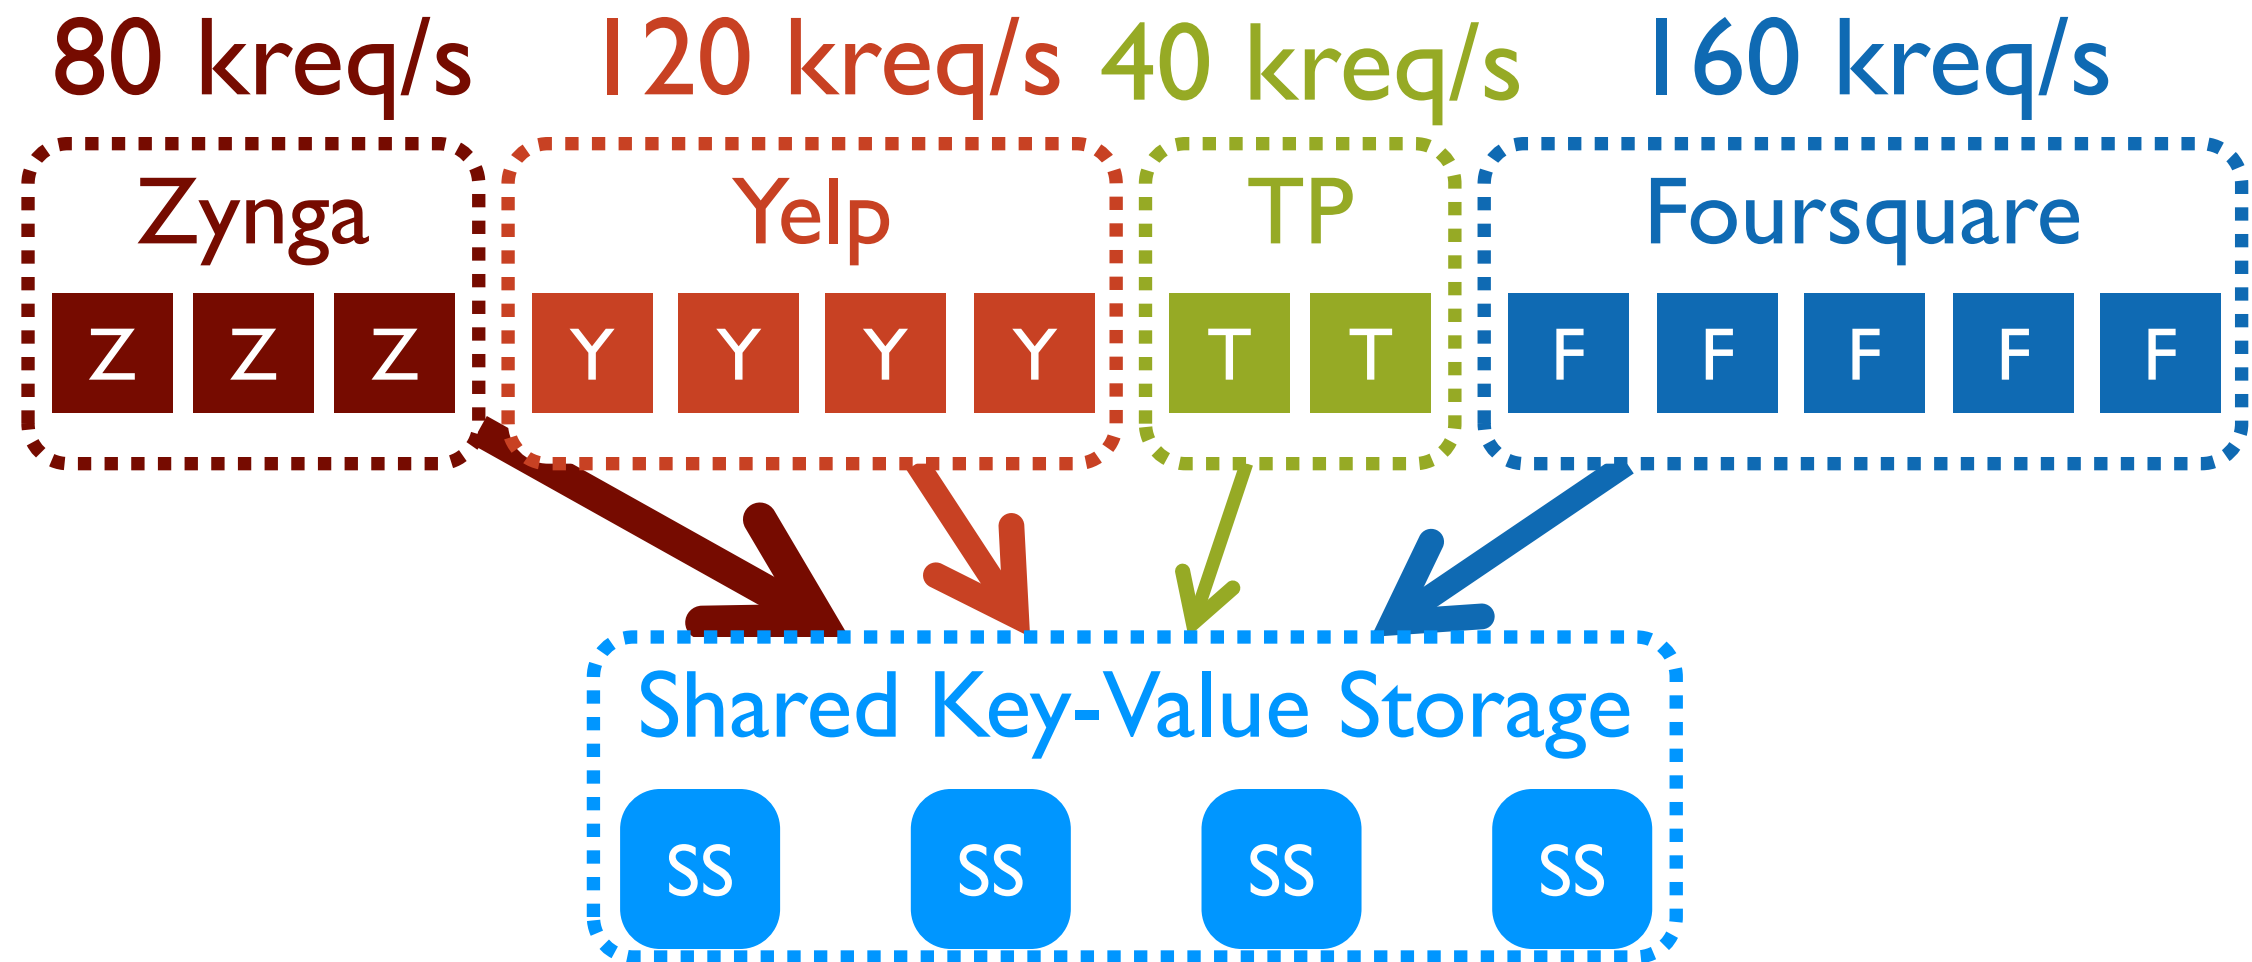
Multiple co-located tenants \Rightarrow resource contention

Distributed system \Rightarrow distributed resource allocation

Skewed object popularity \Rightarrow variable per-node demand

Disparate workloads \Rightarrow different bottleneck resources

Tenants Want System-wide Resource Guarantees



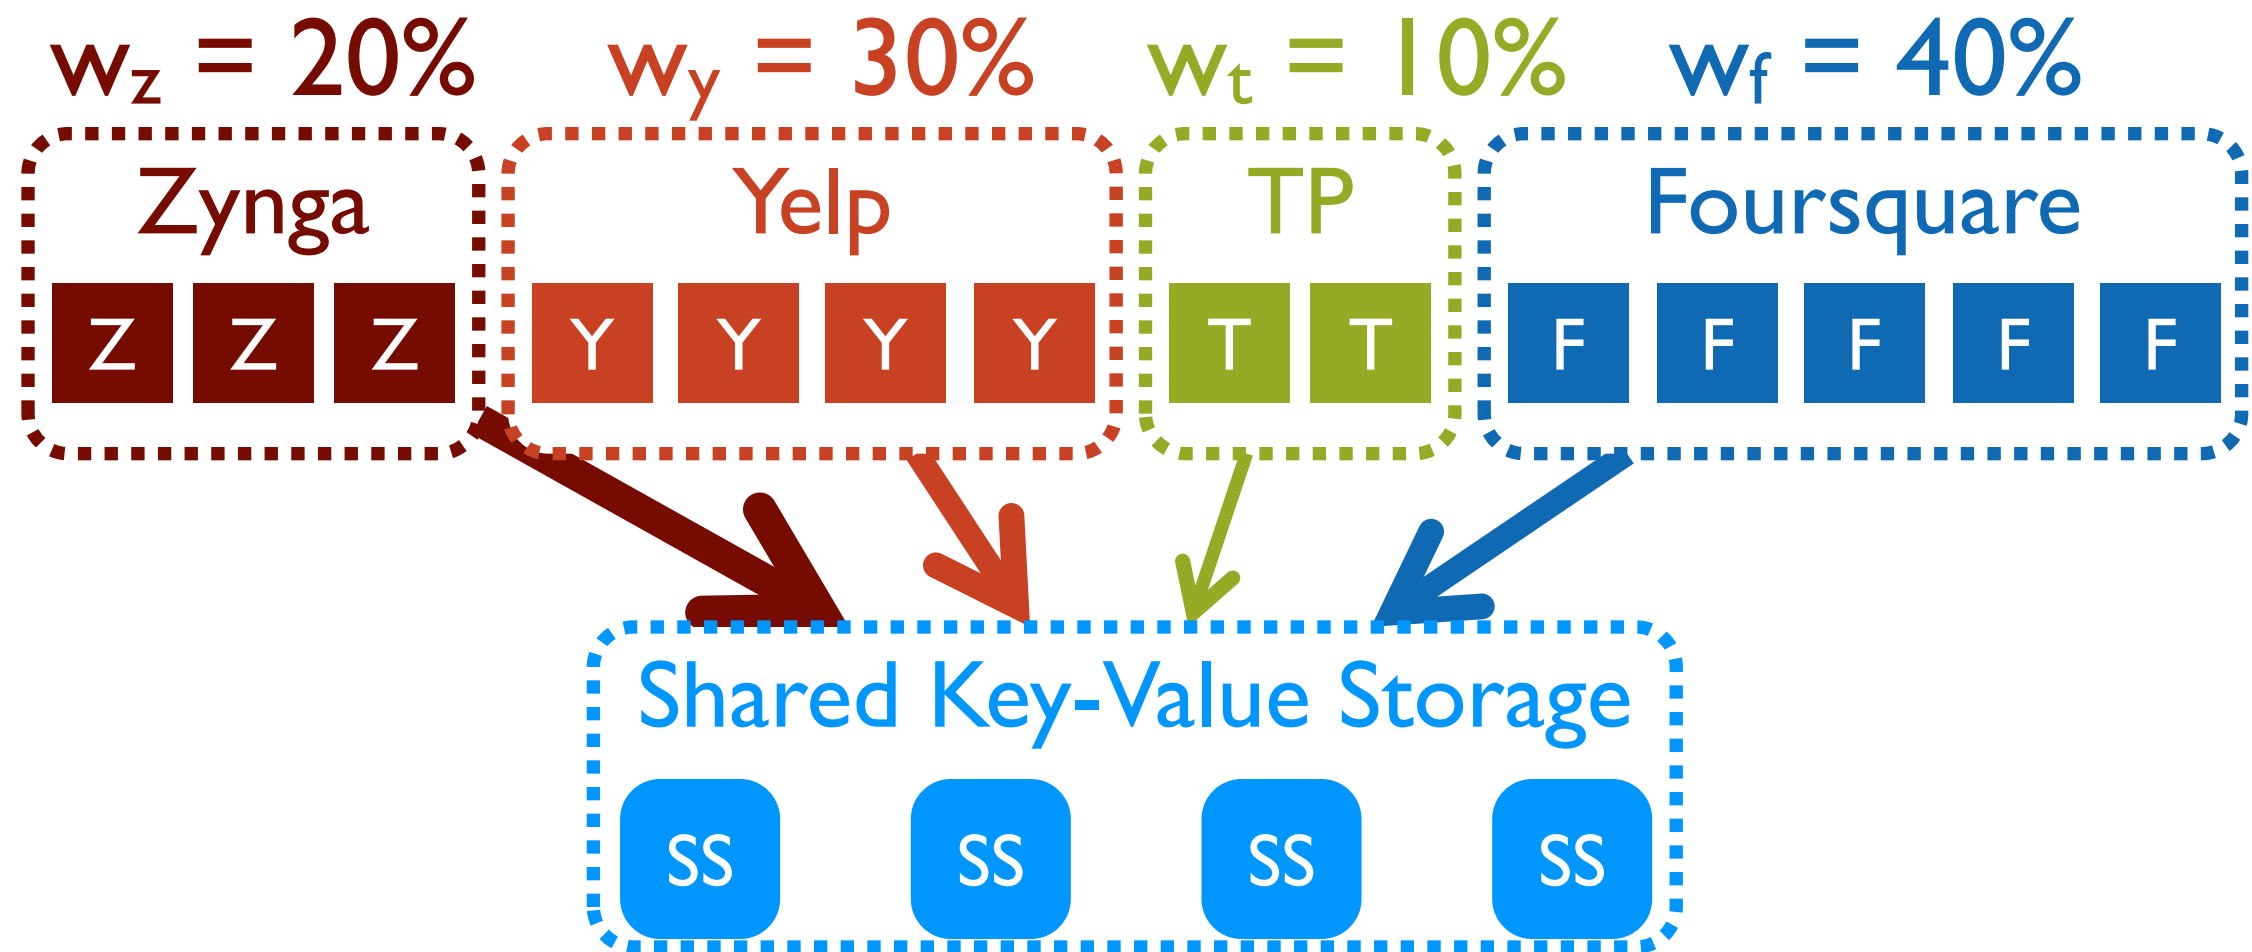
Multiple co-located tenants \Rightarrow resource contention

Distributed system \Rightarrow distributed resource allocation

Skewed object popularity \Rightarrow variable per-node demand

Disparate workloads \Rightarrow different bottleneck resources

Pisces Provides Weighted Fair-shares



Multiple co-located tenants \Rightarrow resource contention

Distributed system \Rightarrow distributed resource allocation

Skewed object popularity \Rightarrow variable per-node demand

Disparate workloads \Rightarrow different bottleneck resources

Pisces: Predictable Shared Cloud Storage

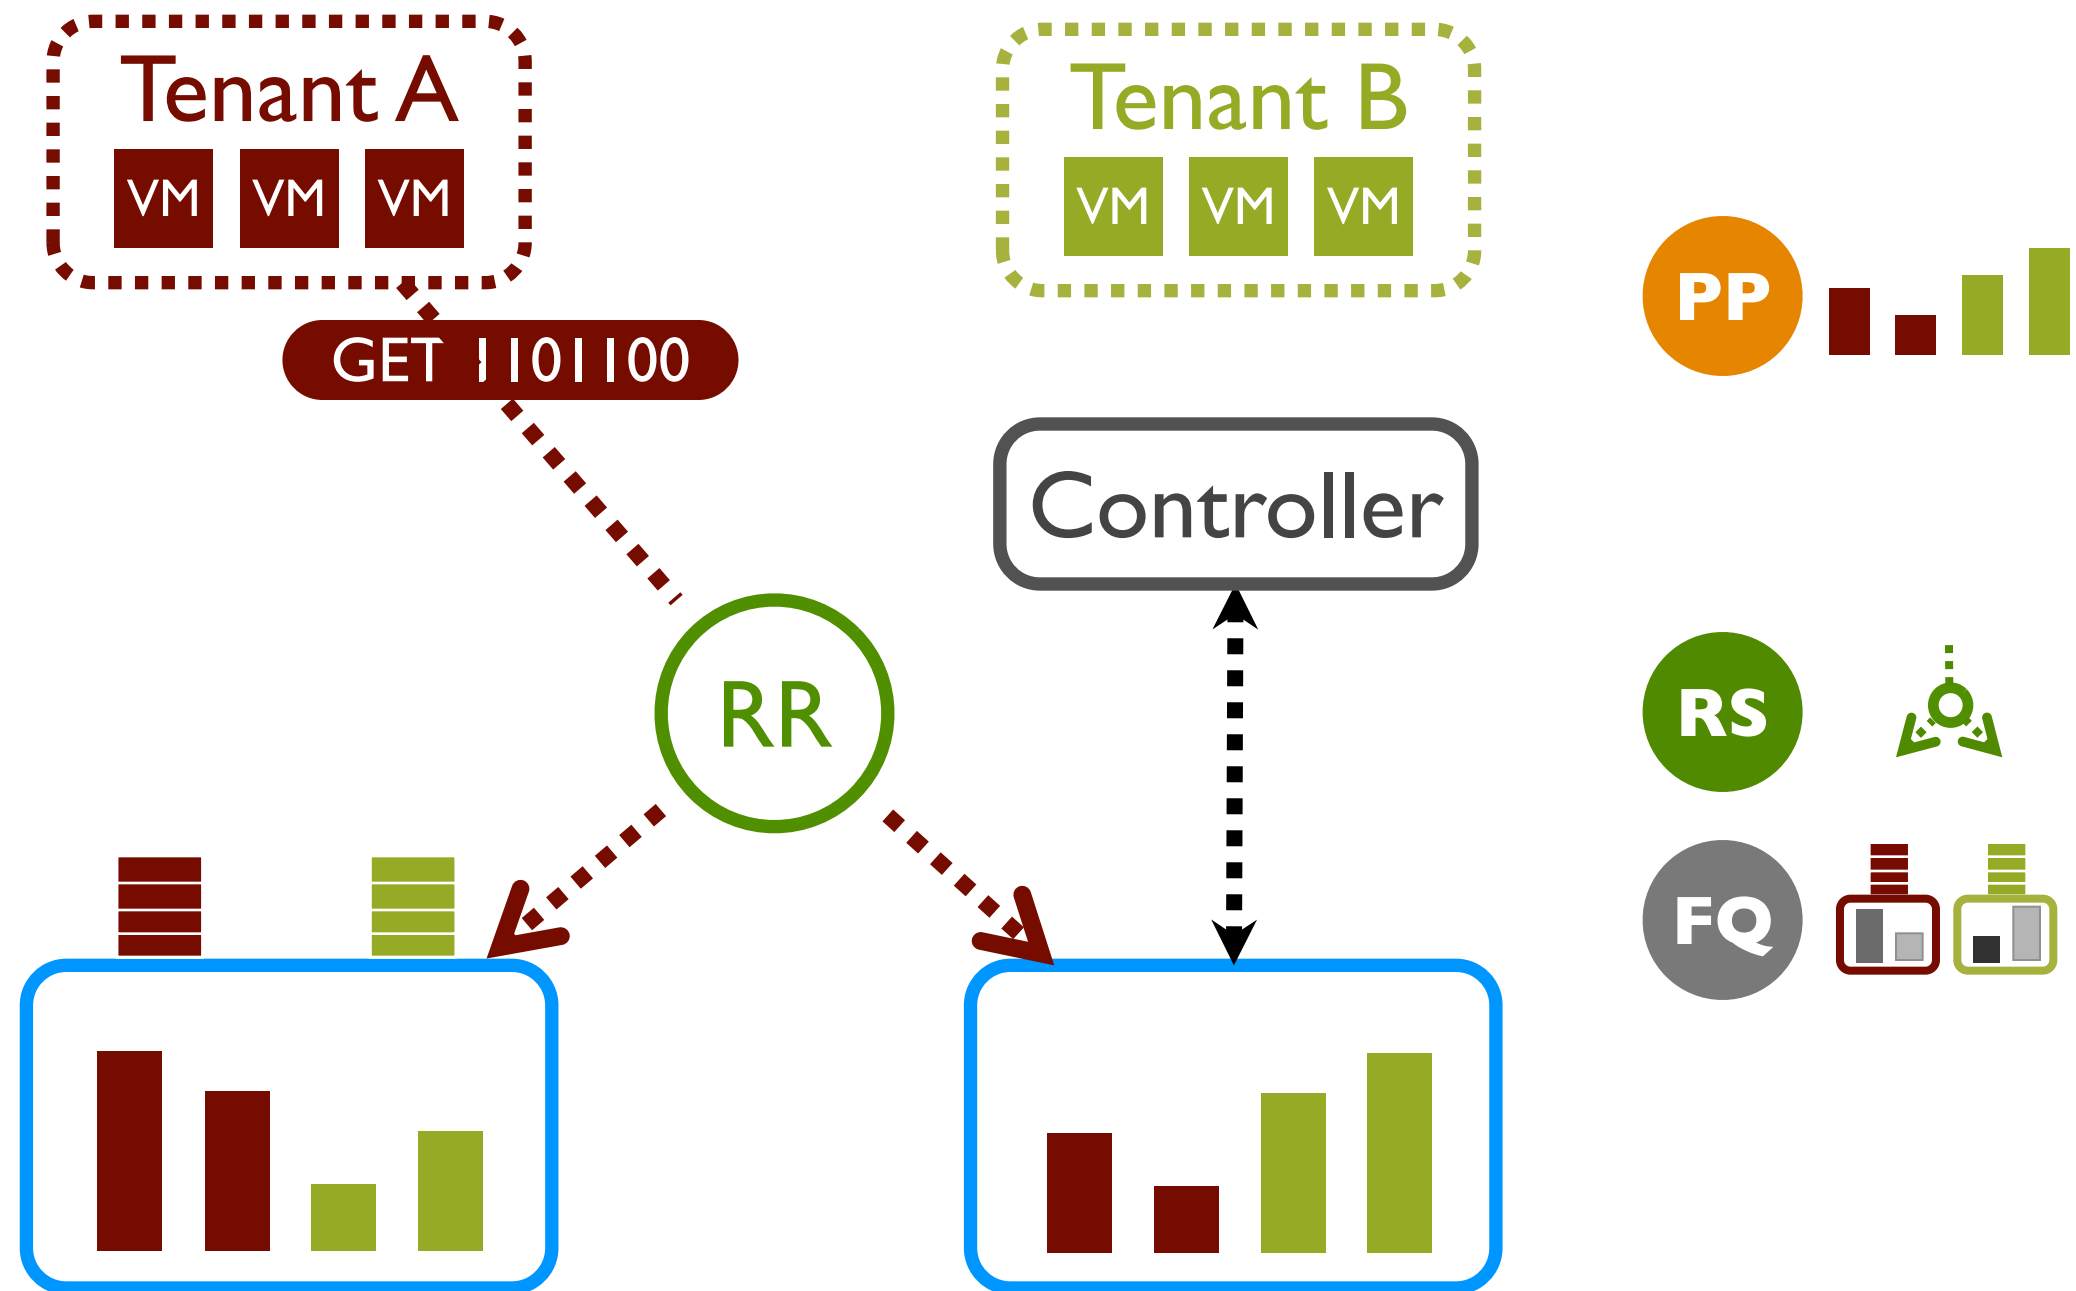
● Pisces

- Per-tenant max-min fair shares of system-wide resources
~ min guarantees, high utilization
- Arbitrary object popularity
- Different resource bottlenecks

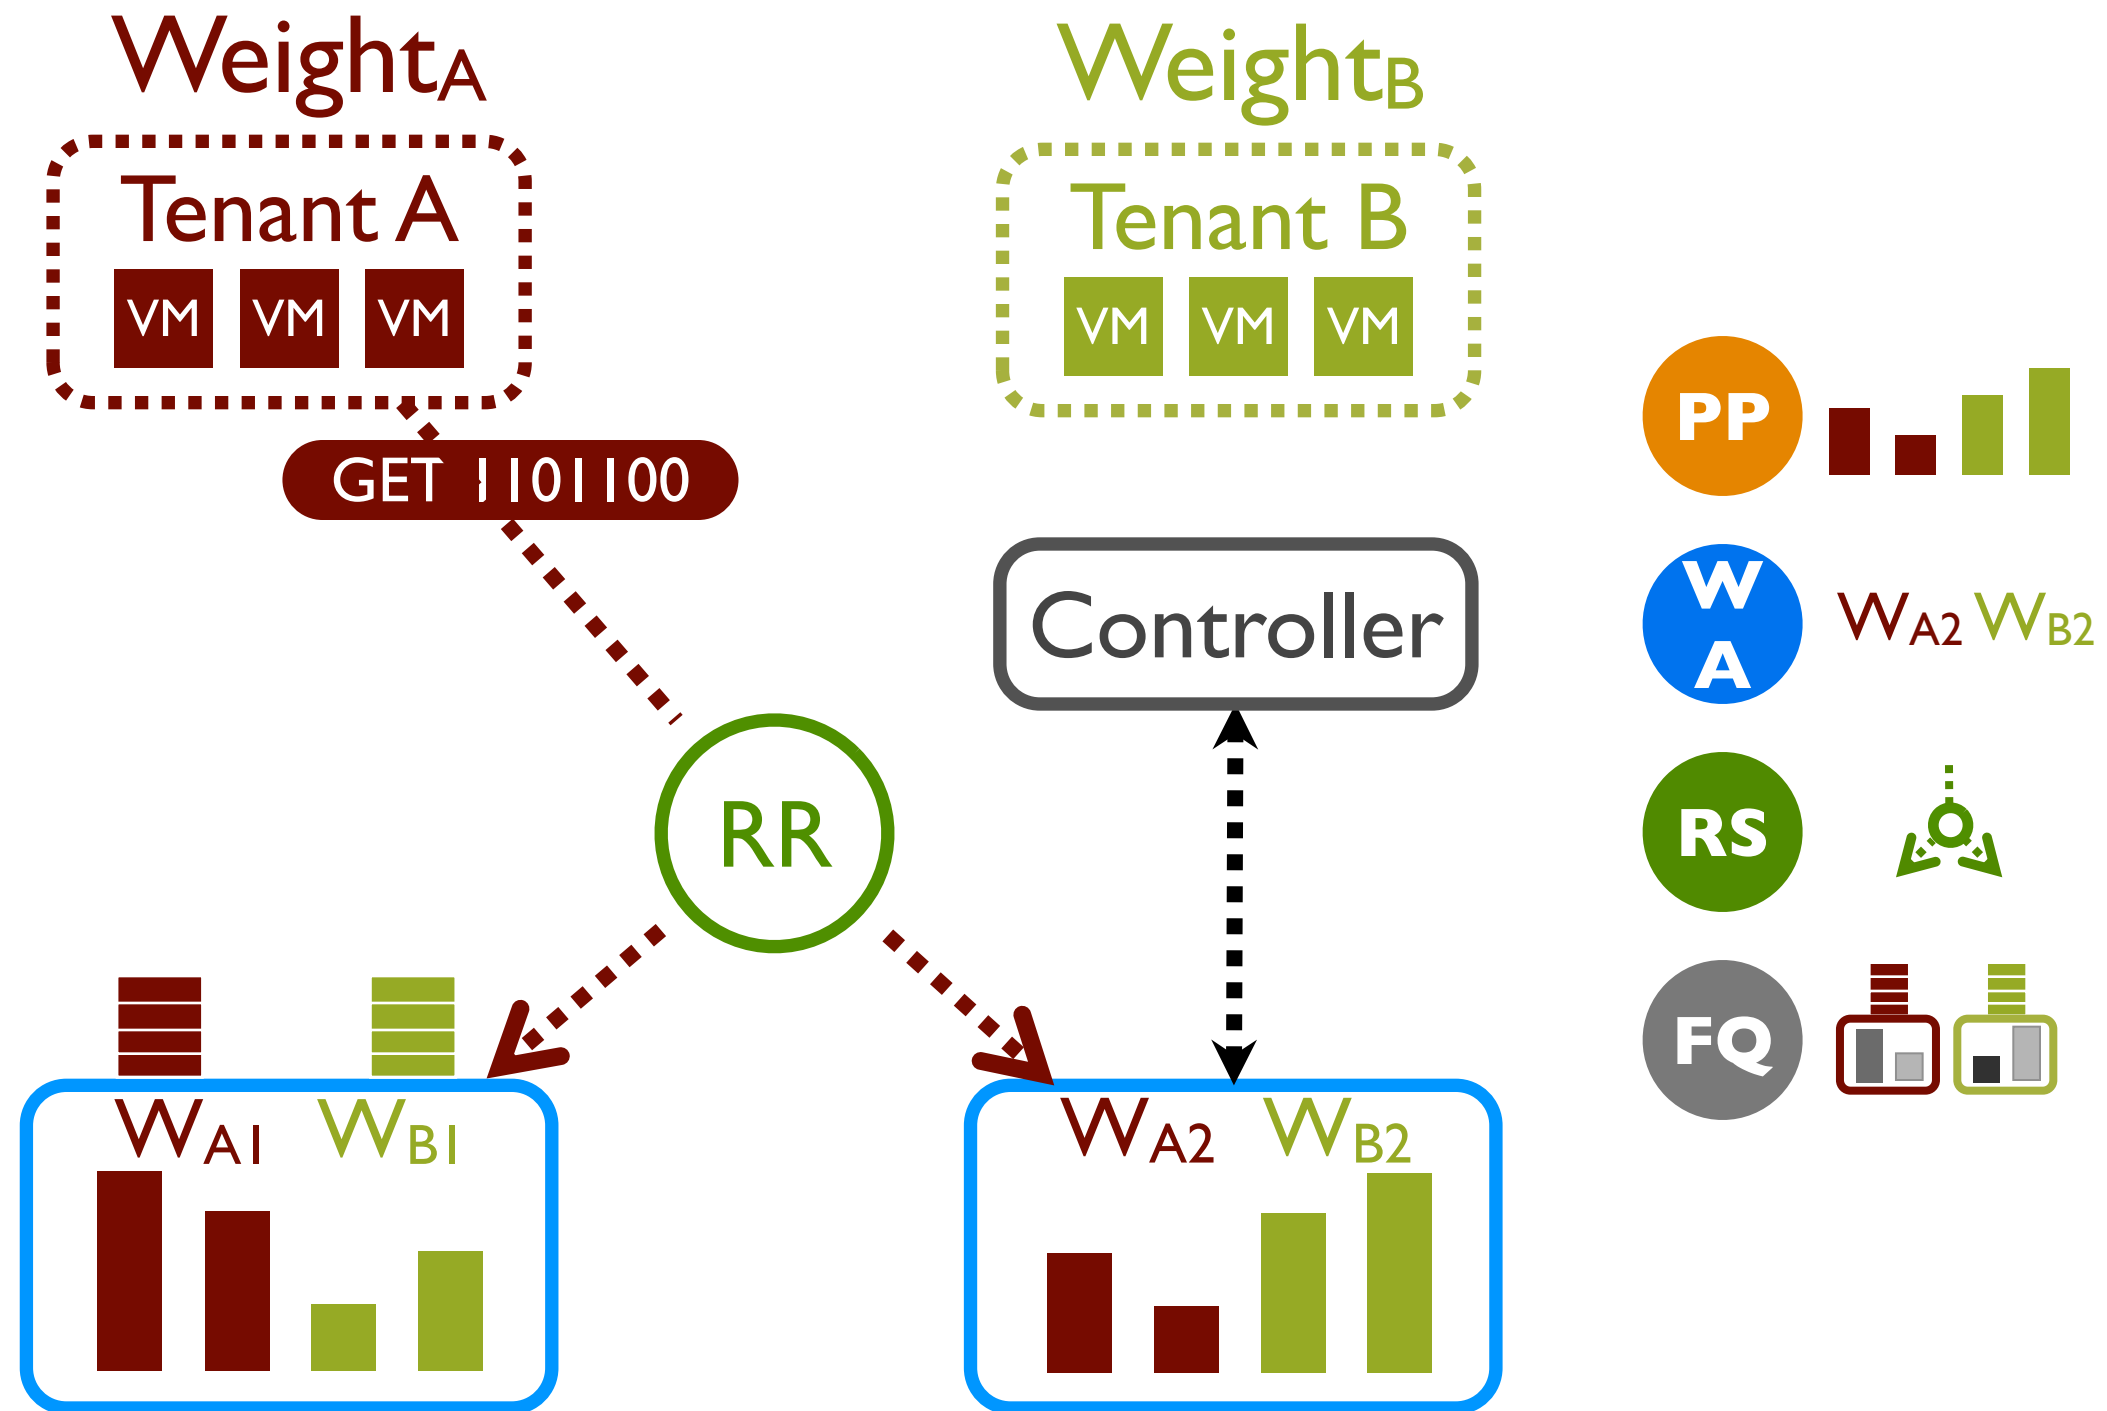
● Amazon DynamoDB

- Per-tenant provisioned rates
~ rate limited, non-work conserving
- Uniform object popularity
- Single resource (1kB requests)

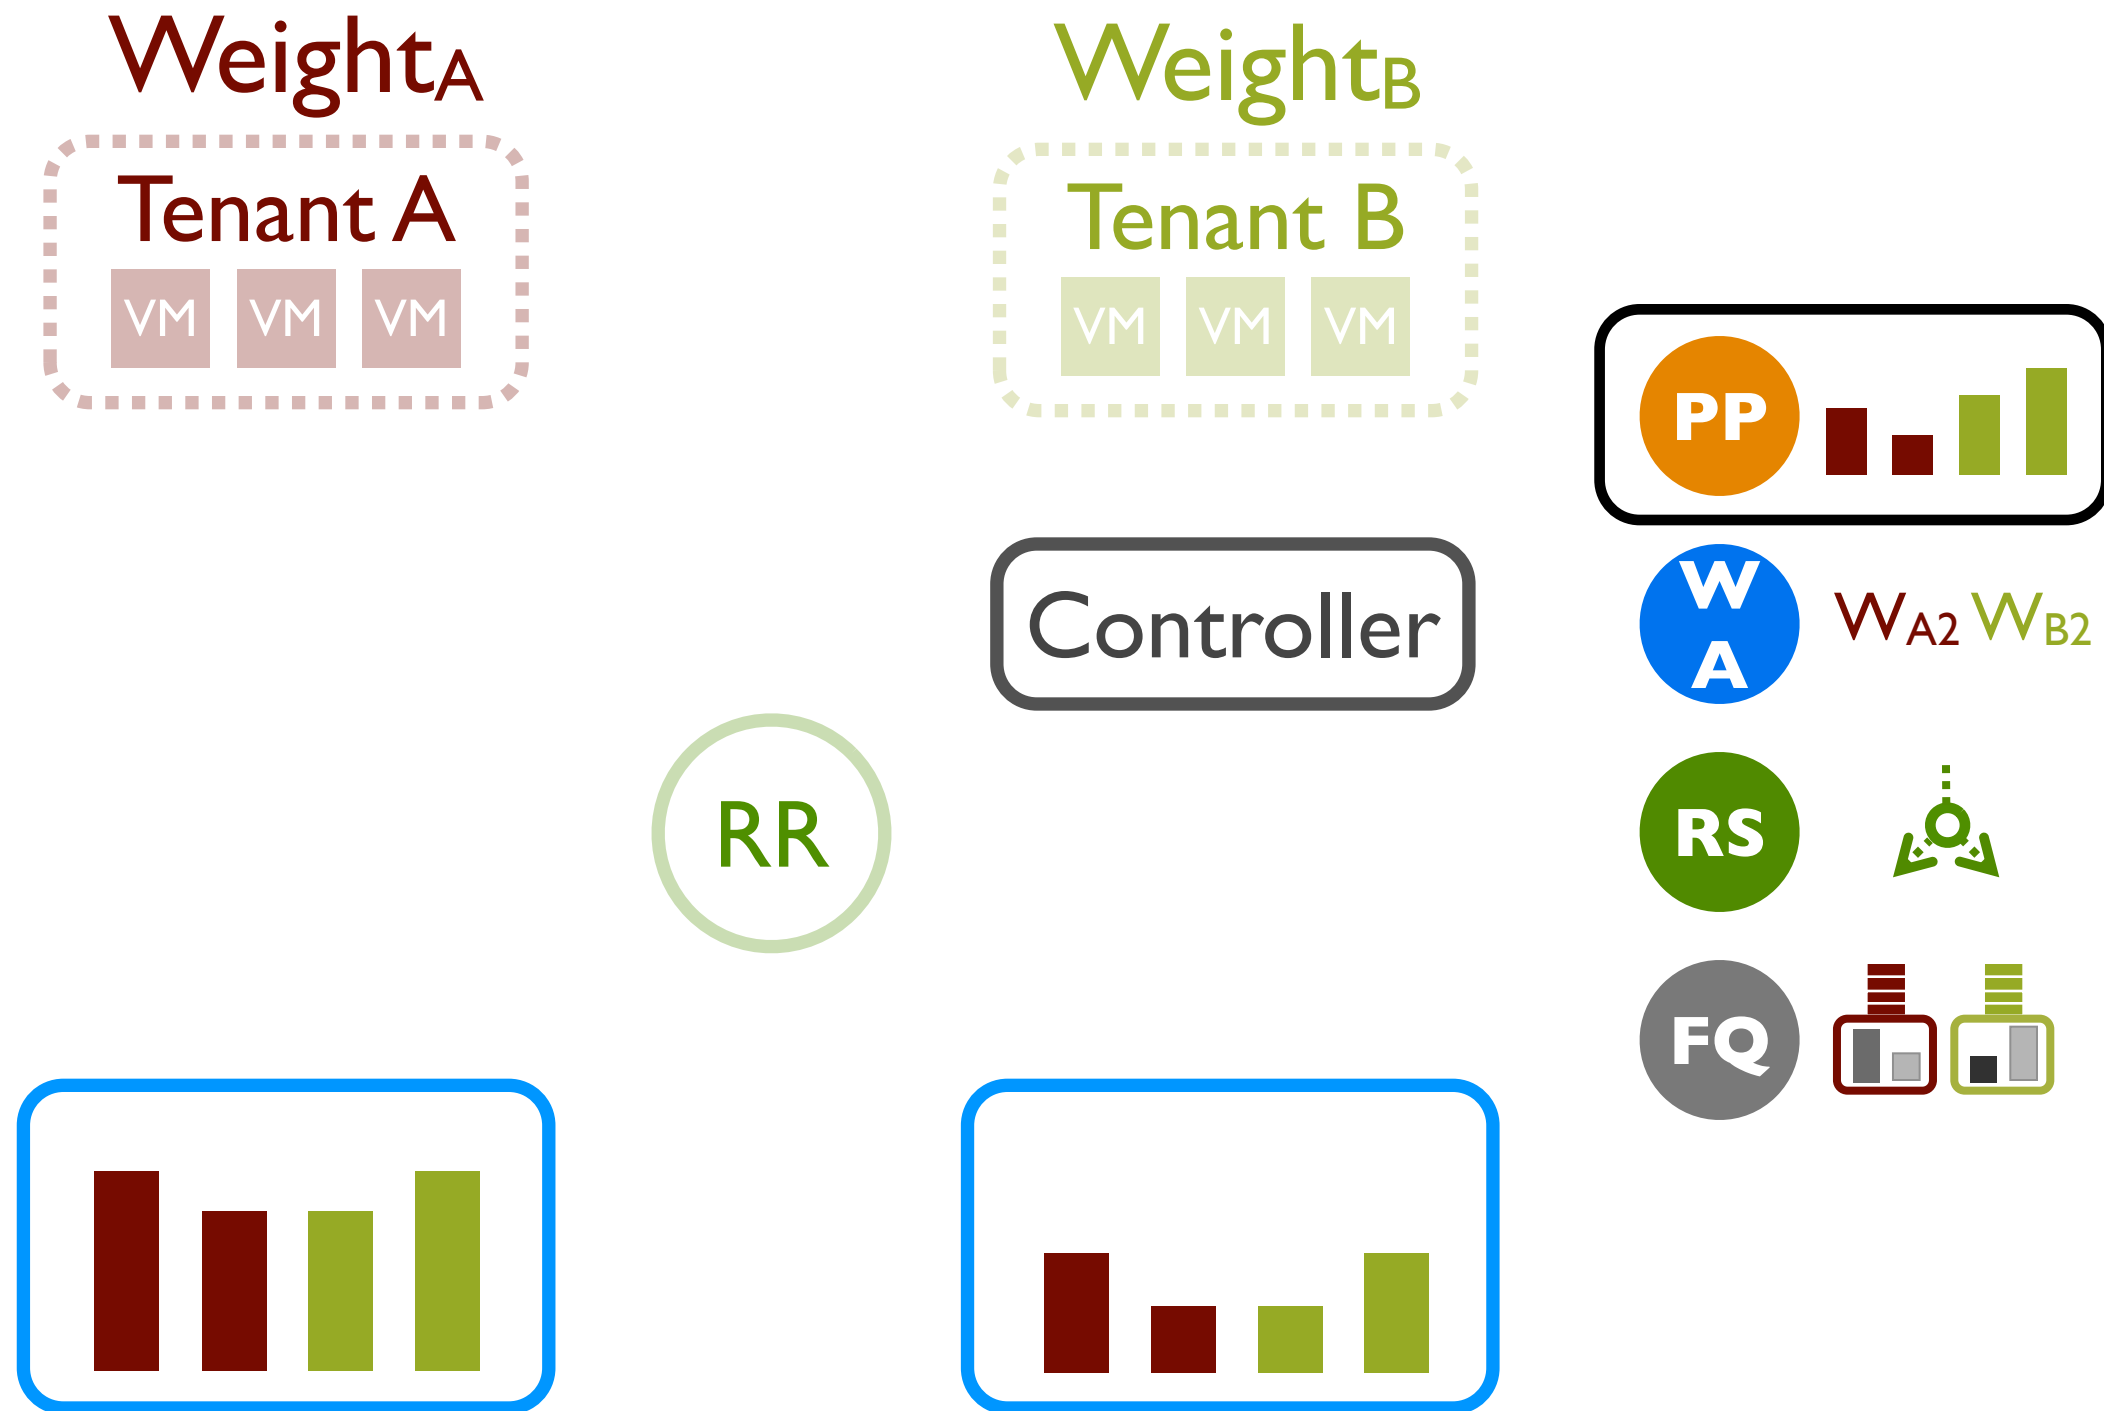
Predictable Multi-Tenant Key-Value Storage



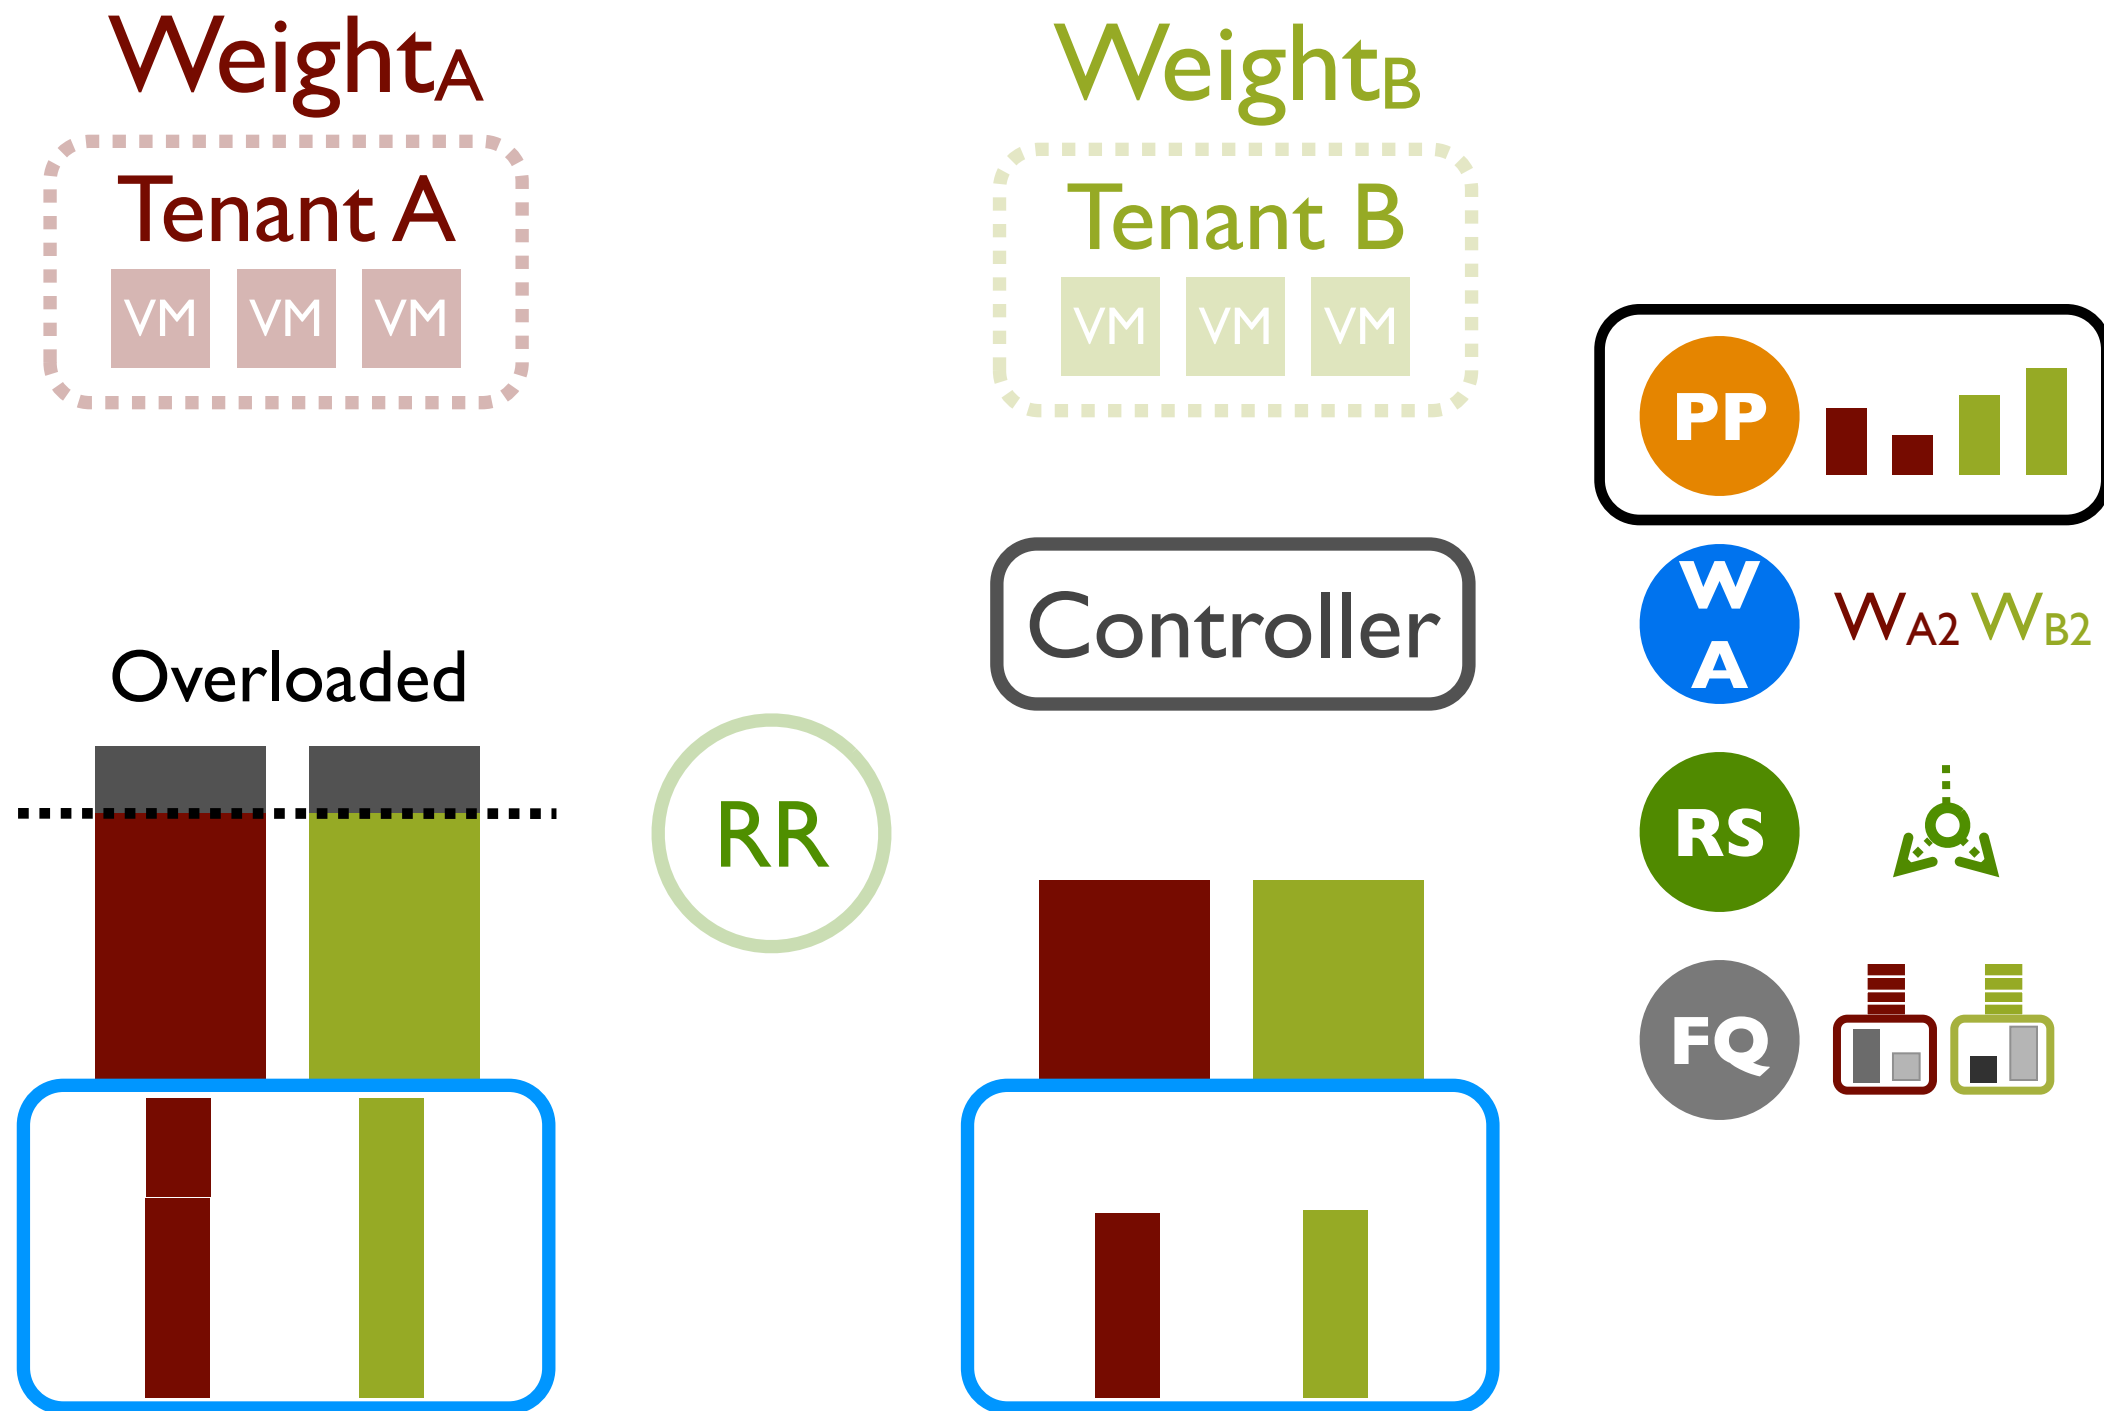
Predictable Multi-Tenant Key-Value Storage



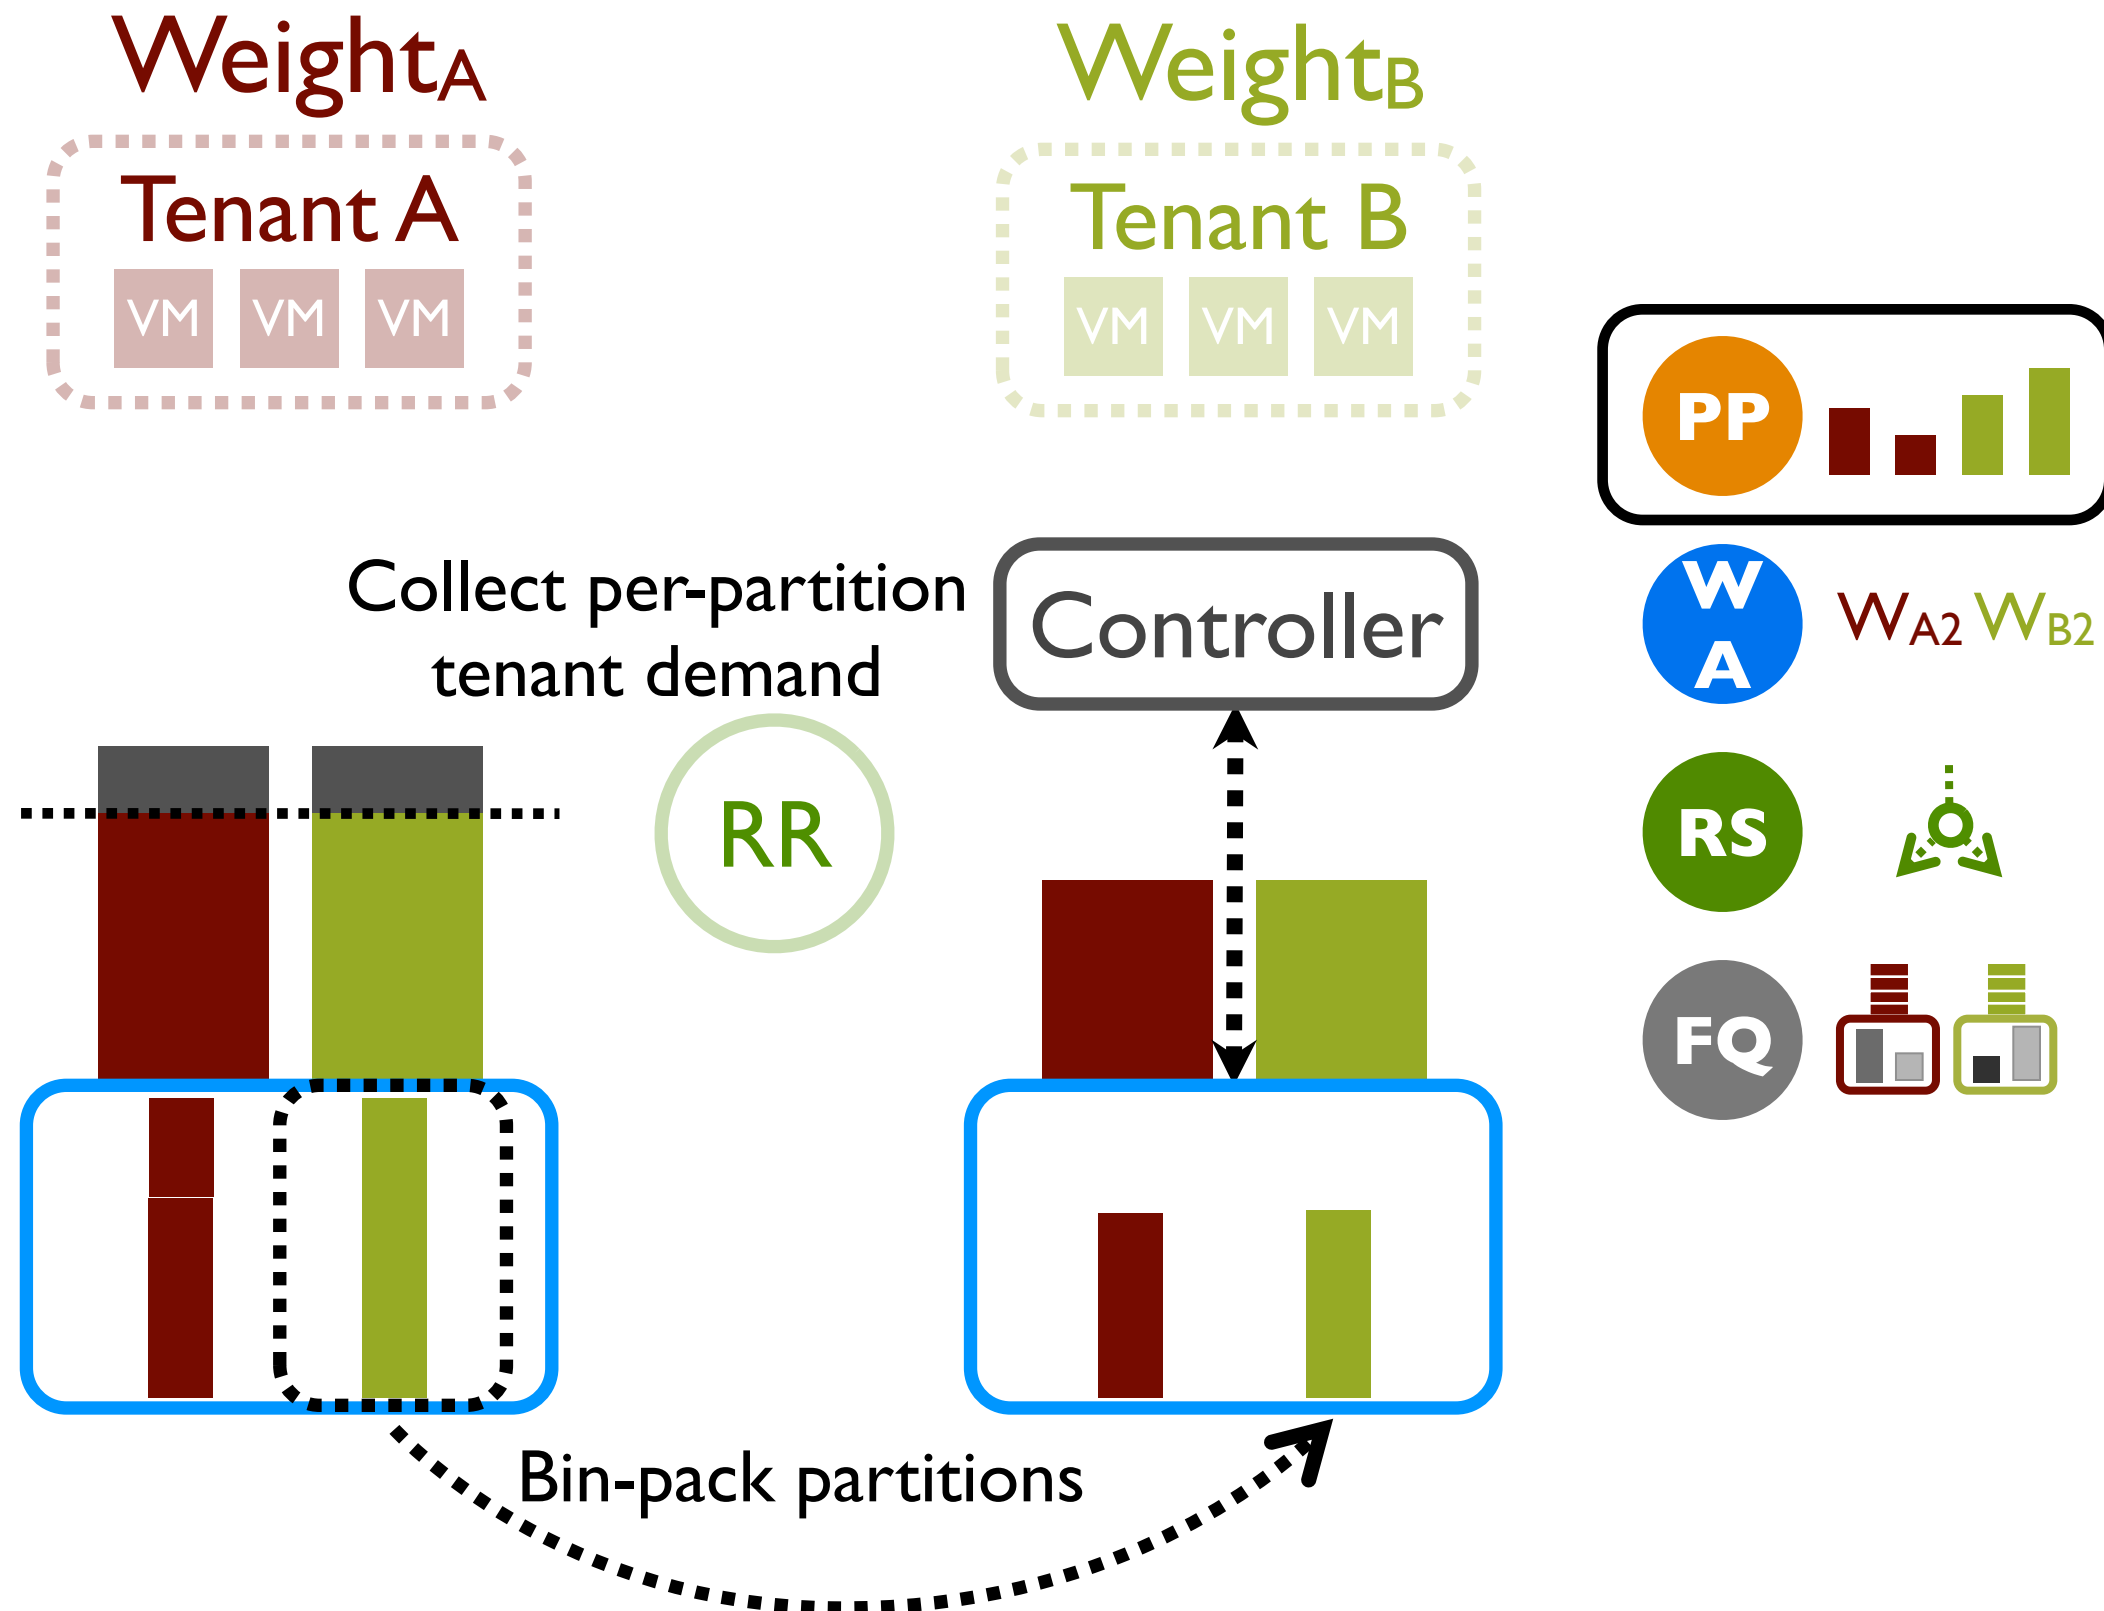
Strawman: Place Partitions Randomly



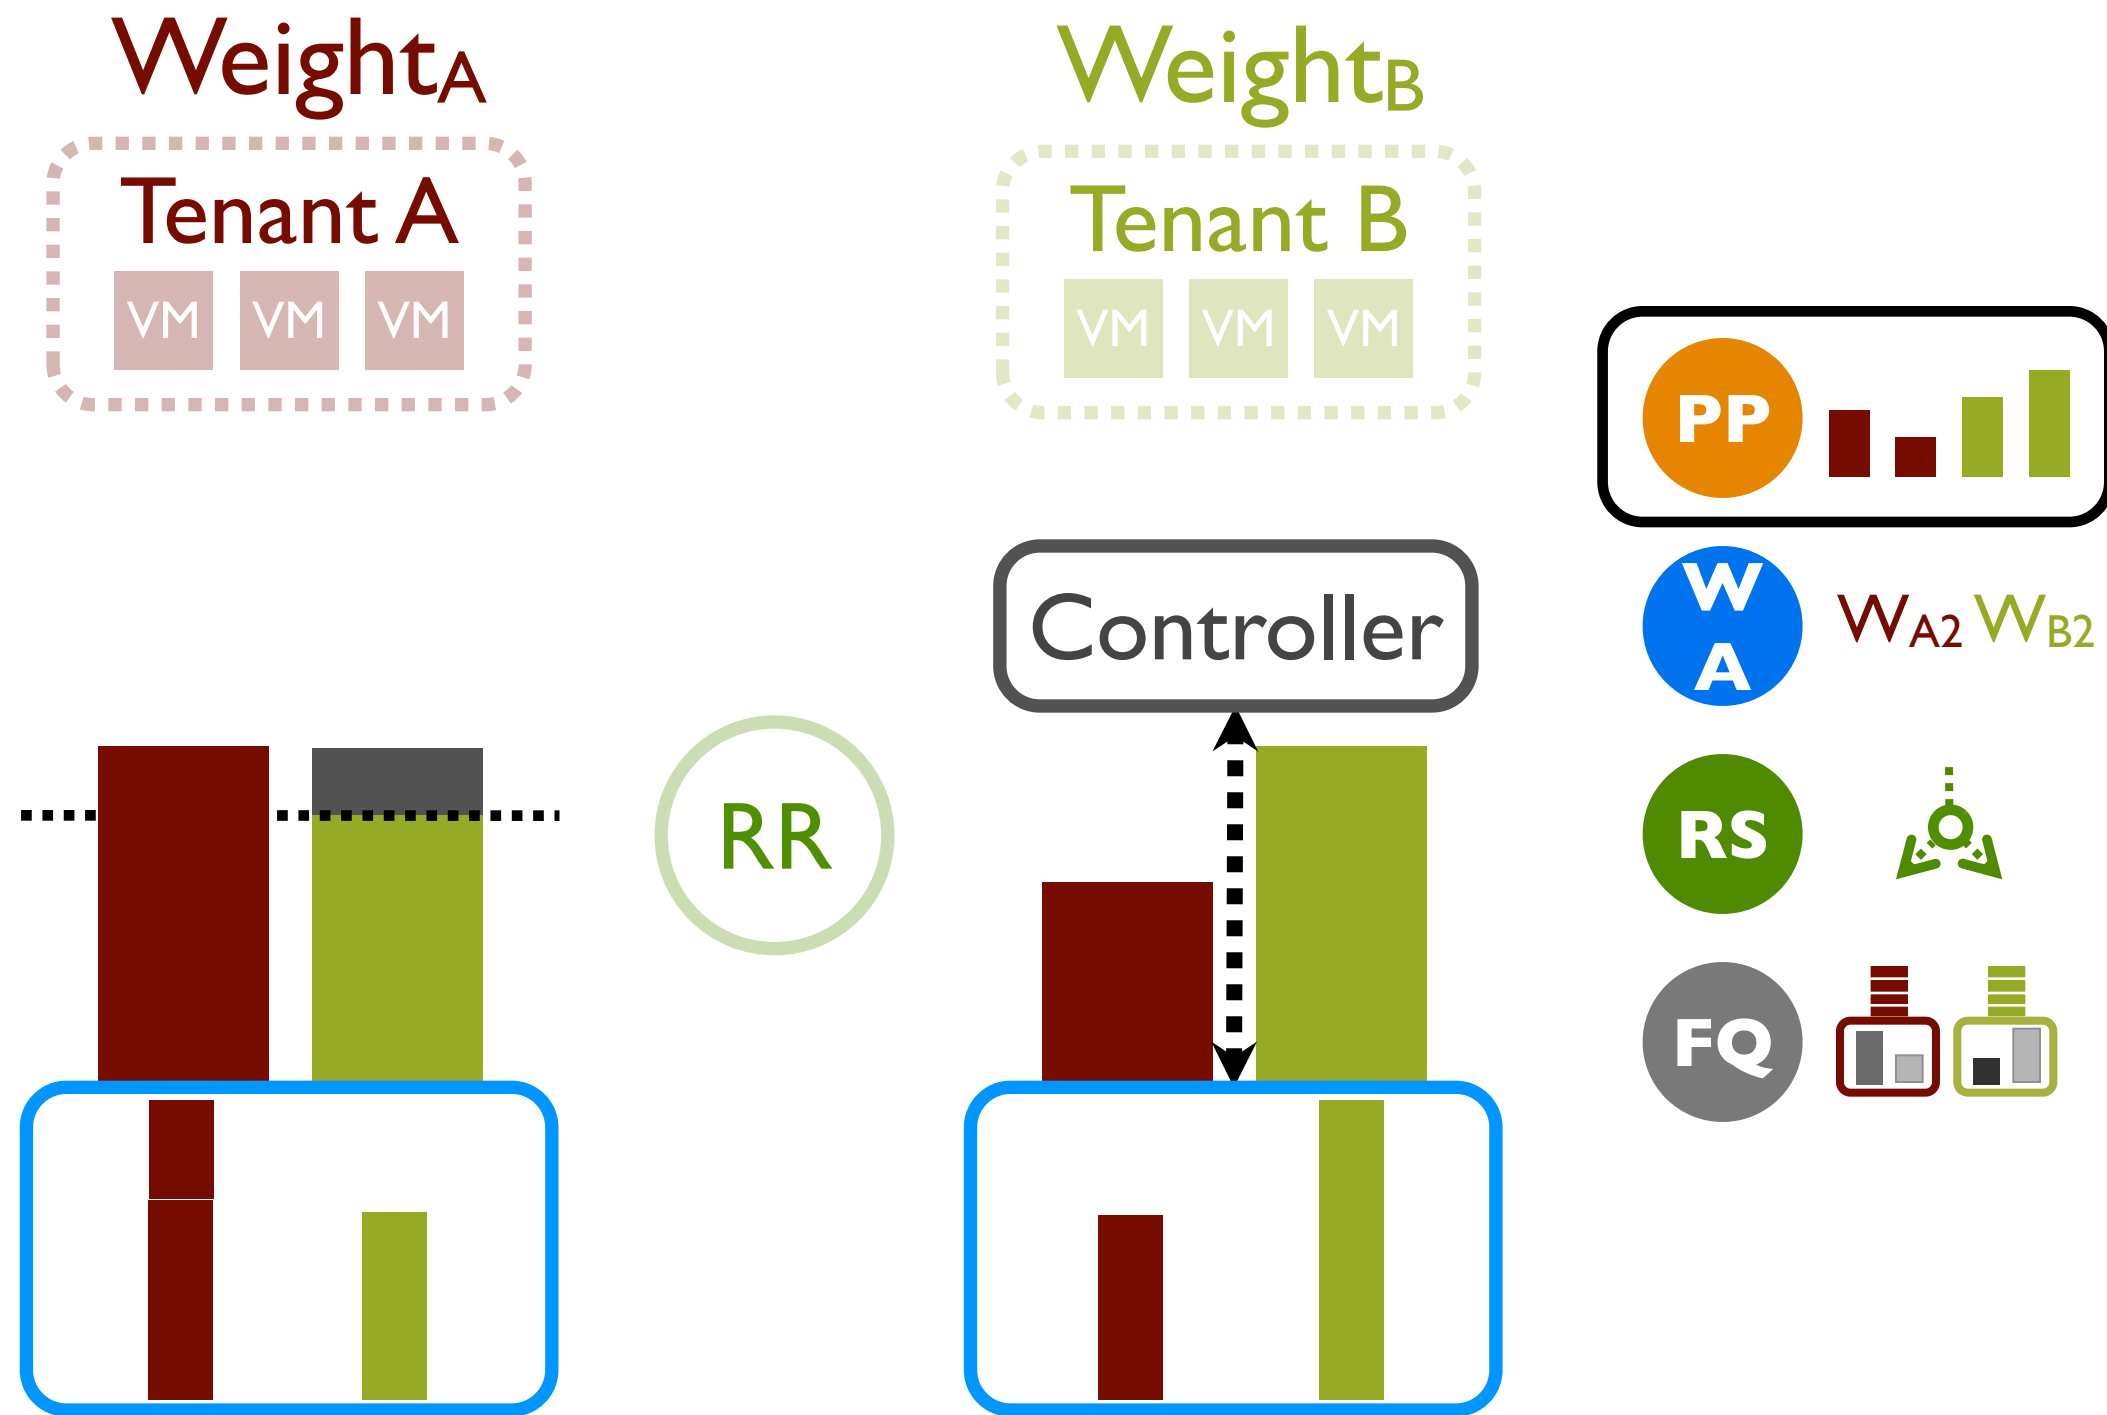
Strawman: Place Partitions Randomly



Pisces: Place Partitions By Fairness Constraints

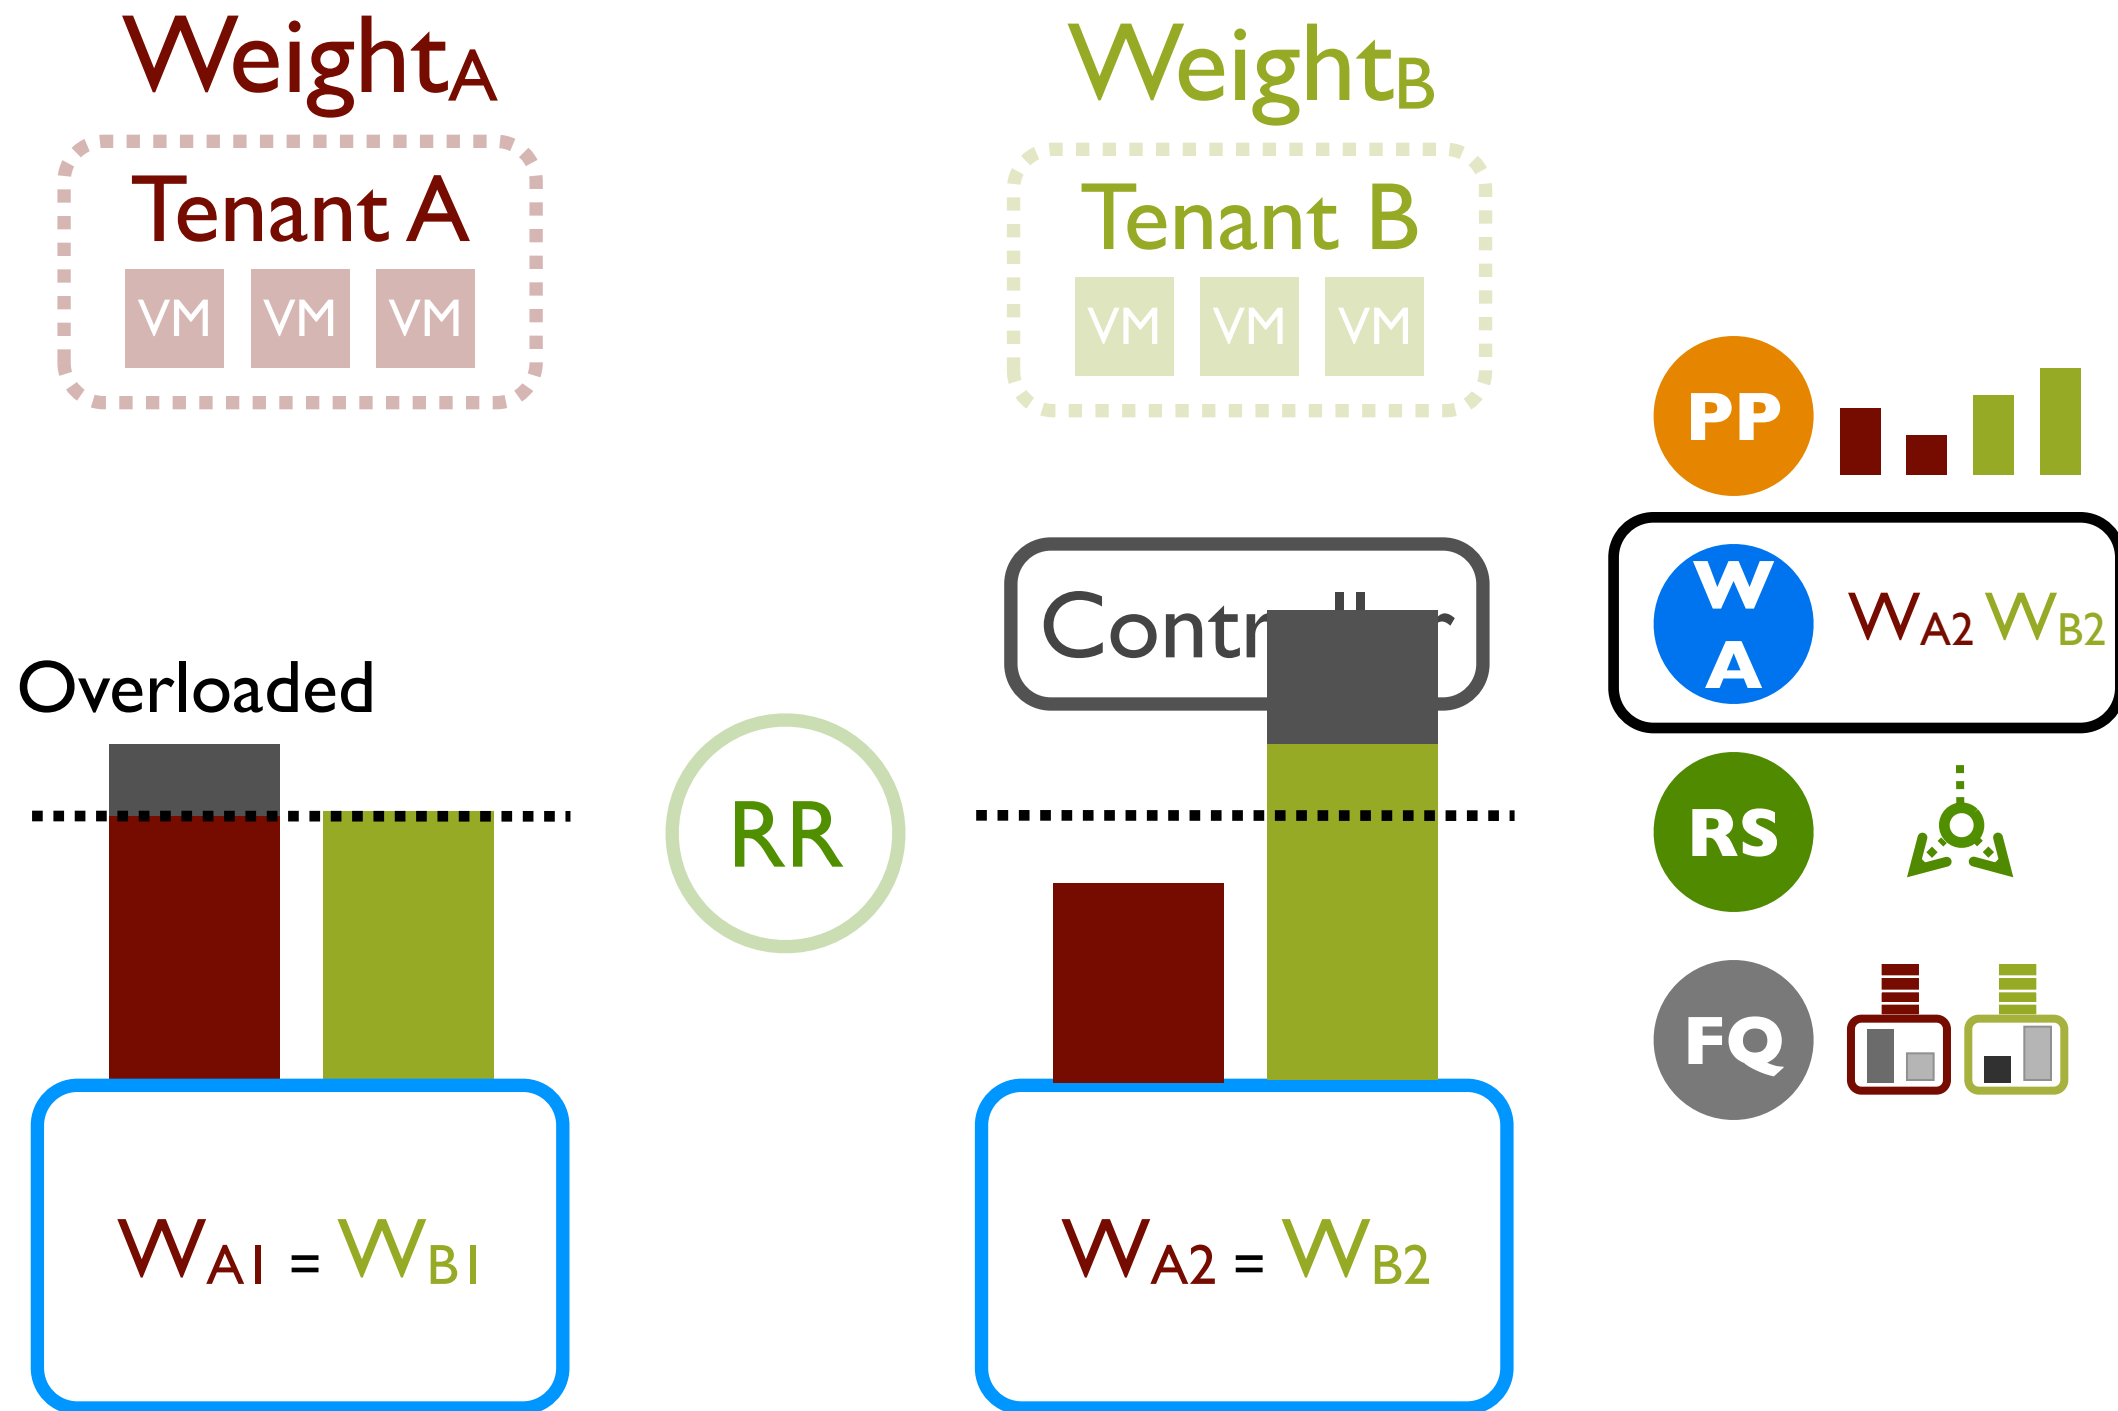


Pisces: Place Partitions By Fairness Constraints

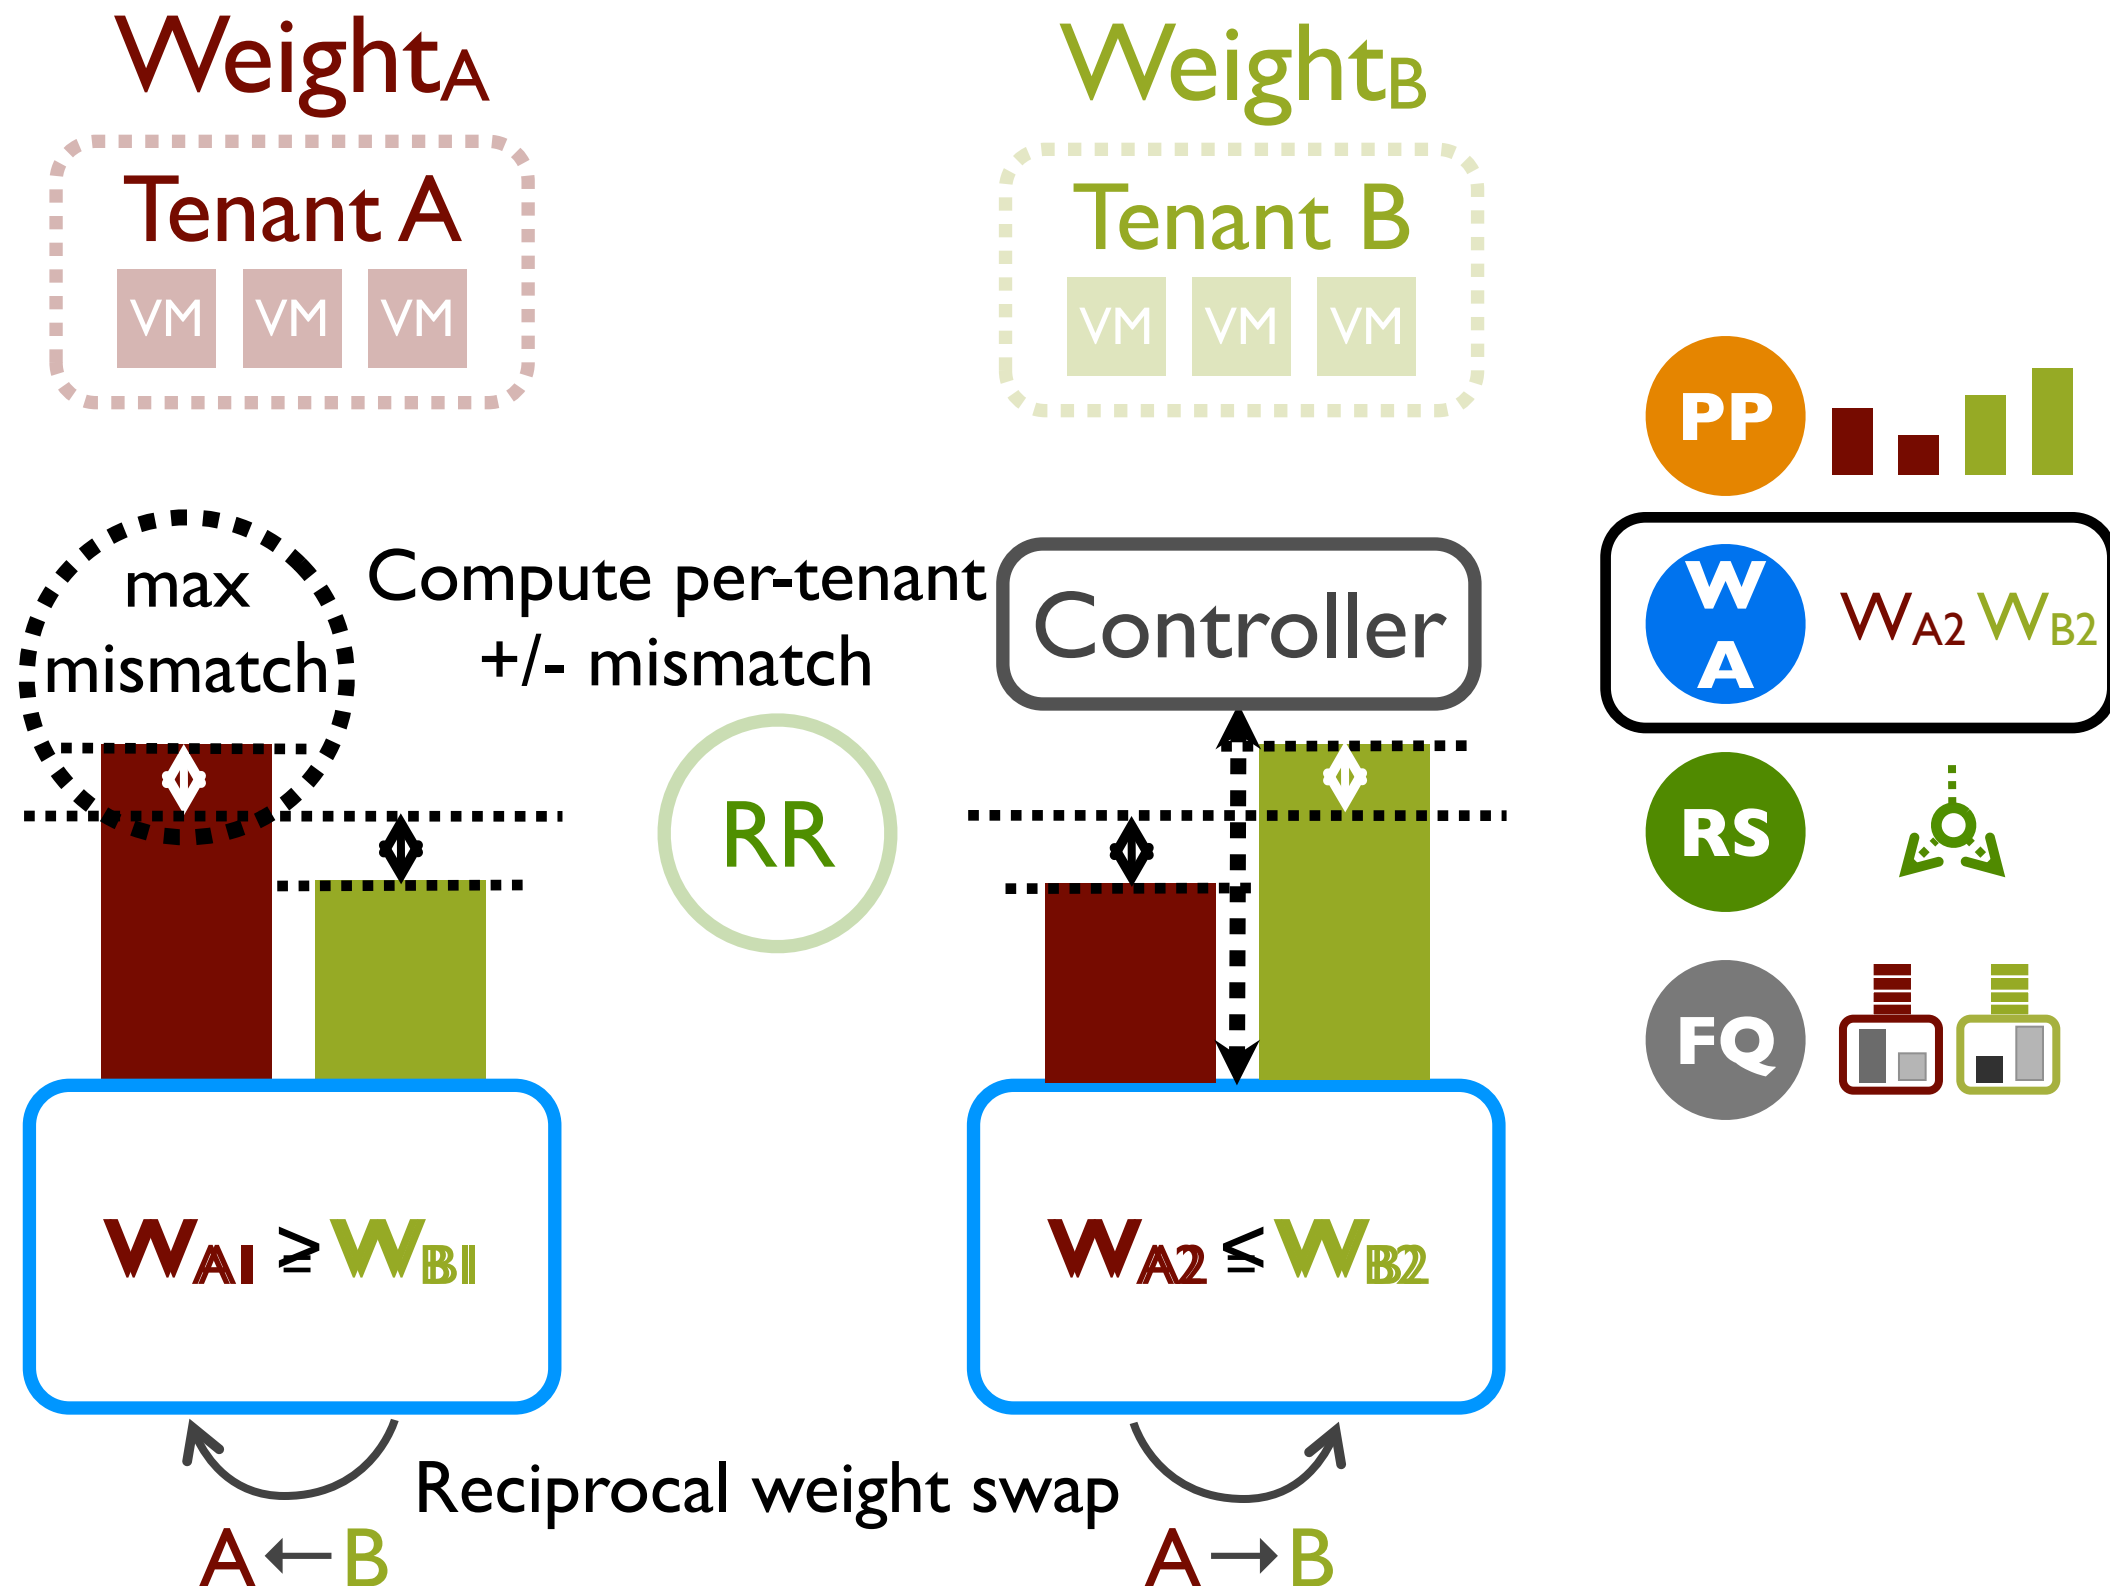


Results in feasible partition placement

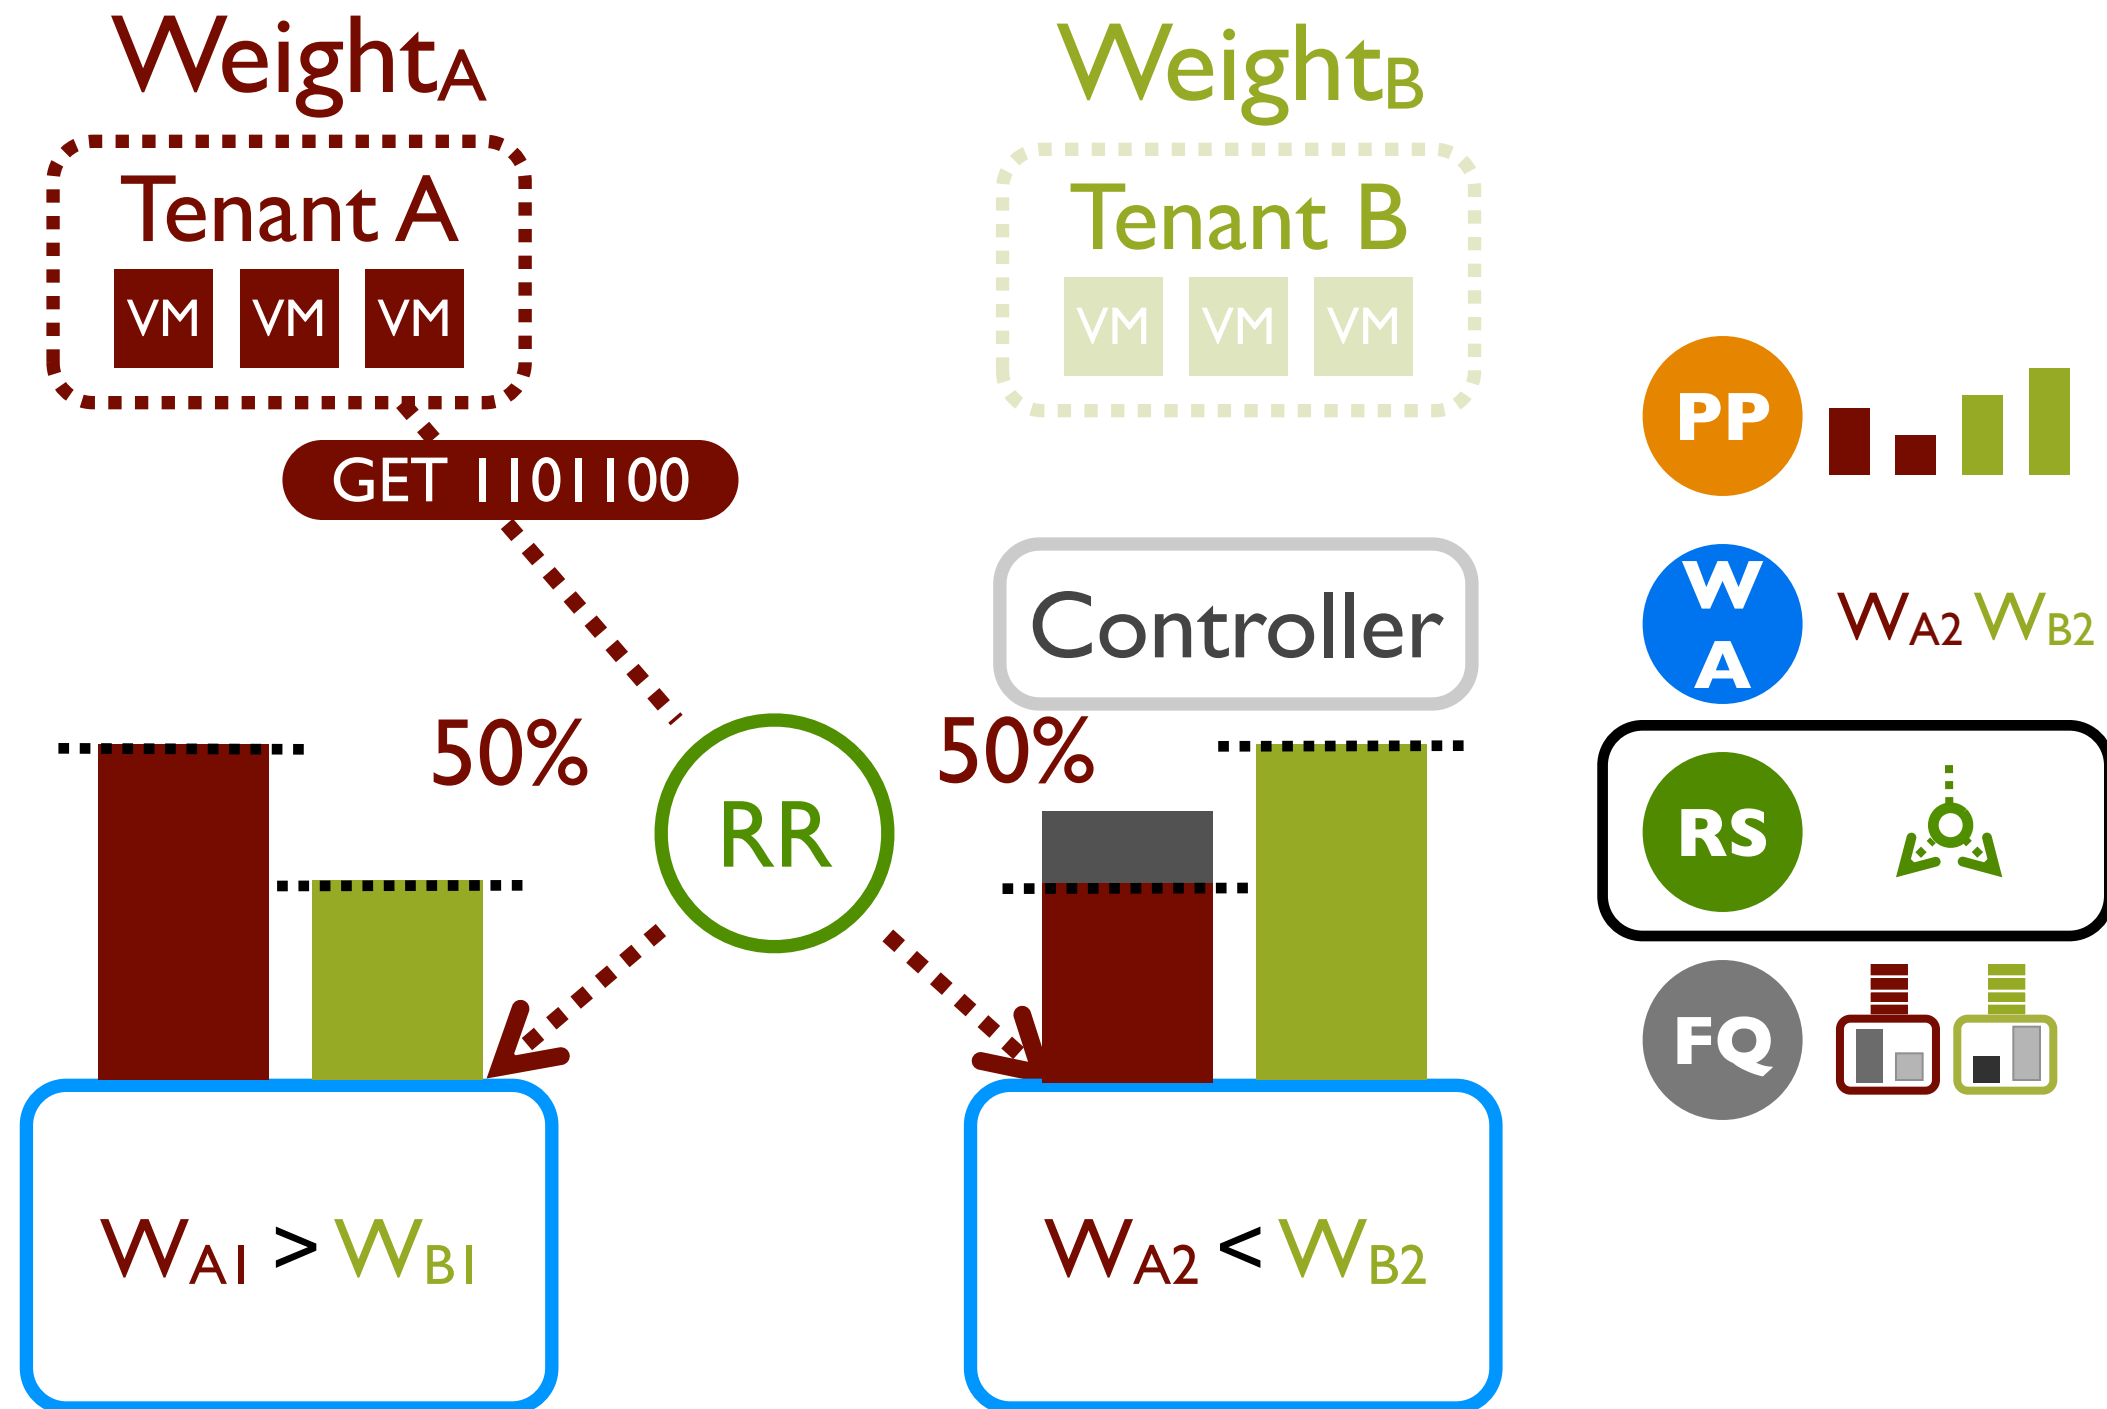
Strawman: Allocate Local Weights Evenly



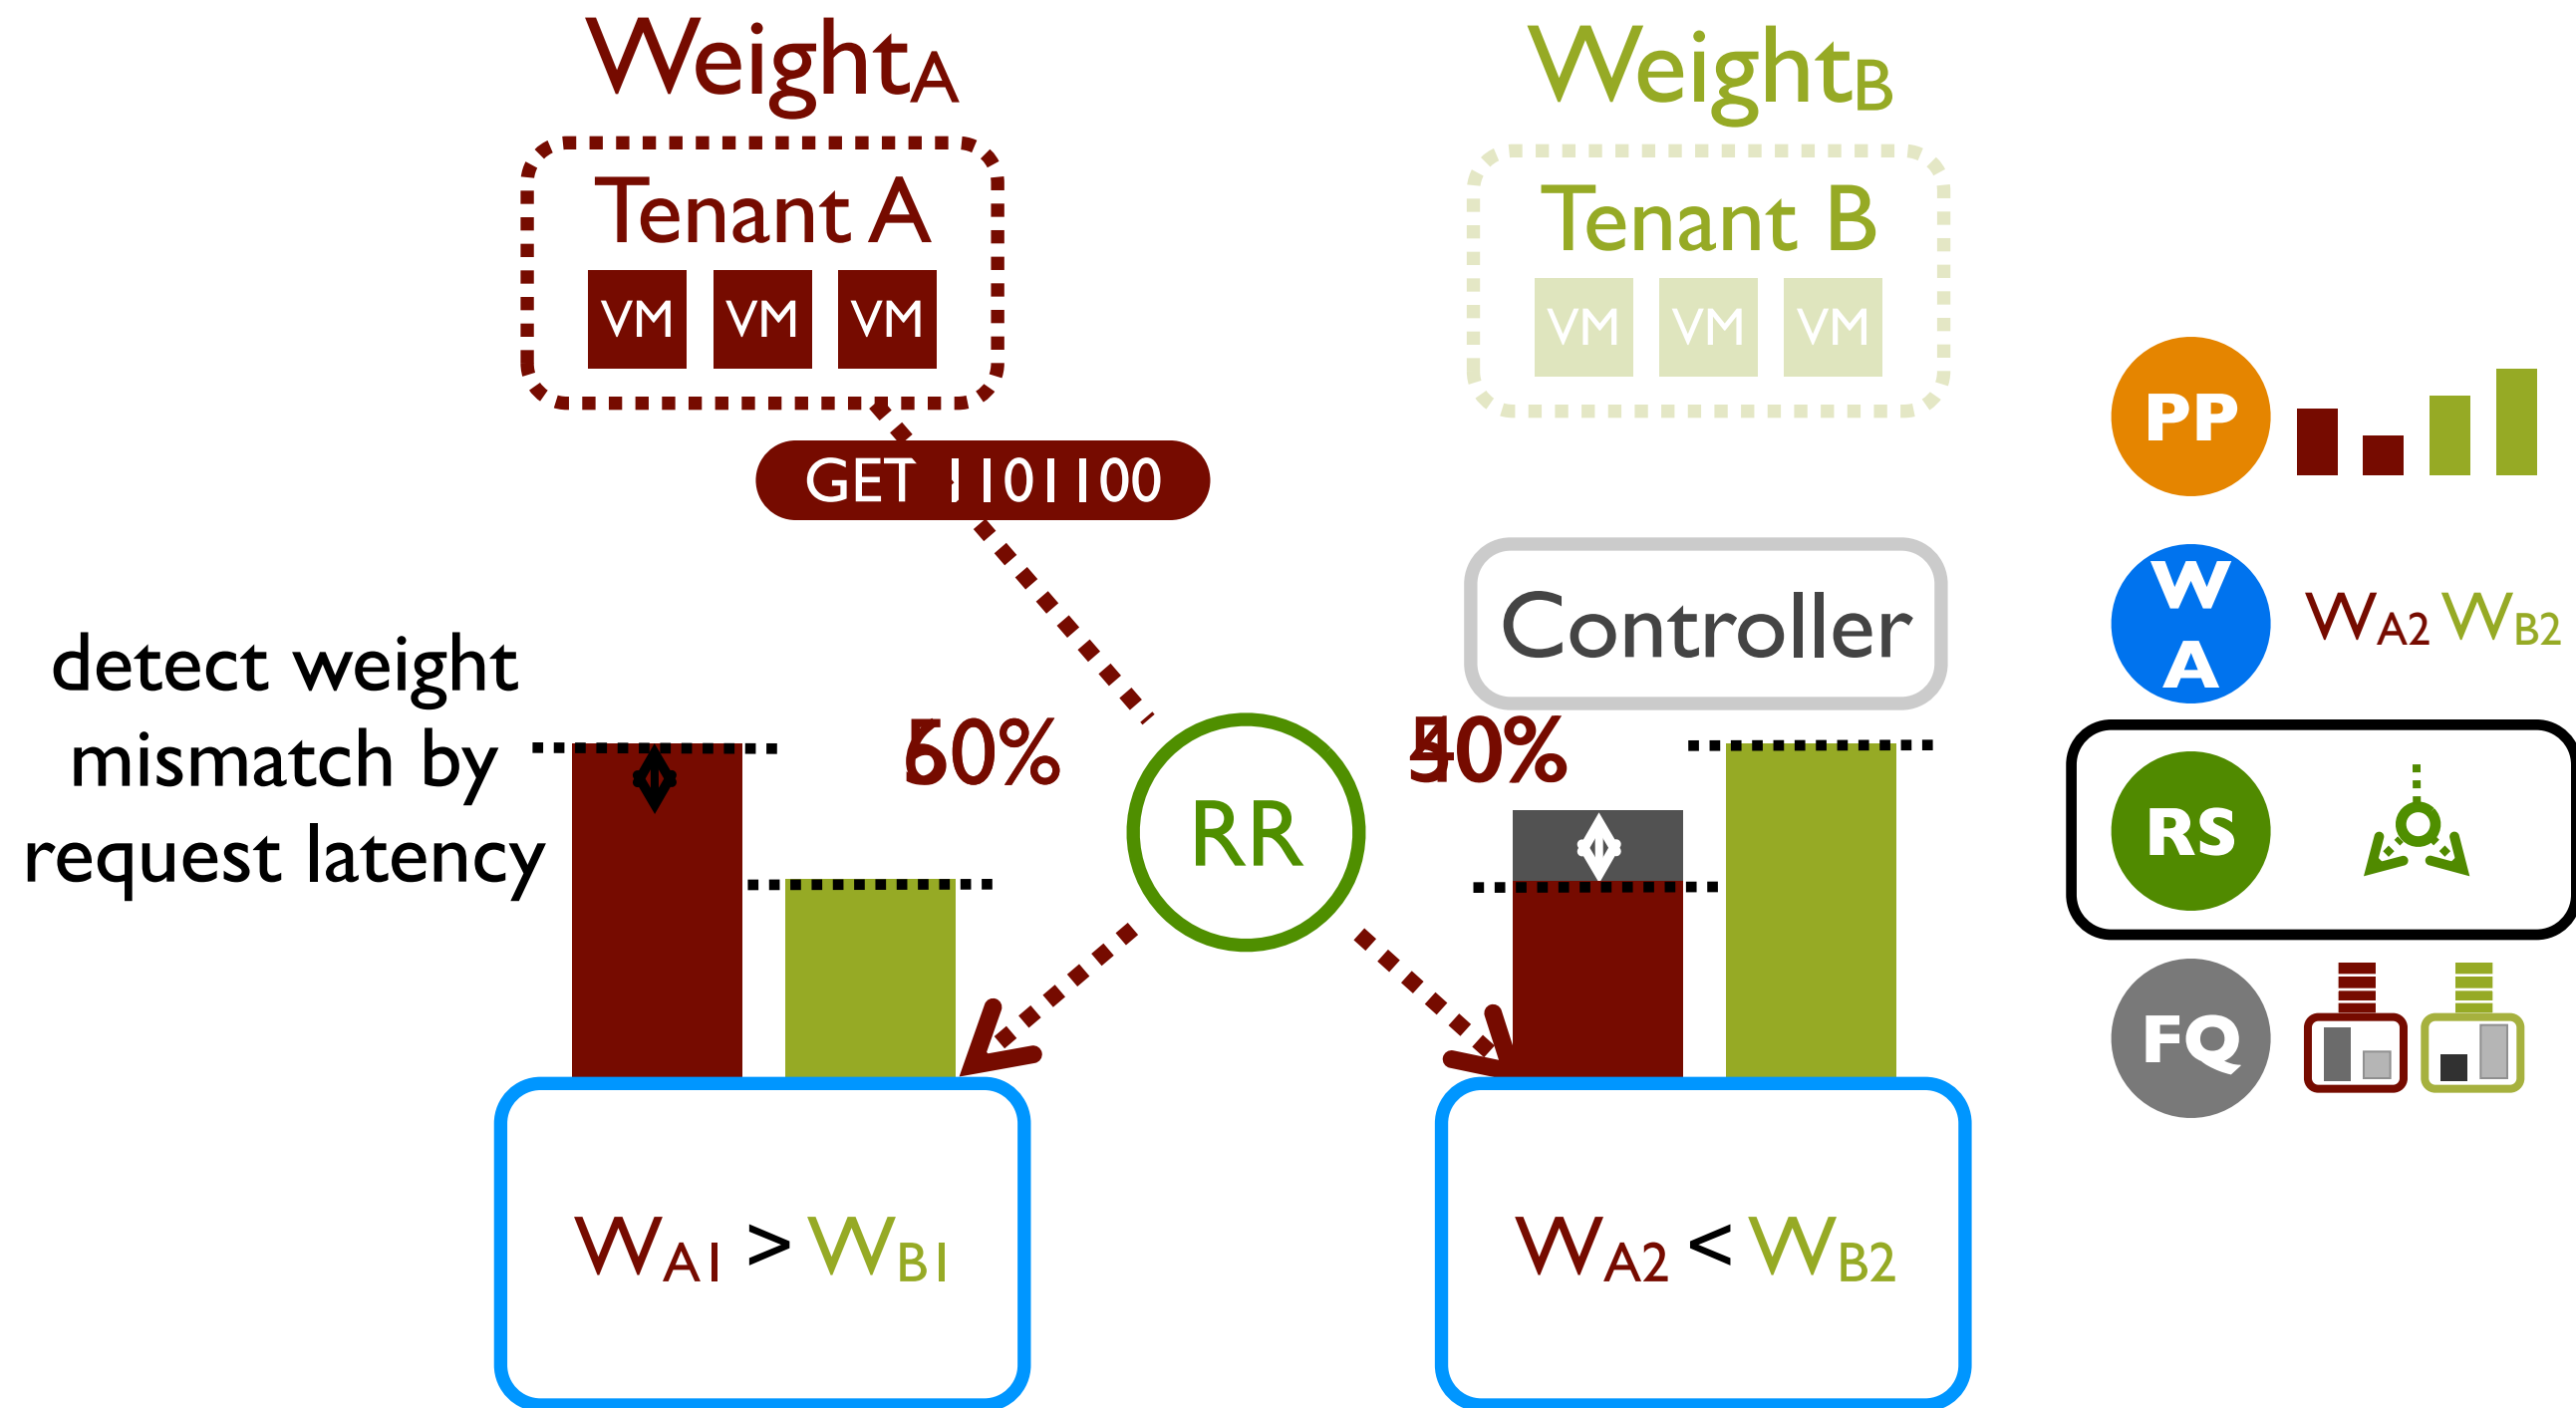
Pisces: Allocate Local Weights By Tenant Demand



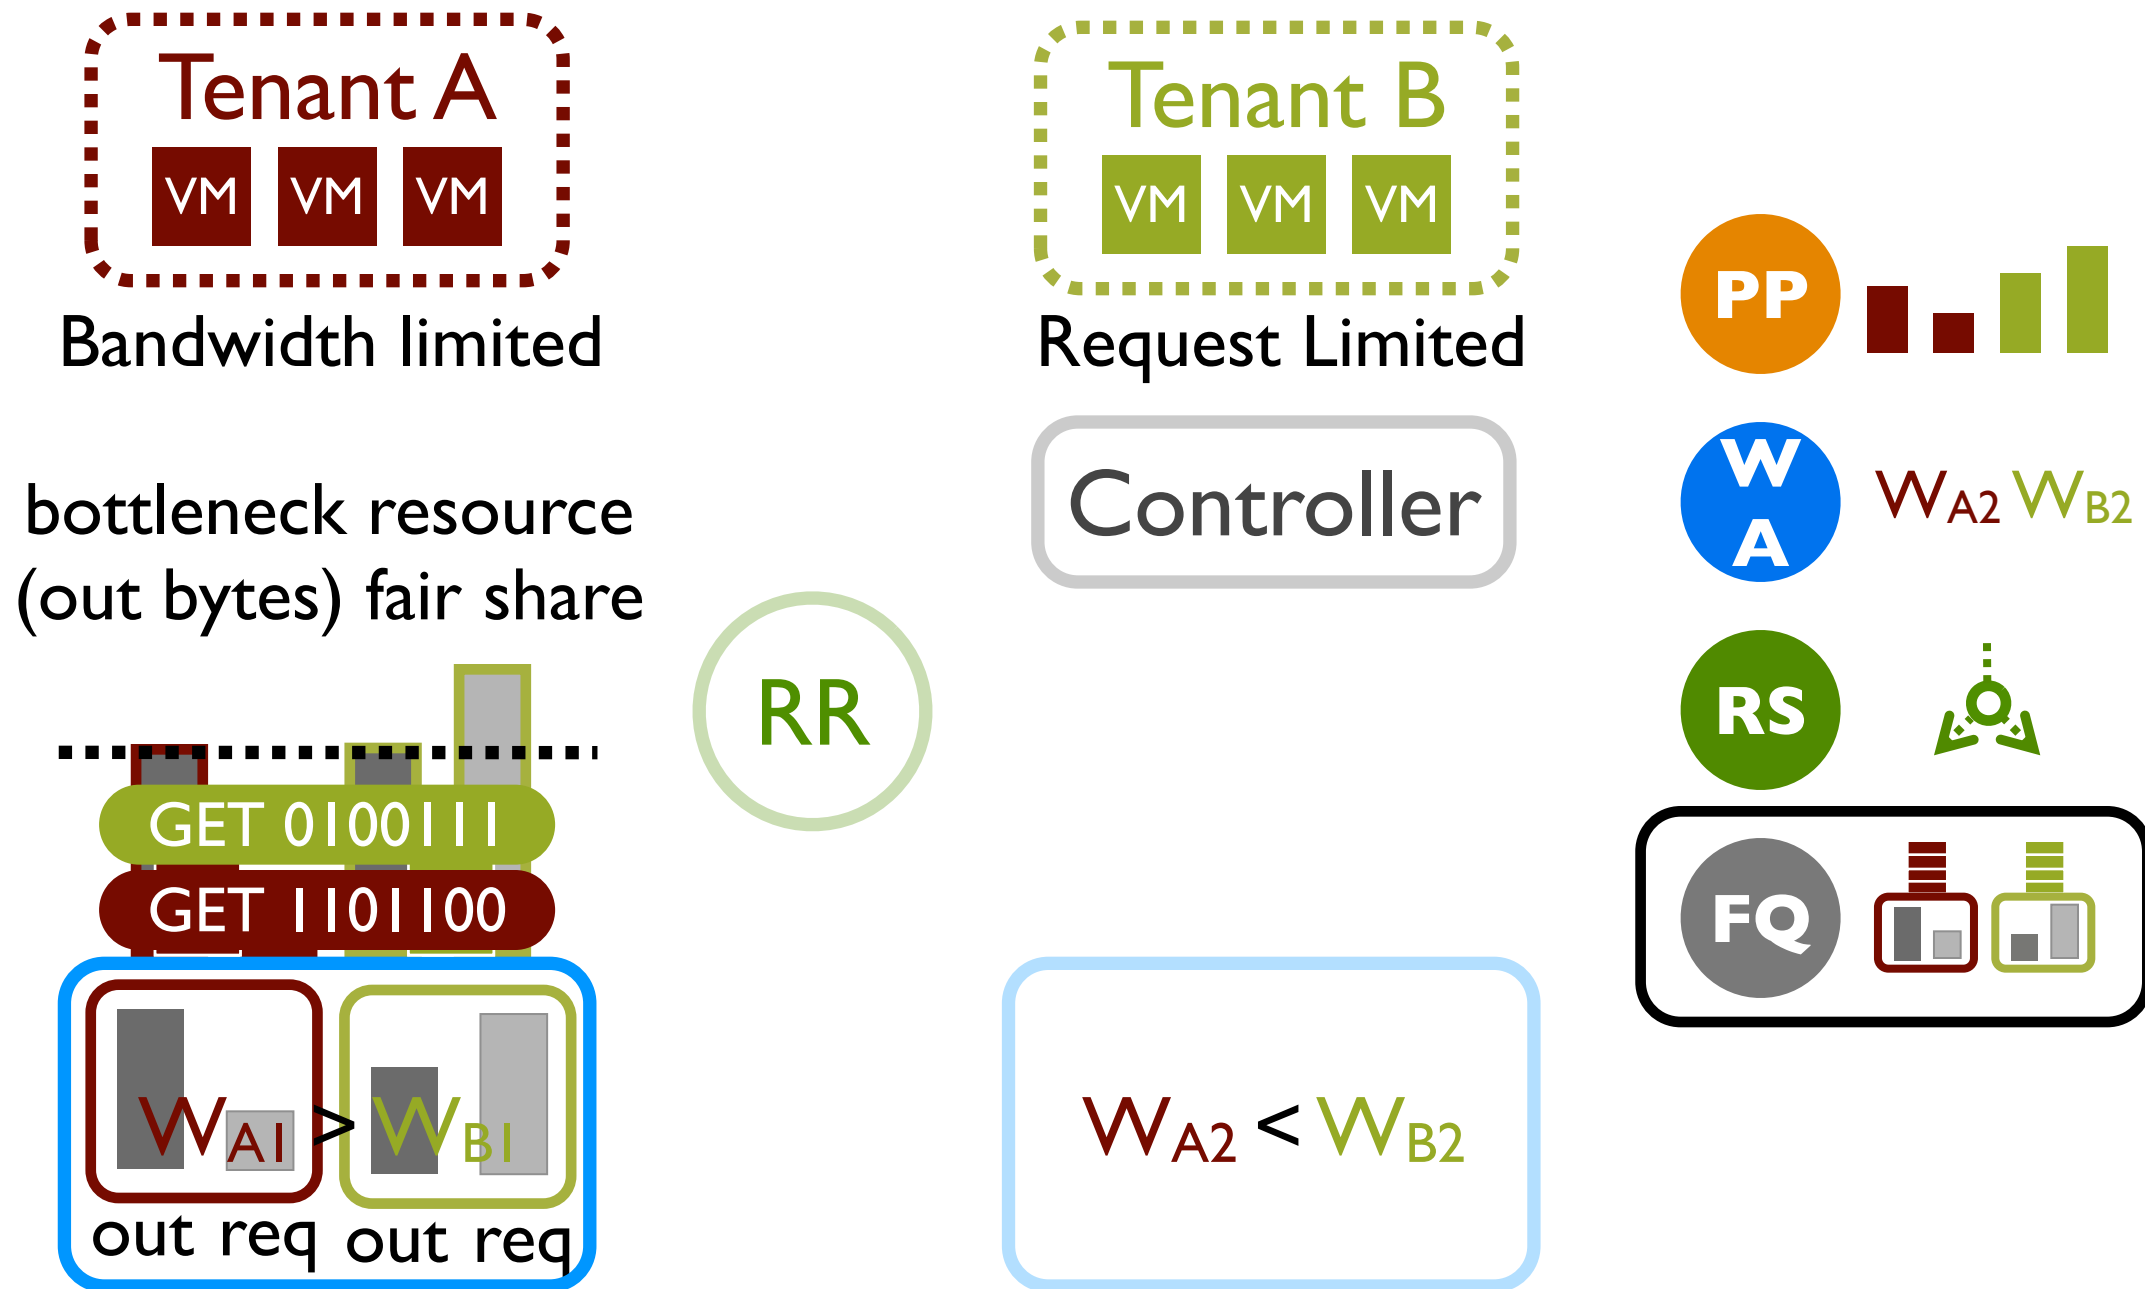
Strawman: Select Replicas Evenly



Pisces: Select Replicas By Local Weight



Strawman: Queue Tenants By Single Resource



Pisces: Queue Tenants By Dominant Resource




Bandwidth limited



Request Limited


dominant resource
fair share



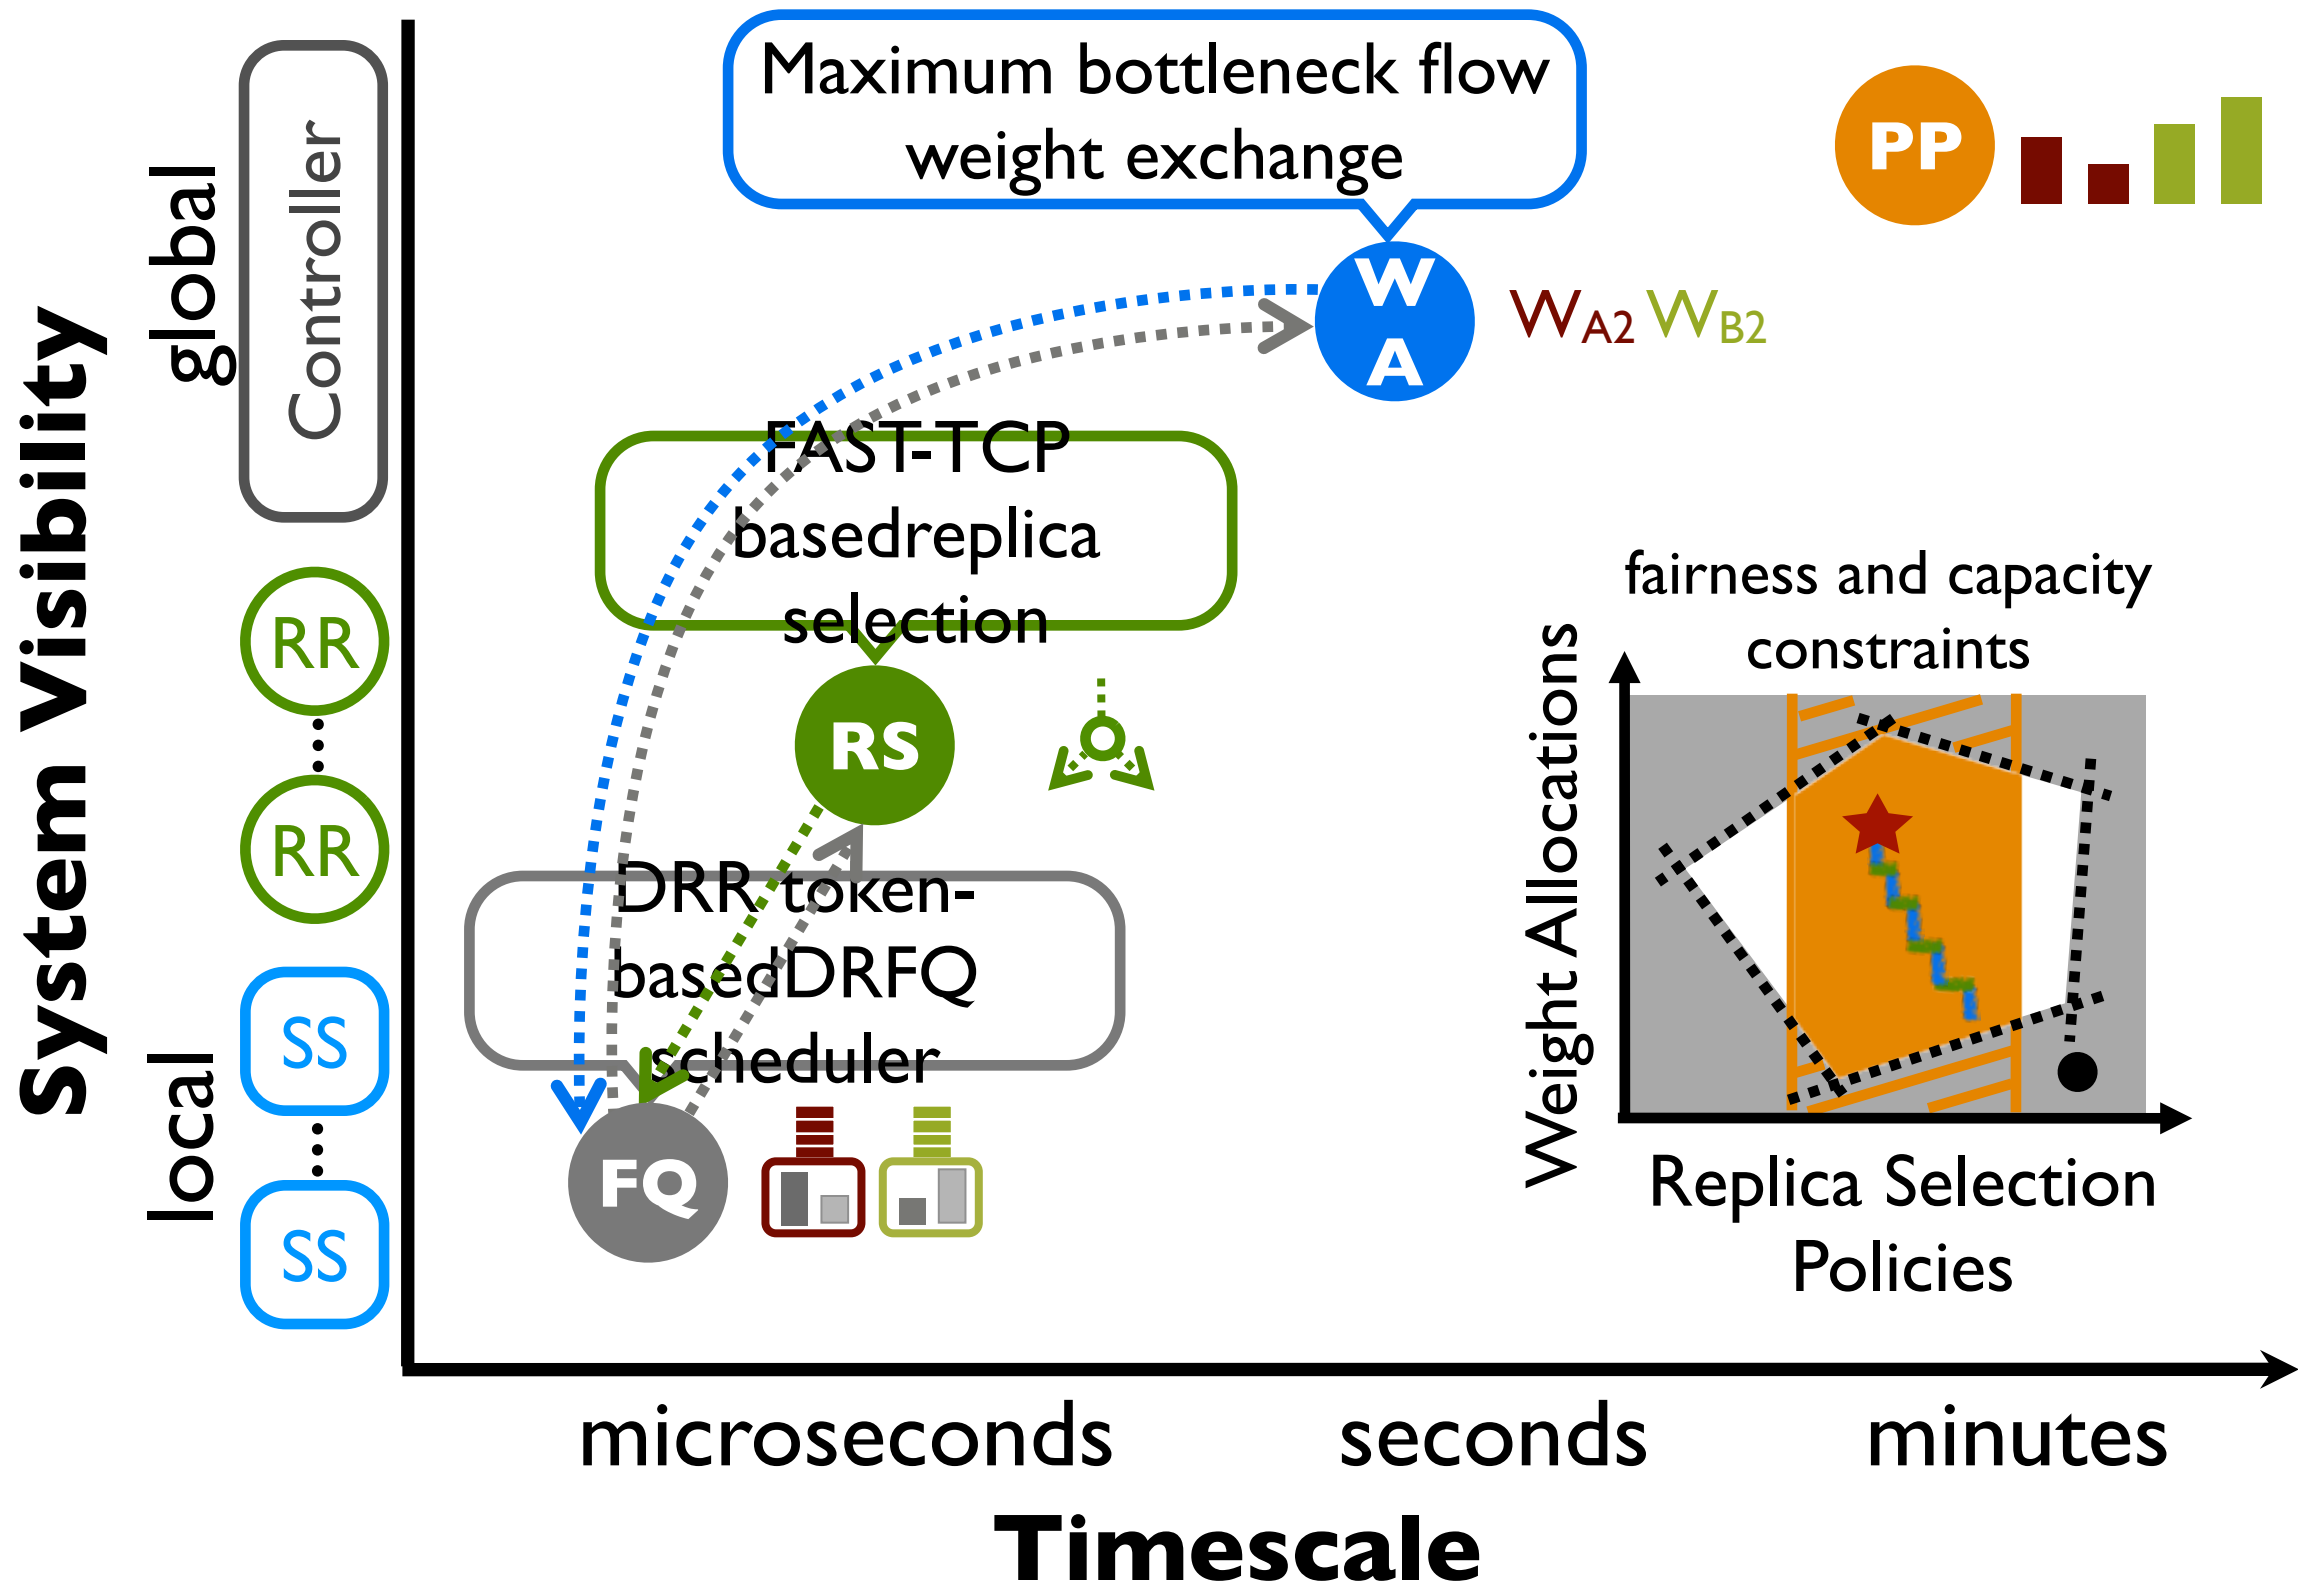
Track per-tenant
resource vector



$$W_{A2} < W_{B2}$$



Pisces Mechanisms Solve For Global Fairness



Evaluation

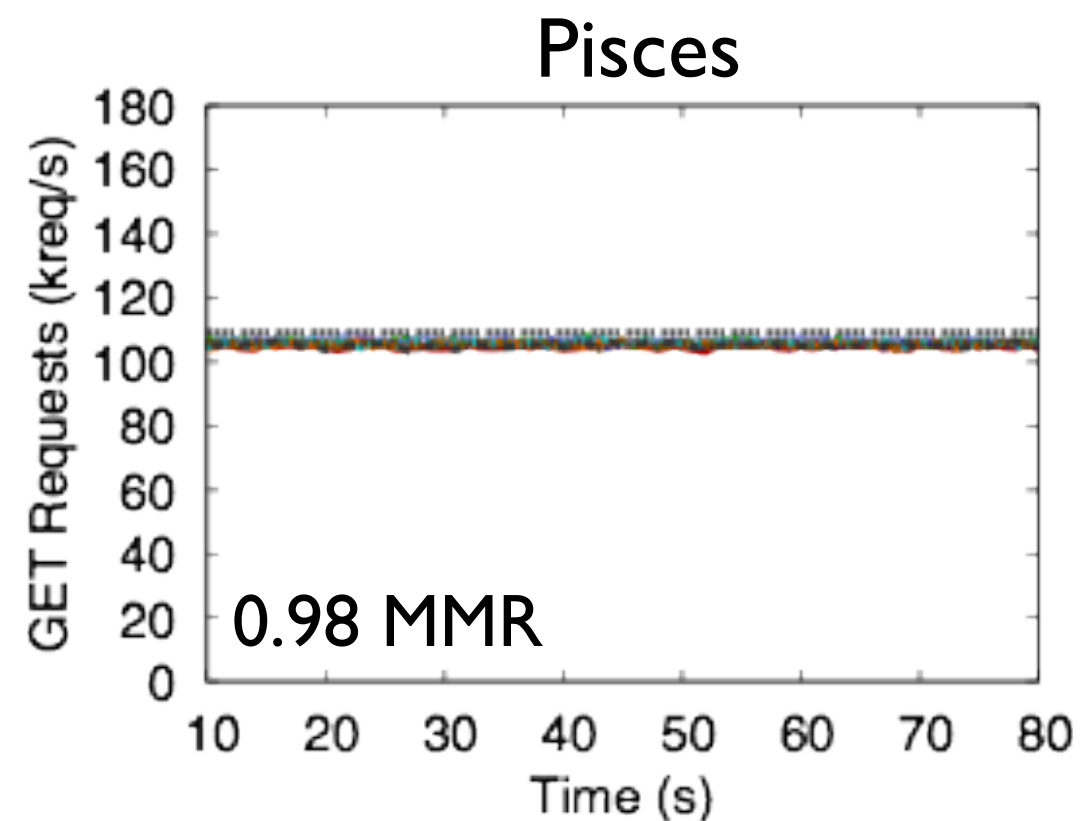
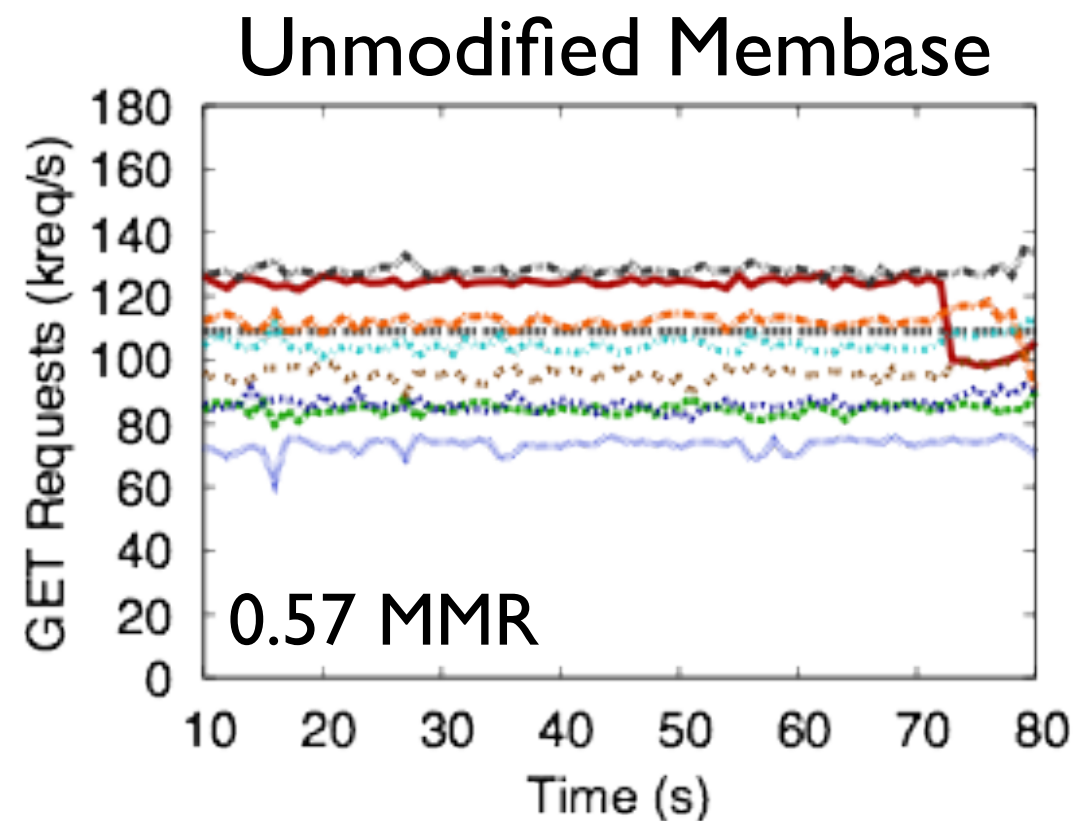
- Does Pisces achieve (even) system-wide fairness?
 - Is each Pisces mechanism necessary for fairness?
 - What is the overhead of using Pisces?
- Does Pisces handle mixed workloads?
- Does Pisces provide weighted system-wide fairness?
- Does Pisces provide local dominant resource fairness?
- Does Pisces handle dynamic demand?
- Does Pisces adapt to changes in object popularity?

Evaluation

- Does Pisces achieve (even) system-wide fairness?
 - Is each Pisces mechanism necessary for fairness?
 - What is the overhead of using Pisces?
- Does Pisces handle mixed workloads?
- Does Pisces provide weighted system-wide fairness?
- Does Pisces provide local dominant resource fairness?
- Does Pisces handle dynamic demand?
- Does Pisces adapt to changes in object popularity?

Pisces Achieves System-wide Per-tenant Fairness

Ideal fair share: 110 kreq/s (1kB requests)



8 Tenants - 8 Client - 8 Storage Nodes

Zipfian object popularity distribution

Min-Max Ratio: min rate/max rate (0,1]

Each Pisces Mechanism Contributes to System-wide Fairness and Isolation

Unmodified Membase

FQ

FQ

PP


W
A

FQ


PP

RS


W
A



2x vs 1x demand




Pisces Imposes Low-overhead



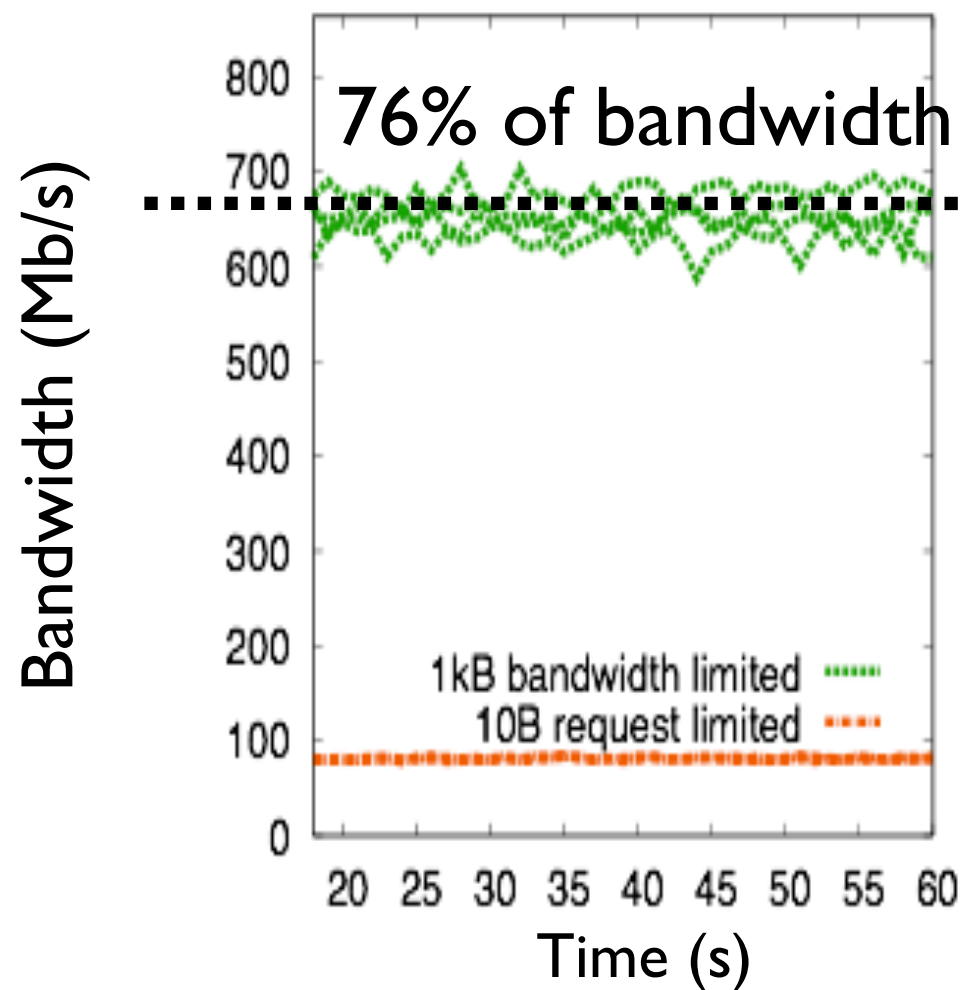
Pisces Achieves System-wide Weighted Fairness

0.98 MMR 0.89 MMR 0.91 MMR
4 heavy hitters 20 moderate demand 40 low demand




Pisces Achieves Dominant Resource Fairness



1kB workload
bandwidth limited



10B workload
request limited



Pisces Adapts to Dynamic Demand



Conclusion

● Pisces Contributions

- Per-tenant weighted max-min fair shares of system-wide resources w/ high utilization
- Arbitrary object distributions
- Different resource bottlenecks
- Novel decomposition into 4 complementary mechanisms



Partition
Placement



Weight
Allocation



Replica
Selection



Fair
Queuing

Thank you

Performance Isolation and Fairness for Multi-Tenant Cloud Storage



David Shue*, Michael Freedman*, and Anees Shaikh♦