

Profit-Maximizing Virtual Machine Trading in a Federation of Selfish Clouds

Presenter: Shengkai Shi

Jan. 10 , 2013

Problem background

- Cloud federation
- When and how should a selfish cloud trade VMs with others
- Correlated sub-problems
 - VM pricing
 - Job scheduling
 - Server provisioning

Algorithm

- Target:
 - Maximize the individual profit
 - Achieve a satisfactory social welfare
- Feature:
 - Lyapunov optimization
 - Truthful double auction

System model

Service demand

- F individual clouds
- M distinct VM types
- S distinct job types
 - m_s : the type of VM required by type- s job
 - g_s : the number of VM required by type- s job
 - d_s : the SLA of type- s job

System model (cont.)

Service demand

- N_i^m homogenous servers to provision type- m VM in cloud i
 - C_i^m : the maximum number of type- m VM in each server
 - $\sum_{m=1}^M N_i^m$: the number of servers in cloud i
- $r_i^s(t)$: the amount of type- s job arriving at cloud i in t
- $p_i^s(t)$: the given price of type- s job in t

System model

Job scheduling

- $Q_s^i(t)$: the length of the queue for type- s job
- $D_s^i(t)$: the number of type- s jobs dropped
- ξ_s^i : the penalty to drop one type- s job
- $\mu_{ij}^s(t)$: the number of type- s jobs scheduled to cloud j by cloud i
- Queue update:

$$Q_i^s(t+1) = \max\{Q_i^s(t) - \sum_{j=1}^F \mu_{ij}^s(t) - D_i^s(t), 0\} + r_i^s(t), \forall s \in [1, S], \forall i \in [1, F]$$

System model (cont.)

Job scheduling

- Z_i^s : a virtual queue reflecting the cumulated response delay of jobs from the respective job queue

$$Z_i^s(t+1) = \max\{Z_i^s(t) + 1_{\{Q_i^s(t) > 0\}} \cdot [\varepsilon_s - \sum_{j=1}^F \mu_{ij}^s(t)] - D_i^s(t) - 1_{\{Q_i^s(t) = 0\}} \cdot \sum_{j=1}^F \frac{C_j^{m_s} \cdot N_j^{m_s}}{g_s}, 0\}$$

$$\forall s \in [1, S], \forall i \in [1, F]$$

System model

Server provisioning

- Electricity cost
 - main component of operation cost
- $\beta_i(t)$: the operation cost of cloud i in t
- $n_i^m(t)$: the number of active servers provisioning type- m VMs at cloud i in t

- Constrain:

$$\sum_{j \in [1, F]} \sum_{s: m_{s=m}, s \in [1, S]} g_s \mu_{ij}^s(t) \leq C_i^m \cdot n_i^m(t)$$

$$n_i^m(t) \leq N_i^m, \forall m \in [1, M], \forall s \in [1, F]$$

Double auction

- Buyer and seller
 - $\langle b_i^m(t), \gamma_i^m(t) \rangle$: the unit price and maximum quantity at which cloud i is willing to buy type- m VMs
 - $\langle s_i^m(t), \eta_i^m(t) \rangle$: the unit price and maximum quantity at which cloud i is willing to sell type- m VMs

Double auction (cont.)

- Auctioneer
 - a broker in the cloud federation
 - $\bar{b}_i^m(t)$ and $\bar{\gamma}_i^m(t)$: the actual unit charge price and actual quantity for cloud i to buy type- m VMs respectively
 - $\bar{s}_i^m(t)$ and $\bar{\eta}_i^m(t)$: the actual unit income and actual quantity for cloud i to sell type- m VMs respectively

Double auction (cont.)

$$\hat{\gamma}_i^m(t) = \sum_{j \in [1, F], j \neq i} \alpha_{ij}^m(t), \forall m \in [1, M], i \in [1, F]$$

$$\hat{\eta}_i^m(t) = \sum_{j \in [1, F], j \neq i} \alpha_{ji}^m(t), \forall m \in [1, M], \forall i \in [1, F]$$

- Here,
 - $\alpha_{ij}^m(t)$: the number of type- m VMs purchased by cloud i by cloud j from cloud

Double auction (cont.)

- Economic properties for double auction
 - truthfulness
 - individual rationality
 - ex-post budget balance

Truthfulness

- Definition:

A truthful double auction is employed at the auctioneer, where sellers and buyers bid their true values of the prices and quantities, in order to maximize their individual utilities

- $\tilde{b}_i^m(t)$ and $\tilde{s}_i^m(t)$: the true value of buying and selling a type- m VM respectively
- $\tilde{\gamma}_i^m(t)$ and $\tilde{\eta}_i^m(t)$: the true value of the quantity to buy and sell type- m VMs respectively

Individual selfishness

- Revenue

- job service charge paid by customers

$$\Phi_1^i = \lim_{T \rightarrow \infty} \frac{1}{T} \sum_{t=0}^{T-1} \sum_{s \in [1, S]} E\{p_i^s(t) \cdot r_i^s(t)\}, \forall i \in [1, F]$$

- income from selling VMs

$$\Phi_2^i = \lim_{T \rightarrow \infty} \frac{1}{T} \sum_{t=0}^{T-1} \sum_{m \in [1, M]} E\{\hat{s}_i^m(t) \cdot \hat{\eta}_i^m(t)\}, \forall i \in [1, F]$$

Individual selfishness (cont.)

- Cost

- cost for running active servers

$$\Psi_1^i = \lim_{T \rightarrow \infty} \frac{1}{T} \sum_{t=0}^{T-1} \sum_{s \in [1, S]} E\{\beta_i(t) \cdot \sum_{m=1}^M n_i^m(t)\}, \forall i \in [1, F]$$

- penalties for dropped jobs

$$\Psi_2^i = \lim_{T \rightarrow \infty} \frac{1}{T} \sum_{t=0}^{T-1} \sum_{s \in [1, S]} E\{\xi_i^s \cdot D_i^s(t)\}, \forall i \in [1, F]$$

- cost for buying VMs from other clouds

$$\Psi_3^i = \lim_{T \rightarrow \infty} \frac{1}{T} \sum_{t=0}^{T-1} E\left\{ \sum_{m \in [1, M]} \hat{b}_i^m(t) \cdot \hat{\gamma}_i^m(t) \right\}, \forall i \in [1, F]$$

Individual selfishness (cont.)

- Profit maximization

$$\max \Phi_1^i + \Phi_2^i - \Psi_1^i - \Psi_2^i - \Psi_3^i$$

- Social welfare
 - overall profit of the cloud federation

$$\sum_{i \in [1, F]} \Phi_1^i + \Phi_2^i - \Psi_1^i - \Psi_2^i - \Psi_3^i$$

- maximization

$$\sum_{i \in [1, F]} \Phi_1^i - \Psi_1^i - \Psi_2^i$$

Dynamic algorithms

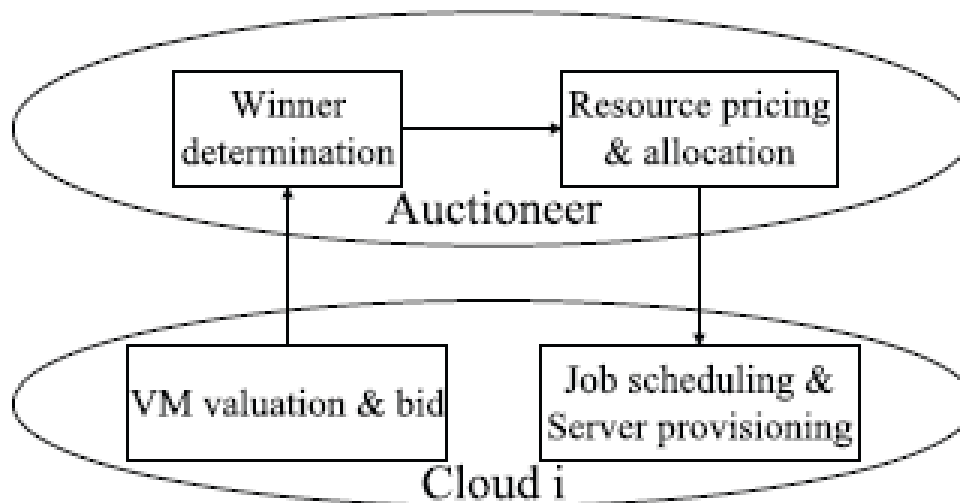


Figure: Key algorithm modules

Individual profit maximization

- Drift-plus-penalty framework in Lyapunov optimization theory
- One-shot optimization problem:

$$\max \varphi_1^i(t) + \varphi_2^i(t) + \varphi_3^i(t)$$

where

$$\varphi_1^i = V \sum_{m \in [1, M]} [\hat{s}_i^m(t) \hat{\eta}_i^m(t) - \hat{b}_i^m(t) \hat{\gamma}_i^m(t) - \beta_i(t) n_i^m(t)]$$

$$\varphi_2^i = \sum_{s \in [1, S]} \sum_{j \in [1, F]} \mu_{ij}^s(t) [Q_i^s(t) + Z_i^s(t)]$$

$$\varphi_3^i = \sum_{s \in [1, S]} D_i^s(t) [Q_i^s(t) + Z_i^s(t) - V \cdot \xi_i^s]$$

Individual profit maximization (cont.)

- Dynamic profit maximization algorithm

Algorithm 1 Dynamic Profit Maximization Algorithm at cloud i in Time Slot t

Input: $r_i^s(t)$, $Q_i^s(t)$, $Z_i^s(t)$, g_s , m_s , ξ_i^s , C_i^m , N_i^m and $\beta_i(t)$, $\forall s \in [1, S]$.

Output: $b_i^m(t)$, $s_i^m(t)$, $\gamma_i^m(t)$, $\eta_i^m(t)$, $D_i^s(t)$, $\mu_{ij}^s(t)$ and $n_i^m(t)$, $\forall m \in [1, M]$, $s \in [1, S]$, $j \in [1, F]$.

- 1: **VM valuation and bid:** Decide $b_i^m(t)$, $s_i^m(t)$, $\gamma_i^m(t)$ and $\eta_i^m(t)$ with Eqn. (18)-(23);
- 2: **Server provisioning, job scheduling and dropping:** Decide $\mu_{ij}^s(t)$, $D_i^s(t)$ and $n_i^m(t)$ with Eqn. (24), (26) and (27);
- 3: Update $Q_i^s(t)$ and $Z_i^s(t)$ with Eqn. (2) and (4).

VM valuation

- True value of the price to buy (sell)
 - equal to this value, profit remains the same
 - lower than this value, a profit gain (loss) occurs
 - higher than this value, a profit loss (gain) occurs

VM valuation (cont.)

- True values of the buy and sell prices

$$\tilde{b}_i^m(t) = \frac{Q_i^{s_m^*}(t) + Z_i^{s_m^*}(t)}{V \cdot g_{s_m^*}}$$

$$\tilde{s}_i^m(t) = \begin{cases} \frac{Q_i^{s_m^*}(t) + Z_i^{s_m^*}(t)}{V \cdot g_{s_m^*}} & \text{if } \frac{Q_i^{s_m^*}(t) + Z_i^{s_m^*}(t)}{V \cdot g_{s_m^*}} > \beta_i(t) / c_i^m \\ \beta_i(t) / c_i^m & \text{Otherwise} \end{cases}$$

where

$$s_m^* = \arg \max_{s' \in [1, S], m_{s'} = m} \{W_i^{s'}(t)\}$$

$$W_i^{s'}(t) = \frac{Q_i^{s'}(t) + Z_i^{s'}(t)}{g_{s'}}$$

VM valuation (cont.)

- True volumes of type- m VMs to buy and sell

$$\tilde{\gamma}_i^m(t) = \sum_{i \in [1, F]} C_i^m N_i^m$$

$$\tilde{\eta}_i^m(t) = C_i^m N_i^m$$

- $\tilde{\gamma}_i^m(t)$: the number of all potential type- m VMs
- Bidding price

$$b_i^m(t) = \tilde{b}_i^m(t), s_i^m(t) = \tilde{s}_i^m(t)$$

$$\gamma_i^m(t) = \tilde{\gamma}_i^m(t), \eta_i^m(t) = \tilde{\eta}_i^m(t)$$

Double auction

- Winner determination
 - sort bids in non-decreasing (for sellers) and non-increasing (for buyers) orders:

$$\theta_1^m(t) \geq \theta_2^m(t) \geq \dots \geq \theta_N^m(t)$$

$$v_1^m(t) \leq v_2^m(t) \leq \dots \leq v_N^m(t)$$

- find $j' = \arg \max v_j^m(t) \leq \theta_j^m(t)$, the largest index in the sorted sequence of sell-bids such that,

$$v_{j'}^m(t) \leq \theta_2^m(t), v_{j'+1}^m(t) > \theta_2^m(t)$$

Price and allocation

- Uniform clearing price
 - the actual charge price to each buyer cloud

$$\hat{b}_i^m(t) = \begin{cases} \theta_2^m(t) & \text{if } b_i^m(t) \text{ wins} \\ 0 & \text{Otherwise} \end{cases}$$

- the actual price paid to each seller cloud

$$\hat{s}_i^m(t) = \begin{cases} v_{j'}^m(t) & \text{if } s_i^m(t) \text{ wins} \\ 0 & \text{Otherwise} \end{cases}$$

Price and allocation (cont.)

- Uniform clearing price
 - the number of type- m VMs bought by cloud

$$\hat{\gamma}_i^m(t) = \begin{cases} \sum_{j=1}^{j'-1} L_j^m(t) & \text{if } b_i^m(t) \text{ wins} \\ 0 & \text{Otherwise} \end{cases}$$

- the number of type- m VMs sold by cloud

$$\hat{\eta}_i^m(t) = \begin{cases} \eta_i^m(t) & \text{if } s_i^m(t) \text{ wins} \\ 0 & \text{Otherwise} \end{cases}$$

Price and allocation (cont.)

- Uniform clearing price
 - the number of type- m VMs sold from cloud j to cloud i

$$\alpha_{ij}^m(t) = \begin{cases} \eta_j^m(t) & \text{if } s_i^m(t) \text{ and } b_i^m(t) \text{ win} \\ 0 & \text{Otherwise} \end{cases}$$

Example

	Cloud 1	Cloud 2	Cloud 3	Cloud 4
Buy-bid	\$10	\$20	\$15	\$8
Sell-bid	\$13	\$22	\$16	\$9

- Buy-bidder winner: Cloud 2
- Sell-bidder winner: Cloud 4
- Actual clearing buy price: \$15
- Actual clearing sell price: \$13

Job scheduling

- The number of type- m jobs scheduled to local servers

$$\mu_{ii}^s(t) = \begin{cases} \frac{C_i^{m_s} \cdot N_i^{m_s} - \sum_{j \neq i} \alpha_{ij}^{m_s}(t)}{g_s} & \text{if } \frac{Q_i^s(t) + Z_i^s(t)}{V \cdot g_s} > \beta_i(t) / C_i^{m_s} \text{ and } s = s_{m_s}^* \\ 0 & \text{Otherwise} \end{cases}$$

- The number of type- m jobs scheduled to other cloud ($j \neq i$)

$$\mu_{ij}^s(t) = \begin{cases} \frac{\alpha_{ij}^{m_s}(t)}{g_s} & \text{if } s = s_{m_s}^* \\ 0 & \text{Otherwise} \end{cases}$$

Job dropping

- The number of type- s jobs dropped by cloud i

$$D_i^s(t) = \begin{cases} D_s^{(\max)}(t) & \text{if } Q_i^s(t) + Z_i^s(t) > V \cdot \xi_i^s \\ 0 & \text{Otherwise} \end{cases}$$

Here,

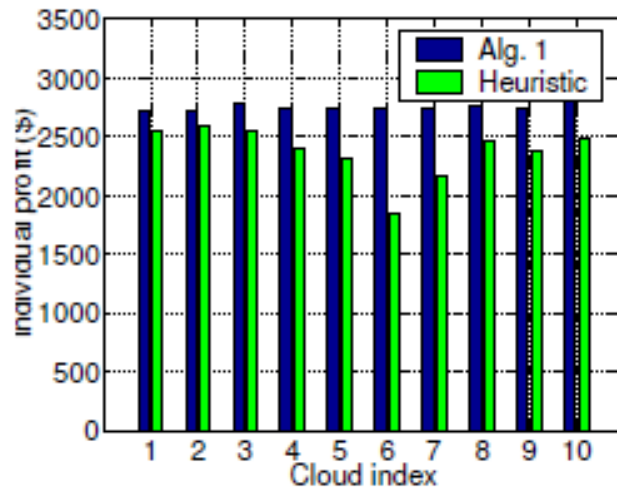
- $D_s^{(\max)}$: the maximum number of type- s jobs dropped by cloud i

Server provisioning

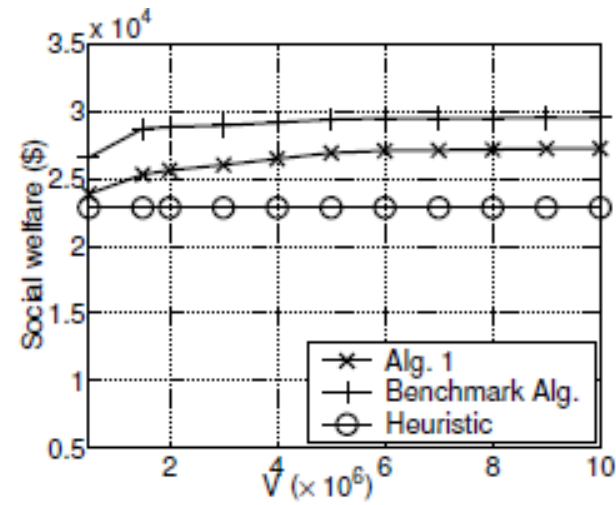
- The number of activated servers at cloud i to provision type- m VM

$$n_i^m(t) = \frac{\sum_{s \in [1, S], m_s = m} \mu_{ii}^s(t) \cdot g_s + \sum_{j \neq i} \alpha_{ji}^m(t)}{C_i^m}$$

Simulation results



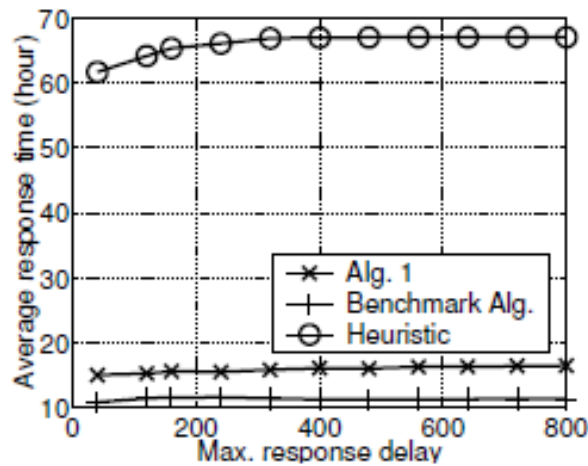
(a) Individual Profit with $V = 10^7$



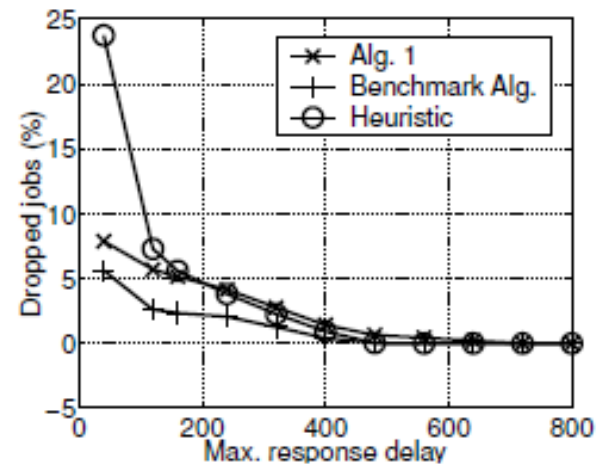
(b) Social Welfare

Figure: Comparisons of individual profit and social welfare

Simulation results (cont.)



(a) Average response delay



(b) Averaged job-drop percentage

Figure: Comparisons of job scheduling delay and drop percentage

Conclusion

- Selfishness of individual cloud is addressed
- A truthful, individual-rational, ex-post budget-balanced double auction
- Online VM pricing, job scheduling, and server provisioning
- Theoretical analysis and simulation study

Thank you!