

mOS

NSDI2017

Simple packet processing

```
example.cc
1
2 void simple_packet_processing_program(){
3     packet* pkts[32];
4
5     // an infinite loop.
6     for(;;){
7         int count = fetch_packets(pkts, 32);
8
9         // the processing loop.
10        for(int i=0; i<count; i++){
11            process_packet(pkts[i]);
12        }
13
14        // the releasing loop
15        for(int i=0; i<count; i++){
16            release_packet(pkts[i]);
17        }
18    }
19 }
20
```

Complicated stream processing

```
1 static bool
2 IsFakeRexmit(mctx_t mctx, int sock, int side, event_t event,
3 struct filter_arg *arg) {
4     struct pkt_info pi;
5     char buf[MSS];
6     struct tcp_ring_fragment frags[MAX_FRAG_NUM];
7     int nfrags = MAX_FRAG_NUM;
8     int i, size, boff, poff;
9
10    // retrieve the current packet information
11    mtcp_getlastpkt(mctx, sock, side, &pi);
12
13    // for full retransmission, compare the entire payload
14    if (mtcp_peek(mctx, sock, side, buf, pi.payloadlen, pi.offset) ==
15        pi.payloadlen)
16        return memcmp(buf, pi.payload, pi.payloadlen);
17
18    // for partial retransmission, compare the overlapping region
19    // retrieve the data fragments and traverse them
20    mtcp_getsockopt(mctx, sock, SOL_MONSOCKET, (side == MOS_SIDE_CLI) ?
21        MOS_FRAGINFO_CLIBUF : MOS_FRAGINFO_SVRBUF, frags, &nfrags);
22
23    for (i = 0; i < nfrags; i++) {
24        if ((size = CalculateOverlapLen(&pi, &(frags[i]), &boff, &poff)))
25            if (memcmp(buf + boff, pi.payload + poff, size))
26                return true; // payload mismatch detected
27    }
28    return false;
29 }
```

An overview of mOS

- Do not directly access and modify packets.
- Hide packet processing logic from the programmer.
- Expose events (**primarily related to TCP**) to the programmer.
- Program middlebox by writing callback functions for events.

The architecture graph

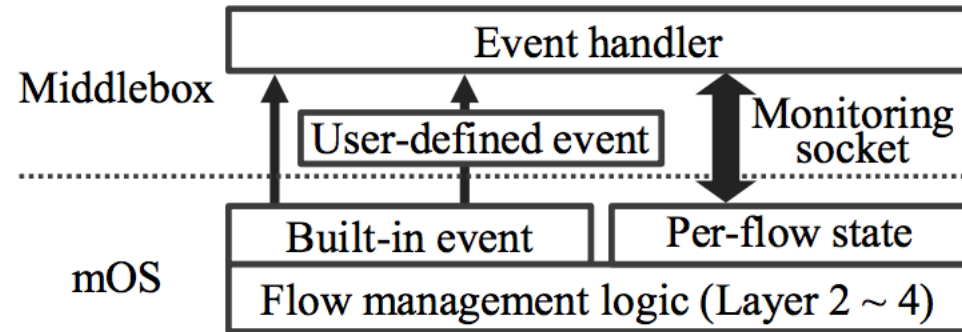
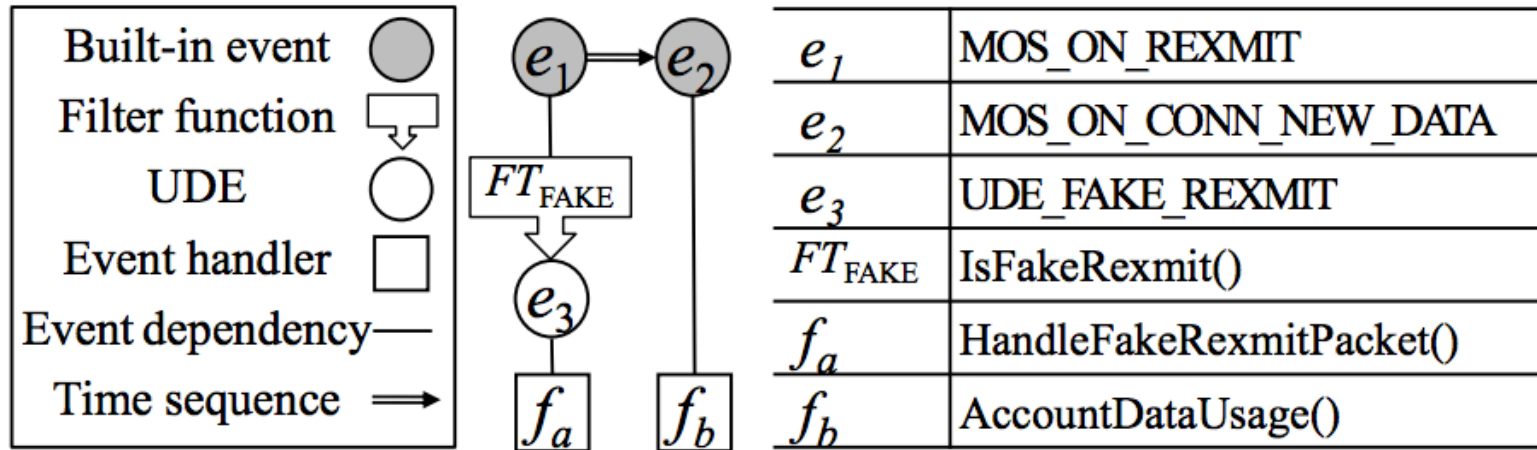


Figure 1: Interaction between mOS and its application

Event action diagram



Eight built-in event

| Event | Description |
|-------------------------|--|
| MOS_ON_PKT_IN | In-flow TCP packet arrival |
| MOS_ON_CONN_START | Connection initiation (the first SYN packet) |
| MOS_ON_REXMIT | TCP packet retransmission |
| MOS_ON_TCP_STATE_CHANGE | TCP state transition |
| MOS_ON_CONN_END | Connection termination |
| MOS_ON_CONN_NEW_DATA | Availability of new flow-reassembled data |
| MOS_ON_ORPHAN | Out-of-flow (or non-TCP) packet arrival |
| MOS_ON_ERROR | Error report (e.g., receive buffer full) |

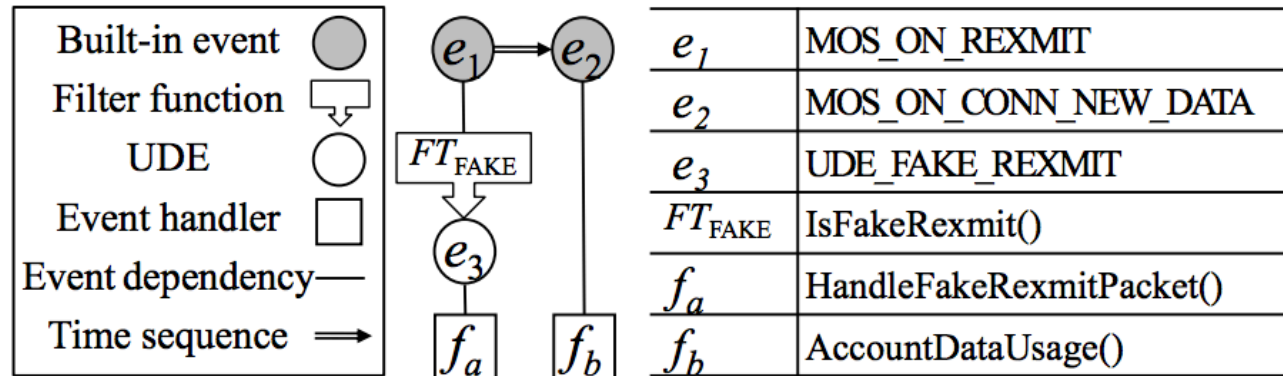
Table 1: mOS built-in events for stream monitoring sockets. Raw monitoring sockets can use only MOS_ON_PKT_IN raised for every incoming packet.

Built-in event triggers filter function

```
1 static bool
2 IsFakeRexmit(mctx_t mctx, int sock, int side, event_t event,
3 struct filter_arg *arg) {
4     struct pkt_info pi;
5     char buf[MSS];
6     struct tcp_ring_fragment frags[MAX_FRAG_NUM];
7     int nfrags = MAX_FRAG_NUM;
8     int i, size, boff, poff;
9
10    // retrieve the current packet information
11    mtcp_getlastpkt(mctx, sock, side, &pi);
12
13    // for full retransmission, compare the entire payload
14    if (mtcp_peek(mctx, sock, side, buf, pi.payloadlen, pi.offset) ==
15        pi.payloadlen)
16        return memcmp(buf, pi.payload, pi.payloadlen);
17
18    // for partial retransmission, compare the overlapping region
19    // retrieve the data fragments and traverse them
20    mtcp_getsockopt(mctx, sock, SOL_MONSOCKET, (side == MOS_SIDE_CLI) ?
21        MOS_FRAGINFO_CLIBUF : MOS_FRAGINFO_SVRBUF, frags, &nfrags);
22
23    for (i = 0; i < nfrags; i++) {
24        if ((size = CalculateOverlapLen(&pi, &(frags[i]), &boff, &poff)))
25            if (memcmp(buf + boff, pi.payload + poff, size))
26                return true; // payload mismatch detected
27    }
28    return false;
29 }
```


Filter function generates user-defined events

- If the filter function evaluates to true, a user-defined event is generated and the corresponding callback is called.



Build event-action diagram.

```
1 static void
2 mOSAppInit(mctx_t m)
3 {
4     monitor_filter_t ft = {0};
5     int s; event_t hev;
6
7     // creates a passive monitoring socket with its scope
8     s = mtcp_socket(m, AF_INET, MOS SOCK_MONITOR_STREAM, 0);
9     ft.stream_syn_filter = "dst net 216.58 and dst port 80";
10    mtcp_bind_monitor_filter(m, s, &ft);
11
12    // sets up an event handler for MOS_ON_REXMIT
13    mtcp_register_callback(m, s, MOS_ON_REXMIT, MOS_HK_RCV, OnRexmitPkt);
14
15    // defines a user-defined event that detects an HTTP request
16    hev = mtcp_define_event(MOS_ON_CONN_NEW_DATA, IsHTTPRequest, NULL);
17
18    // sets up an event handler for hev
19    mtcp_register_callback(m, s, hev, MOS_HK_RCV, OnHTTPRequest);
20 }
```

Figure 2: Initialization code of a typical mOS application. Due to space limit, we omit error handling in this paper.

Per-core, share nothing design

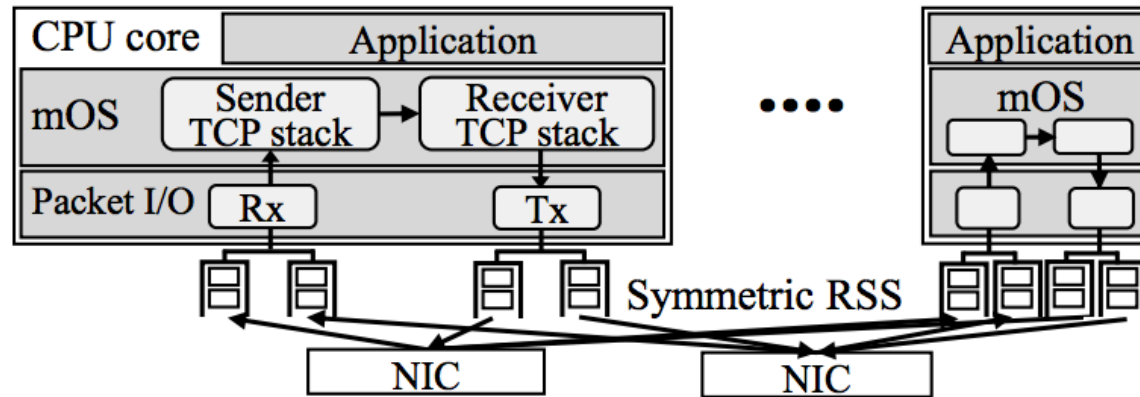


Figure 9: mOS application threading model

Per-flow TCP context

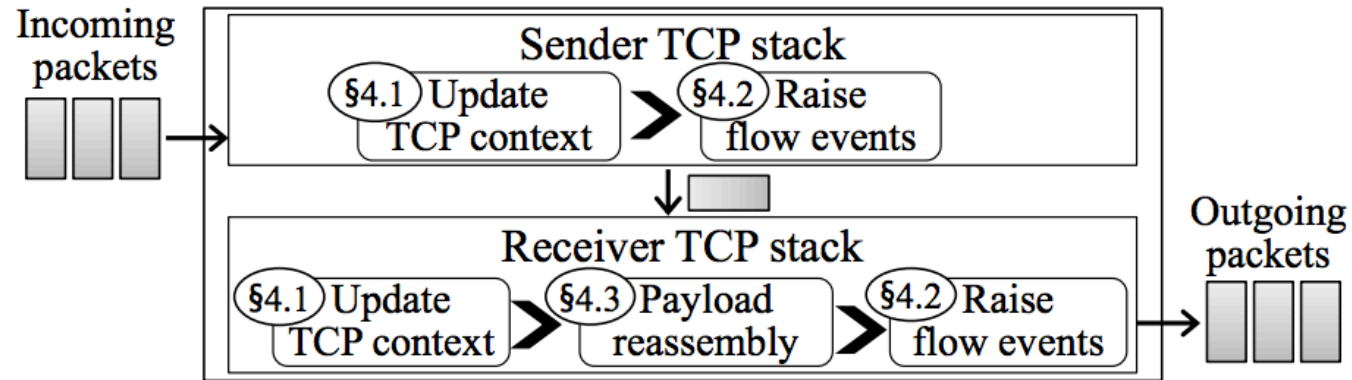
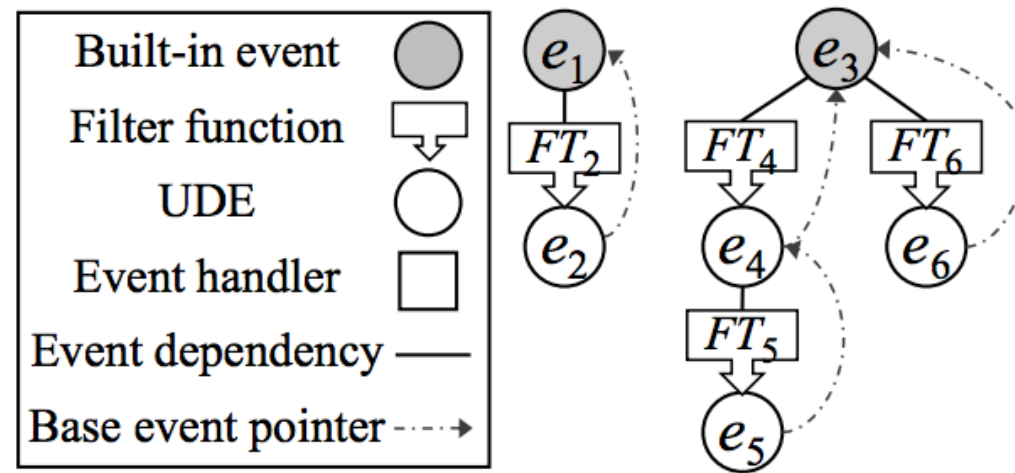


Figure 5: Packet processing steps in mOS

Scalable event management


- Problem:
- The flow can dynamically register event by itself.
- Maintaining a dedicated event diagram for each flow introduces huge overheads.
- Solution: Share event diagram, allocate new event diagram when needed.

Global event dependency forest, d-forest




(a) Global event dependency forest (d-forest)

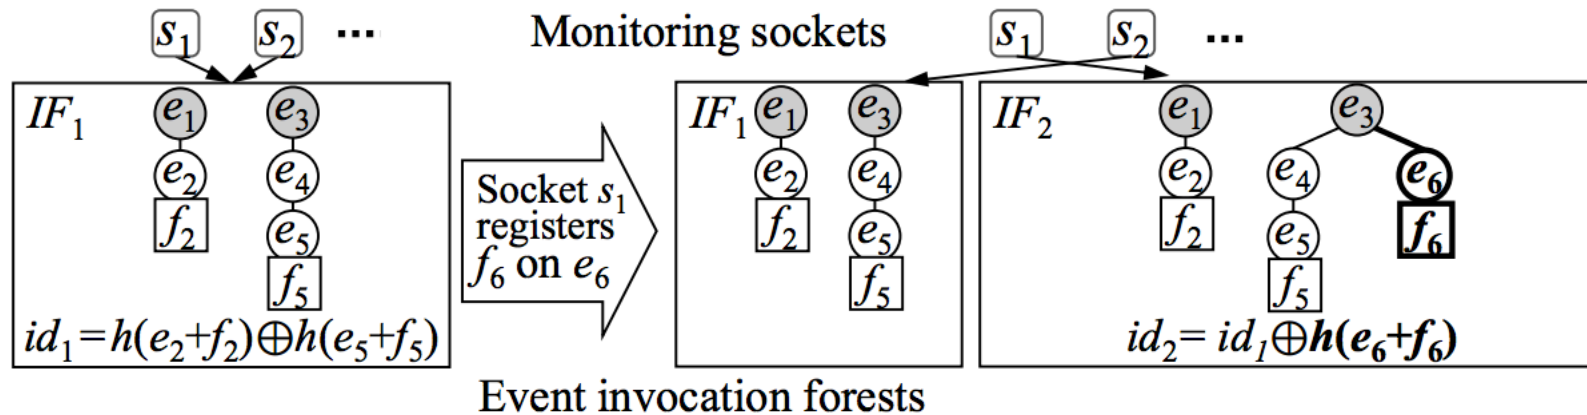
Event invocation forest, i-forest




When a new flow context is created



Then the flow may register additional event



Efficient querying of existing i-forest



Port existing middleboxes

- Snort3
- Abacus (re-implement in 400 LoC)
- Halfback
- nDPI.
- PRADs

Performance

| Application | original + pcap | original + DPDK | mOS port |
|-------------|-----------------|-----------------|------------|
| Snort-AC | 0.51 Gbps | 8.43 Gbps | 9.85 Gbps |
| Snort-DFC | 0.78 Gbps | 10.43 Gbps | 12.51 Gbps |
| nDPIReader | 0.66 Gbps | 29.42 Gbps | 28.34 Gbps |
| PRADS | 0.42 Gbps | 2.05 Gbps | 2.02 Gbps |
| Abacus | - | - | 28.48 Gbps |

Table 3: Performance of original and mOS-ported applications under a real traffic trace. Averaged over five runs.

Micro benchmark




Figure 12: Performance at dynamic event registration