

Kubernetes

Peng Yanghua

Agenda


- ▶ Background
 - ▶ Architecture
 - ▶ Core Concepts
 - ▶ Kubernetes for distributed machine learning
- 

Background



Old way: applications on host


- ▶ Applications and OS share filesystem.
 - Libraries, configurations, resources entangled with each other and the host
- ▶ VM
 - Heavy-weight



New way: applications in container

▶ Container

- OS-level virtualizations
- Isolated from each other and the host
- Small and fast
- 1 app to 1 image



Need a container-centric platform

- ▶ Management unit: host → container
- ▶ Automate orchestration for scale, just like Openstack for virtual machines.


Kubernetes

- ▶ A platform for automating container deployment, scaling and operations.
- ▶ Based on Google's 15-year experience on Borg.
- ▶ Open source, written in Go language.



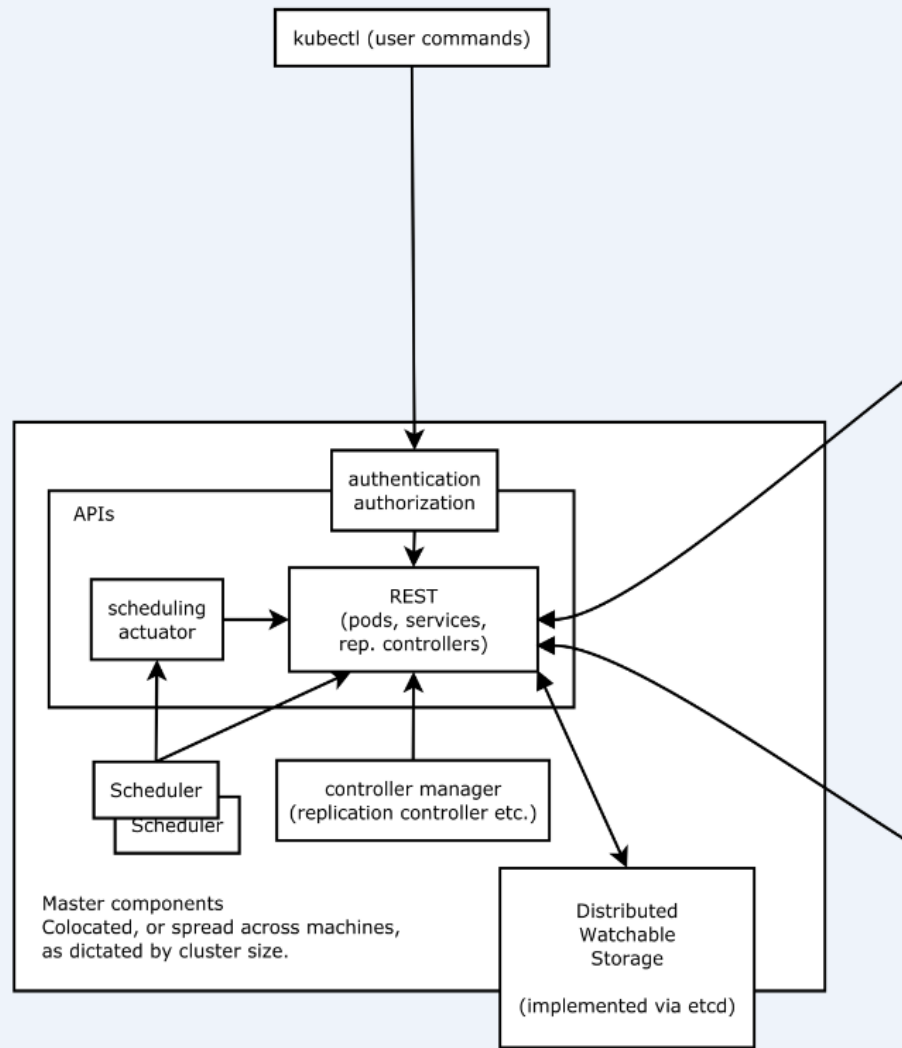
Architecture

High-level architecture



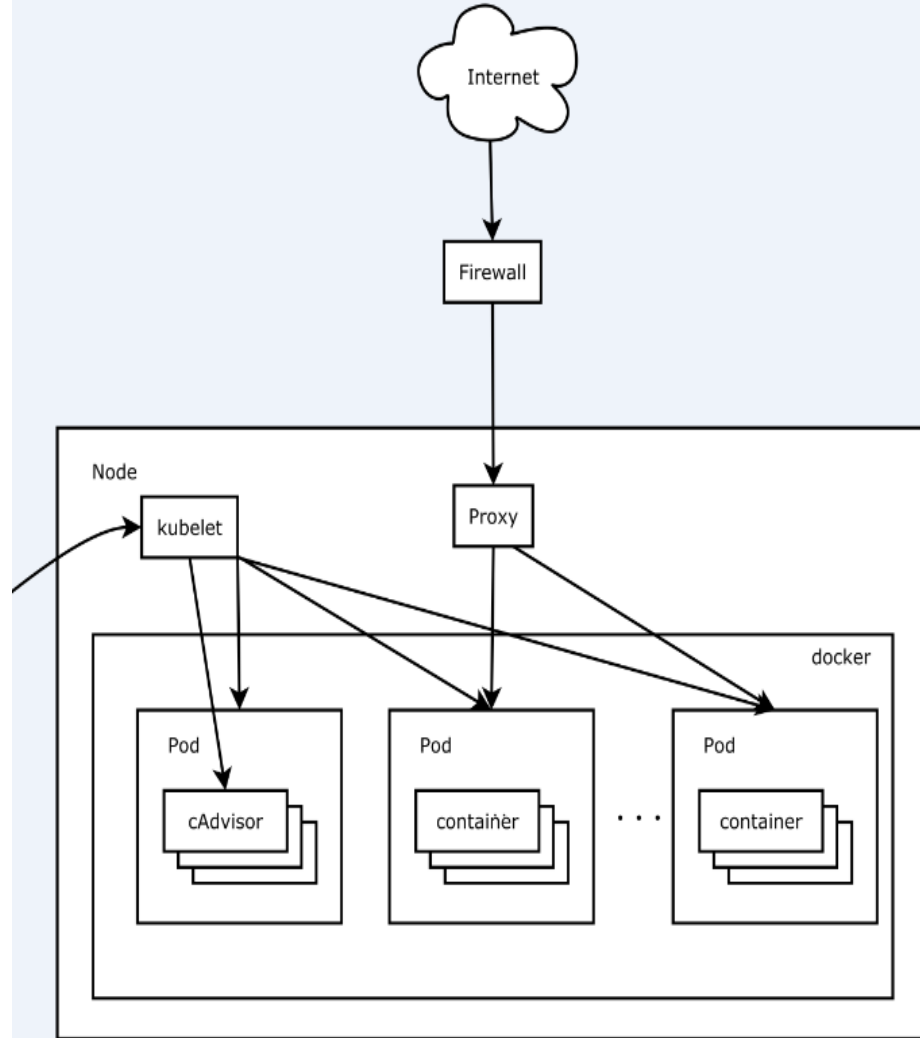
Master

- ▶ REST API
 - object operations
 - authentication, authorization
- ▶ Scheduler
 - random
 - round-robin
- ▶ Controller manager
 - node controller
 - replication controller
- ▶ Storage
 - etcd, distributed reliable



Node

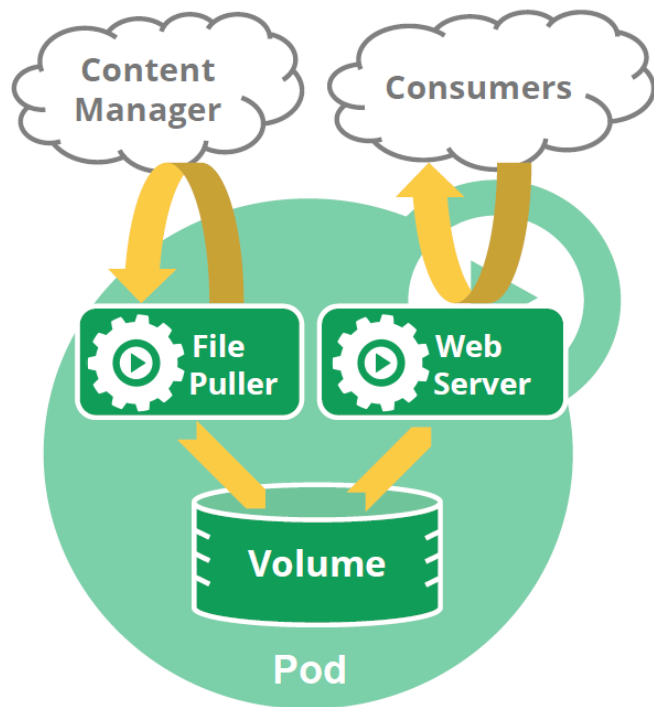
- ▶ kubelet
 - node agent
- ▶ kube proxy
 - user requests
- ▶ Docker
 - container runtime
- ▶ cAdvisor
 - container monitoring
- ▶ Flannel
 - inter-node overlay



Core concepts

Pods

- ▶ A group of containers sharing common resources
 - e.g., IP, filesystem.
 - usually different components of an app.
- ▶ Pod-level scheduling
 - mostly one container one pod




Volumes

- ▶ Pod-level shared storage
 - Any containers within the pod can access it.
 - Communication



Labels

- ▶ A key-value pair attached to an object
 - query-able by **selectors**
 - the **only** grouping mechanism




app = my-app, tier = FE

Replications

- ▶ Ensure N copies of a pod
 - or auto-scaling
 - grouped by a label selector
- ▶ Rolling updates
 - +1 / -1
 - without downtime


ReplicationController

- selector = {"app": "my-app"}
- template = { ... }
- replicas = 4



Services

- ▶ A group of same pods that work together
 - grouped by a label selector
 - a virtual IP for client access
 - discover via DNS



Jobs

- ▶ A group of pods that run to completion
 - pods do not always restart on failures
 - usually no external IP for access


Job

- parallelism: 3
- completions: 6
- selector:
 - job: my-work




DaemonSets

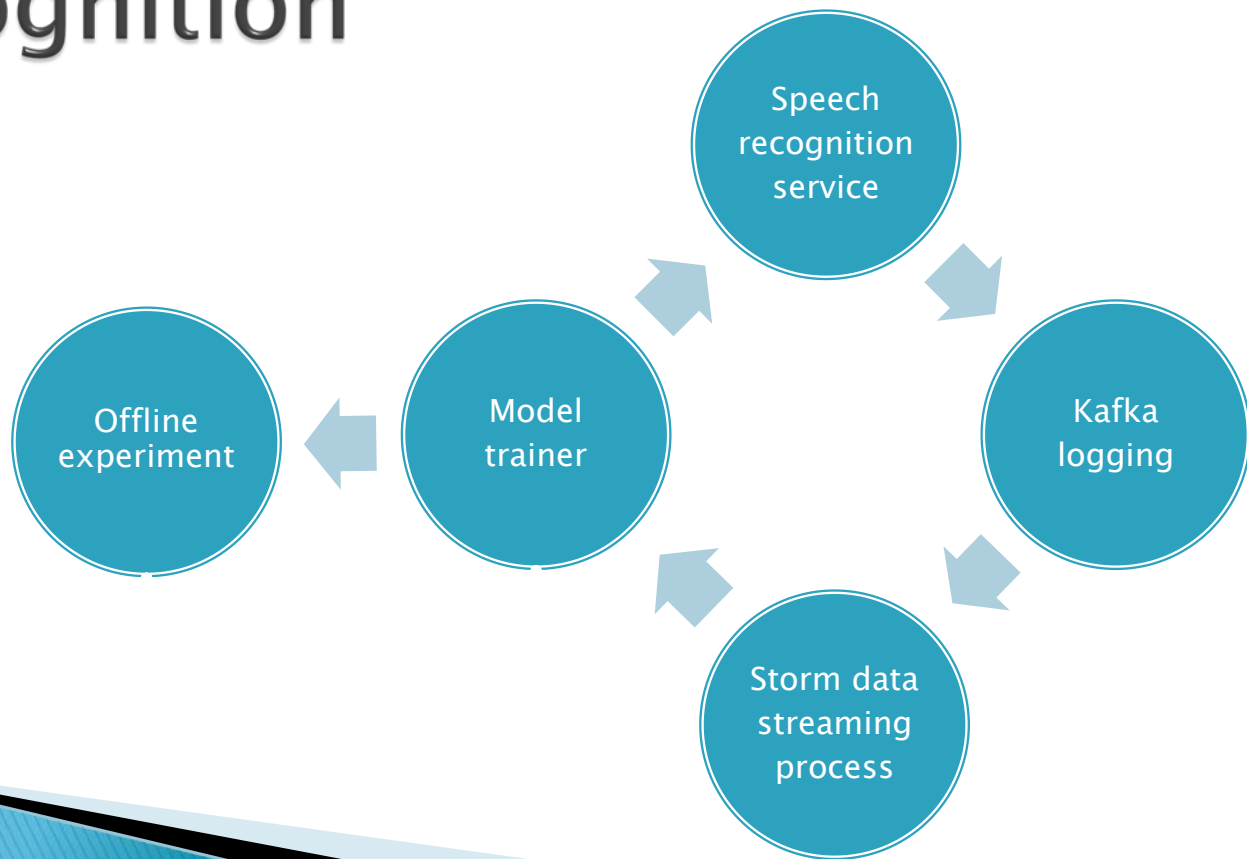
- ▶ Run a pod on every node
 - created as nodes come and go
 - useful for running cluster-wide services, e.g., logging, storage.



Kubernetes for Distributed Machine Learning



Workload example: speech recognition



Why kubernetes?


- ▶ A platform

- Deep learning system
- A lot of others, such as web server, the logger, data processor etc.

- ▶ Online and offline jobs

- Online: serve production traffic
- Offline: experiments

kubernetes: running different kinds of workloads efficiently in a cluster.



But kubernetes is not so perfect

- ▶ Resource allocation is static.
 - Manually change resource requirements if having more resources
- ▶ Deep learning job configuration
 - Best number of parameter servers and trainers
 - Where to deploy
- ▶ GPU support is not enough.
- ▶ **We need a job scheduler running on kubernetes in machine learning clusters for better utilization and efficiency.**

Thanks

Backup

Comparison with Mesos

- ▶ Resource isolation and sharing across distributed frameworks, e.g., spark, kubernetes.

Kubernetes Architecture

