

PCC Vivace: Online-Learning Congestion Control

Mo Dong, Tong Meng, Doron Zarchy, Engin Arslan, Yossi Gilad, P. Brighten Godfrey and Michael Schapira

Why replace TCP

- ▶ TCP performs poorly on lossy links, penalizes high-RTT flows, underutilizes high bandwidth-delay product (BDP) connections
- ▶ TCP's rate control architecture made hardwired mapping: certain predefined packet-level events are hardwired to certain predefined control responses.
- ▶ TCP assumes that the loss indicates congestion in the network
 - ▶ -> a packet loss halves the congestion window.

Performance-oriented Congestion Control (PCC)

- ▶ Goal: understanding what rate control actions improve performance based on live experimental evidence, avoiding TCP's assumptions about the network
- ▶ Sends at a rate r for a short period of time
- ▶ Observes the results
 - ▶ e.g. SACKs indicating delivery, loss, and latency of each packet
- ▶ Maximize the utility function
 - ▶ Achieving the objective like “high throughput and low loss rate”

Original PCC

- ▶ Utility function
 - ▶ $u_i(x_i) = T_i \cdot \text{Sigmoid}_\alpha(L_i - 0.05) - x_i \cdot L_i$
 - ▶ x_i : rate, T_i : throughput, L_i : Loss rate
- ▶ Starting: Doubles its rate at each monitor interval (MI)
 - ▶ Until utility decreases
- ▶ Decision Making State: Try $(1 + \varepsilon)x$ and $(1 - \varepsilon)x$ each for 2 MI
 - ▶ Choose a changing direction
- ▶ Rate Adjusting State:
 - ▶ Keep increase or decrease rate $x \leftarrow (1 + \varepsilon)x$ or $x \leftarrow (1 - \varepsilon)x$ until utility decreases
 - ▶ Then enter Decision Making State again

PCC Vivace

- ▶ Utility function considering change of RTT
 - ▶
$$u \left(x_i, \frac{d(RTT_i)}{dT}, L_i \right) = x_i^t - bx_i \frac{d(RTT_i)}{dT} - cx_i \times L_i,$$
 - ▶ x_i : rate, L_i : Loss rate
 - ▶ t, b, c : pre-defined parameters.
- ▶ A stable global rate configuration (Nash equilibrium) always exists for homogeneous agents
- ▶ “no-regret” guarantee: No worse than the best fixed rate policy
- ▶ Uses gradient-ascent algorithms for quick convergence

Fairness

- ▶ Regarded as a game, this problem is “socially-concave”
- ▶ Theoretically, consider a network model with n senders competing on a bottleneck link with a FIFO queue

Theorem 1. *When n Vivace-senders share a bottleneck link, and each Vivace-sender i 's utility function is defined as in Eq. 1, the senders' sending rates converge to a fixed configuration (x_1^*, \dots, x_n^*) such that $x_1^* = x_2^* = \dots = x_n^*$.*

Rate-control algorithm

- ▶ Start state:
 - ▶ Double sending rate every MI
 - ▶ Exit when utility drops
- ▶ Online Learning State:
 - ▶ Compute gradient
 - ▶ In the next two MIs, try $(1 + \varepsilon)r$ and $(1 - \varepsilon)r$, get two utility u_1, u_2 .
$$\gamma = \frac{u_1 - u_2}{2 \varepsilon r}$$
 - ▶ Choose step size
 - ▶ Monotonically non-decreasing function $m(\tau)$ for integer τ . $m(0) = 1$
 - ▶ $\Delta r = m(\tau)\theta_0\gamma$, for τ consecutive increases or consecutive decreases
 - ▶ Clip the gradient, *dynamic change boundary*
 - ▶ $\Delta r \leftarrow \min\{\omega r, \Delta r\}$
 - ▶ $\omega = \omega_0 + k \cdot \delta$ for k consecutive changes exceeding ωr

Contending with Unreliable Statistics

- ▶ Estimate the RTT gradient via linear regression
 - ▶ Use the linear-regression-generated slope as the RTT gradient
- ▶ Low-pass filtering of RTT gradient
 - ▶ Ignore small RTT gradient to ignore Non-congestion- induced latency jitters
- ▶ Double checking abnormal measurements
 - ▶ Sending faster would not lead to lower loss
 - ▶ Thus re-calculate the gradient.
- ▶ MI timeout
 - ▶ Network condition changes may cause large number of lost packets.
 - ▶ Halves the rate when no packets can provide information in $T_{timeout}$

Implementation, Evaluation

Vivace Parameters

MI duration	1 RTT
sampling step ε	0.05
initial conversion factor θ_0	1
initial dynamic boundary ω_0	0.05
dynamic boundary increment δ	0.1
RTT gradient filter threshold $flt_{latency}$	0.01
MI timeout $T_{timeout}$	4 RTT
confidence amplifier $m(\tau)$	$\tau \quad (\tau \leq 3)$ $2\tau - 3 \quad (\tau > 3)$

Table 1: Vivace's rate control default parameters

Consistent High Performance

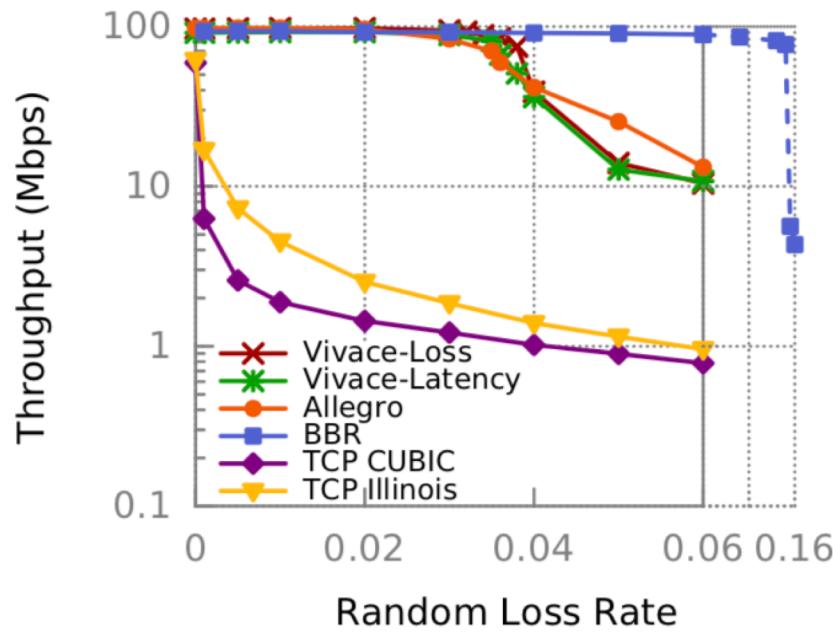


Figure 2: Random loss resilience

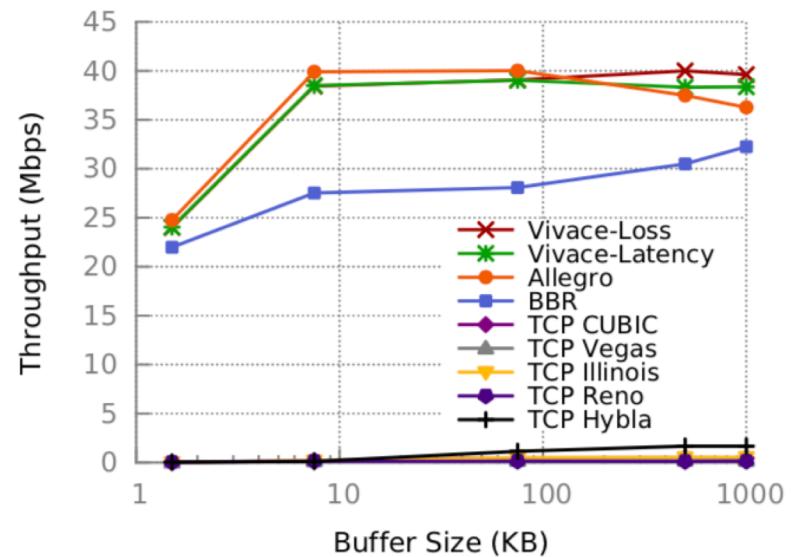
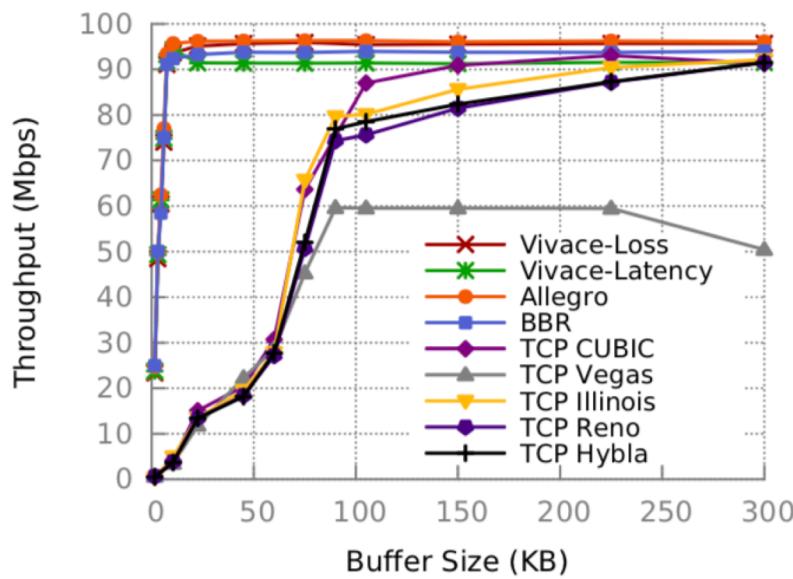
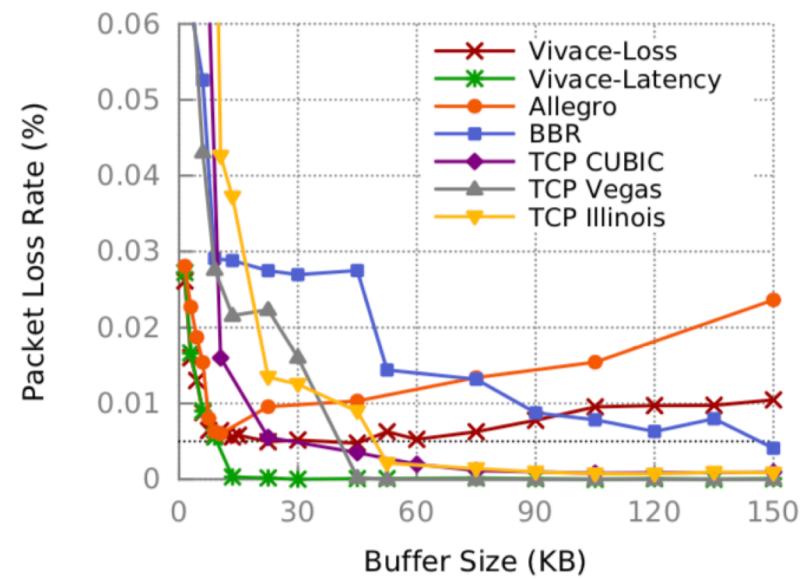


Figure 3: Long RTT tolerance

Influence of Buffer Size

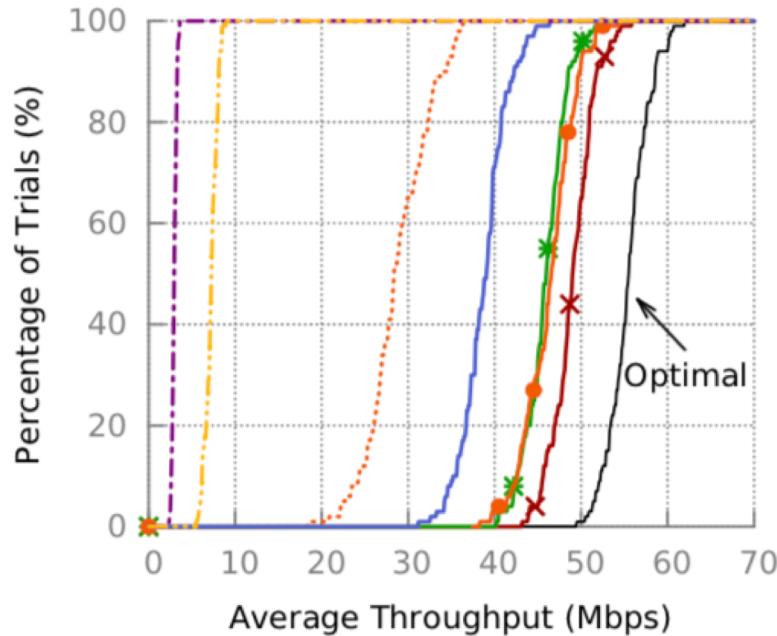


(a) High capacity

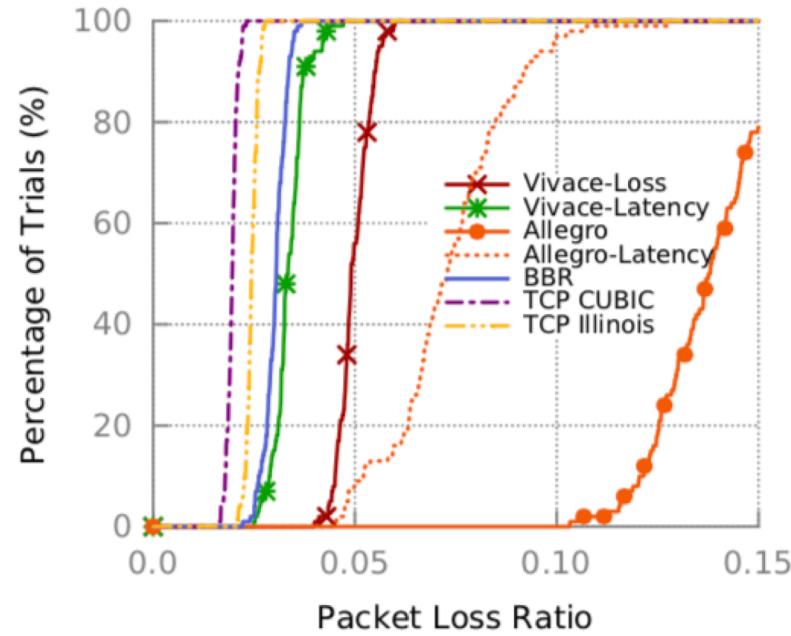


(b) Low packet losses

In rapid changing network condition



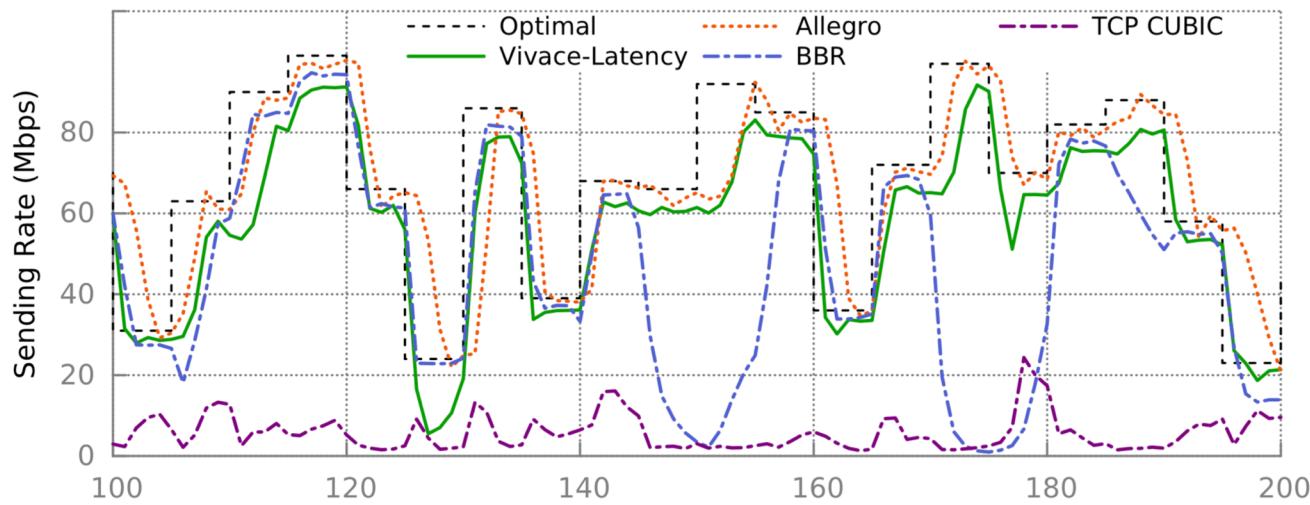
(a) Achieving high capacity



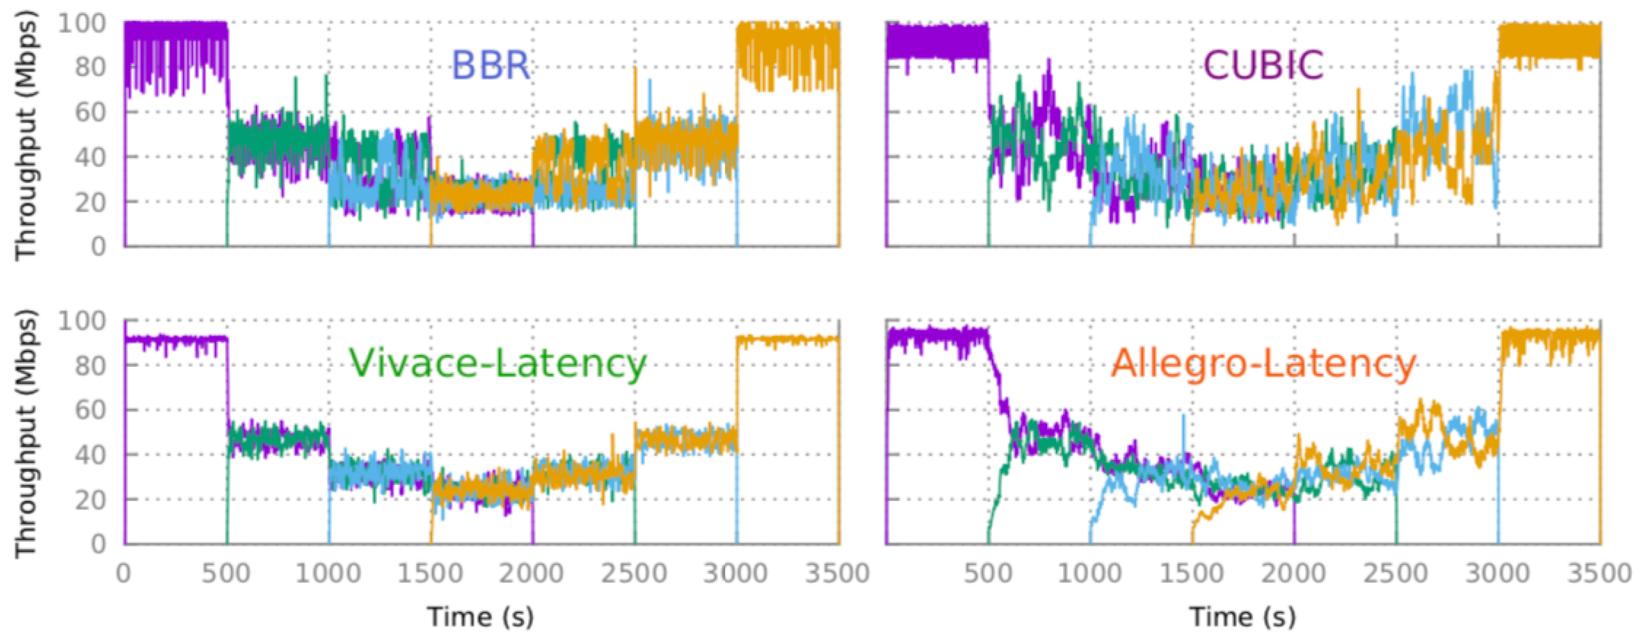
(b) Moderate packet loss

RTT, bottleneck bandwidth, and random loss rate all change every 5 seconds with uniform distribution ranging from 10-100 ms, 10-100 Mbps, and 0-1%

Cont'd

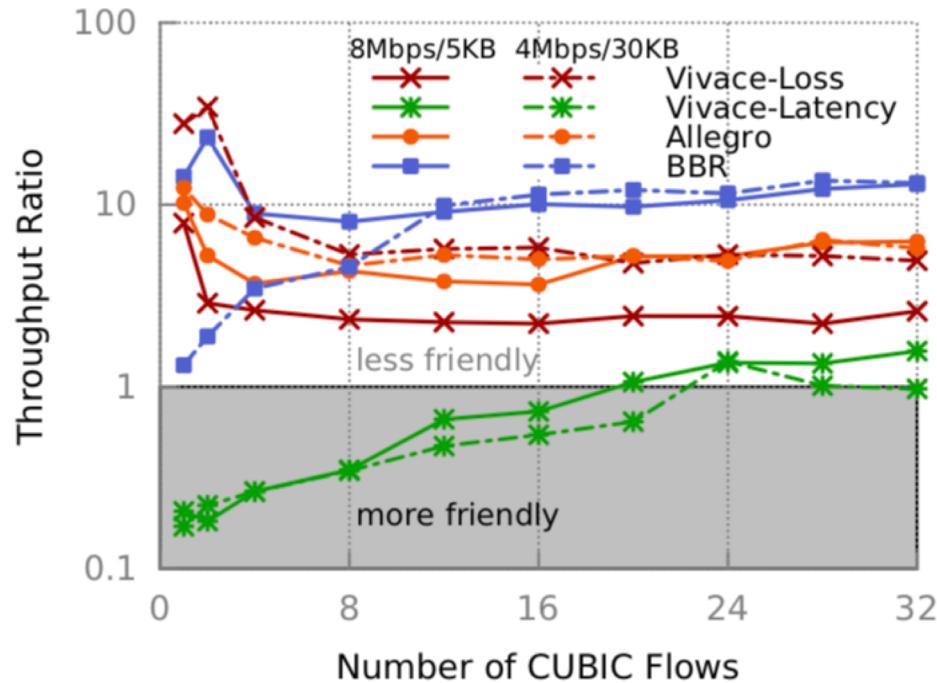


Convergence Properties



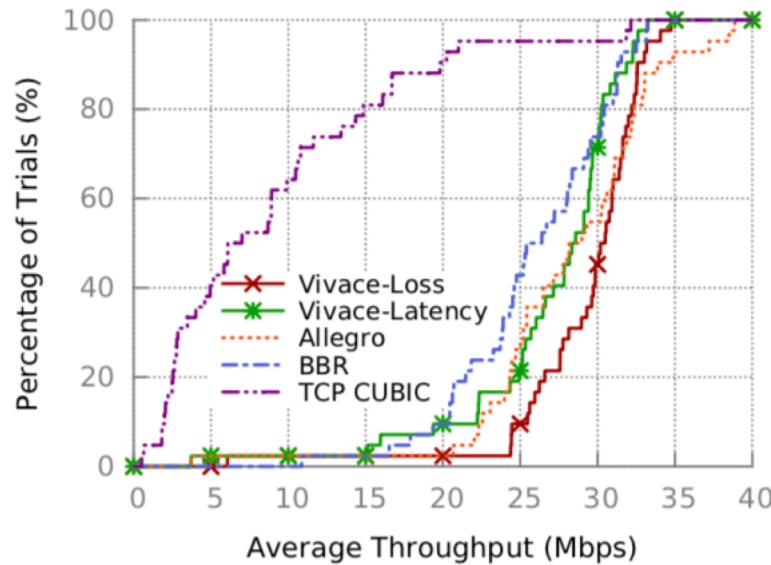
4 flows sharing a link with 100 Mbps bandwidth, 30 ms RTT,
and 75 KB buffer

Improved Friendliness to TCP



Increase number of CUBIC flows, compare the ratio between the throughput of the new-protocol flow and average throughput per CUBIC flow

Performance in real world



(c) Throughput from home networks to AWS

3 senders in different Wi-Fi network, 14 receivers in AWS

Summary

- ▶ A PCC protocol considering latency
 - ▶ Gradient-ascent-based online learning rate control to achieve provably fast convergence and fairness guarantees.
 - ▶ Sufficient Experiments
-
- ▶ Online Learning can be powerful when handling production environment.
 - ▶ There can be lots of implementation details that need consideration when implementing algorithms.