

Device Placement Optimization with Reinforcement Learning

Google Brain

ICML 2017

Background


- Neural networks
 - Image classification, speech recognition, machine translation
- Training and inference with neural networks
 - A growth in size and computational requirements
- Heterogeneous distributed environment
 - A mixture of hardware devices, CPUs and GPUs
- Placing certain operations of the neural models on devices

Challenges

- When the network has many branches
- When the mini batches get larger
- Dynamic environment with many interferences

Overview

- Which parts of the model should be placed on which device
- How to arrange the computations to optimize the communication
- Reward signal: execution time



Method

- A TensorFlow computational graph G
- M operations $\{o_1, o_2, \dots, o_M\}$
- D available devices
- Placement $P = \{p_1, p_2, \dots, p_M\}$ is an assignment of an operation o_i to a device p_i
- $r(P)$ is the time to perform a complete execution of G under the placement P
- Goal: minimize the execution time $r(P)$

Method

- Noisy measurements of $r(P)$ in the beginning of the training process
 - The bad placements sampled
 - Inappropriate learning signals
- RL model gradually converges
 - The sampled placements become more similar to each other
 - Less distinguishable training signals
- Empirically use the square root of running time
$$R(p) = \sqrt{r(P)}$$

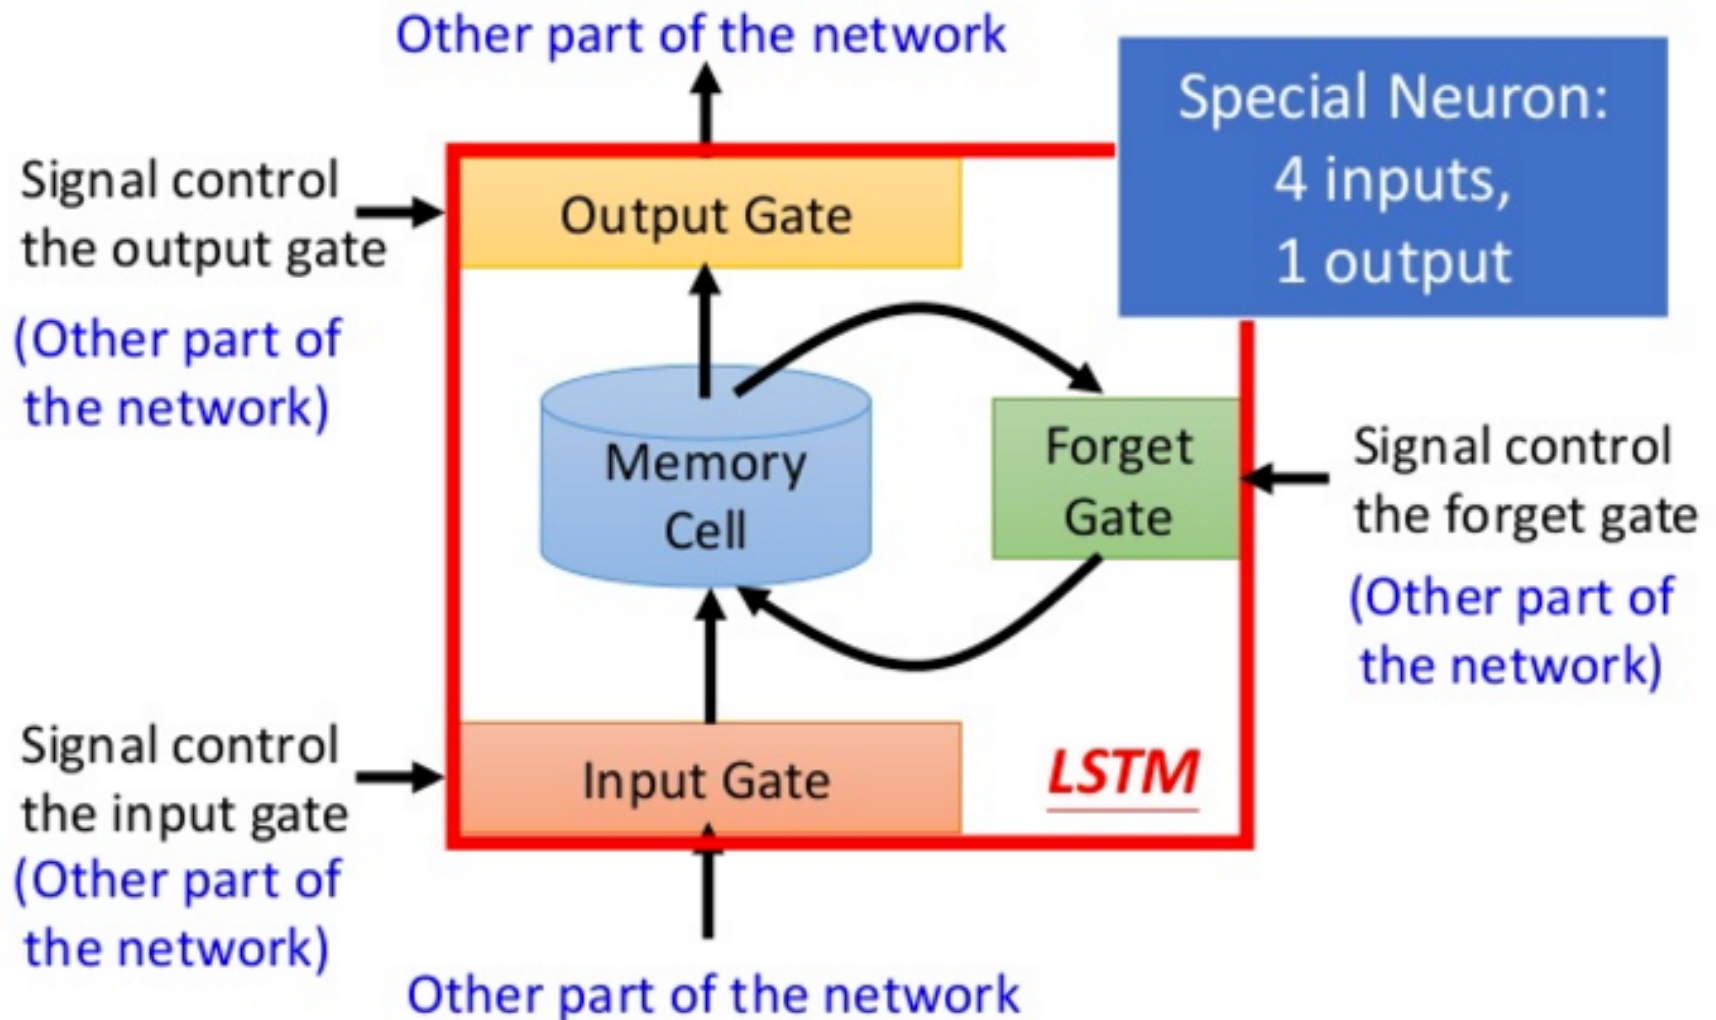
Method

- Train a stochastic policy $\pi(P|G; \theta)$ to minimize the objective $J(\theta) = E_{P \sim \pi(P|G; \theta)}[R(P)|G]$
- Policy gradient
 - Optimize parametrized policies with respect to the expected return by gradient descent
 - Learn network parameters
- $\nabla_{\theta} J(\theta) = E_{P \sim \pi(P|G; \theta)}[R(P) \cdot \nabla_{\theta} \log_p(P|G; \theta)]$
- Estimate $\nabla_{\theta} J(\theta)$ by drawing some placement samples
- $\nabla_{\theta} J(\theta) = \frac{1}{K} \sum_{i=1}^K (R(P_i) - B) \cdot \nabla_{\theta} \log_p(P_i|G; \theta)$

LSTM (Long Short Term Memory)

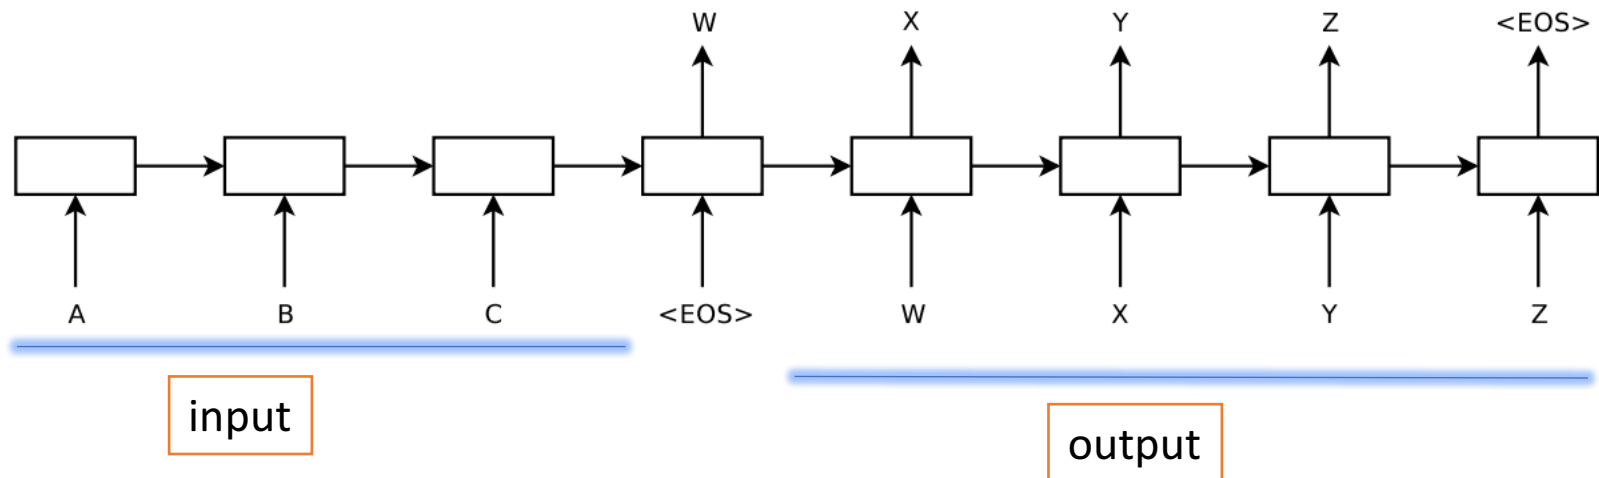
- Learn long-term dependencies
- What information to throw away from the cell state
- What information to store in the cell state
 - Which value to update
 - Create a vector of new candidates values
- Decide what information to output
 - A filtered version based on cell state

Long Short-term Memory (LSTM)



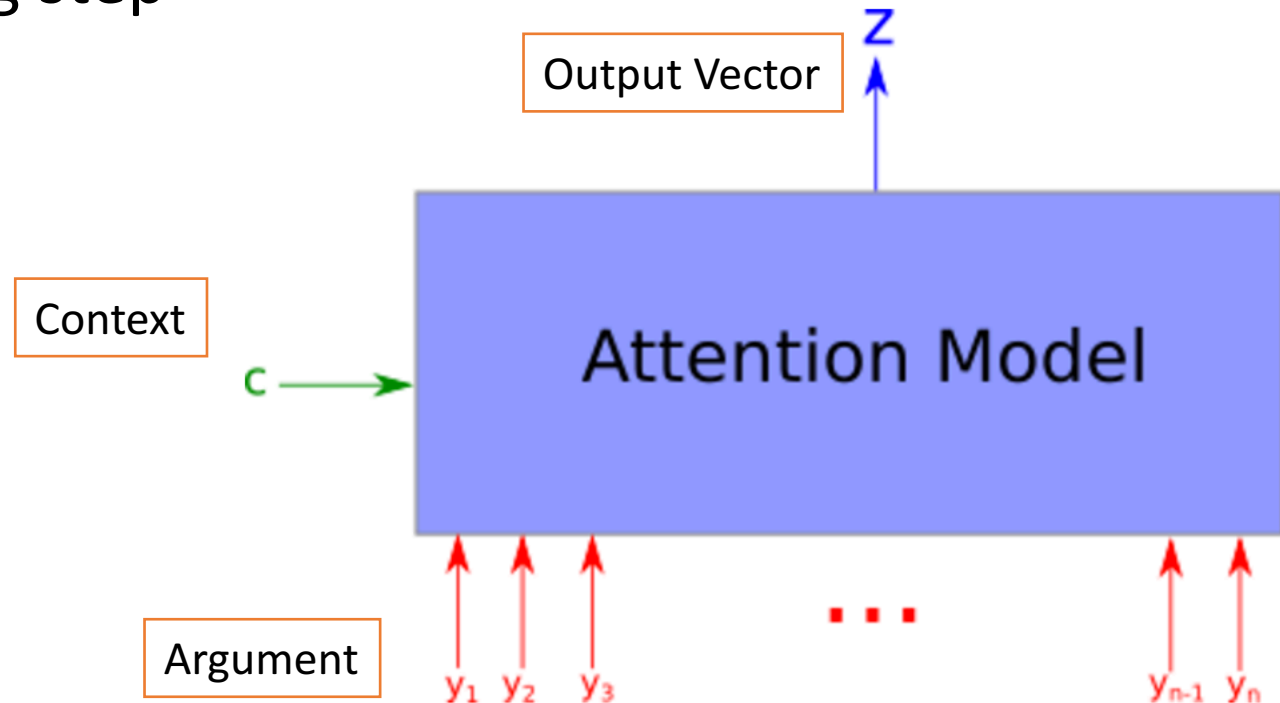
Sequence-to-sequence model

- Deep neural networks require the dimensionality of the inputs and outputs is known and fixed
- Can not map sequences to sequences
- Encoder: process the input
- Decoder: generate the output

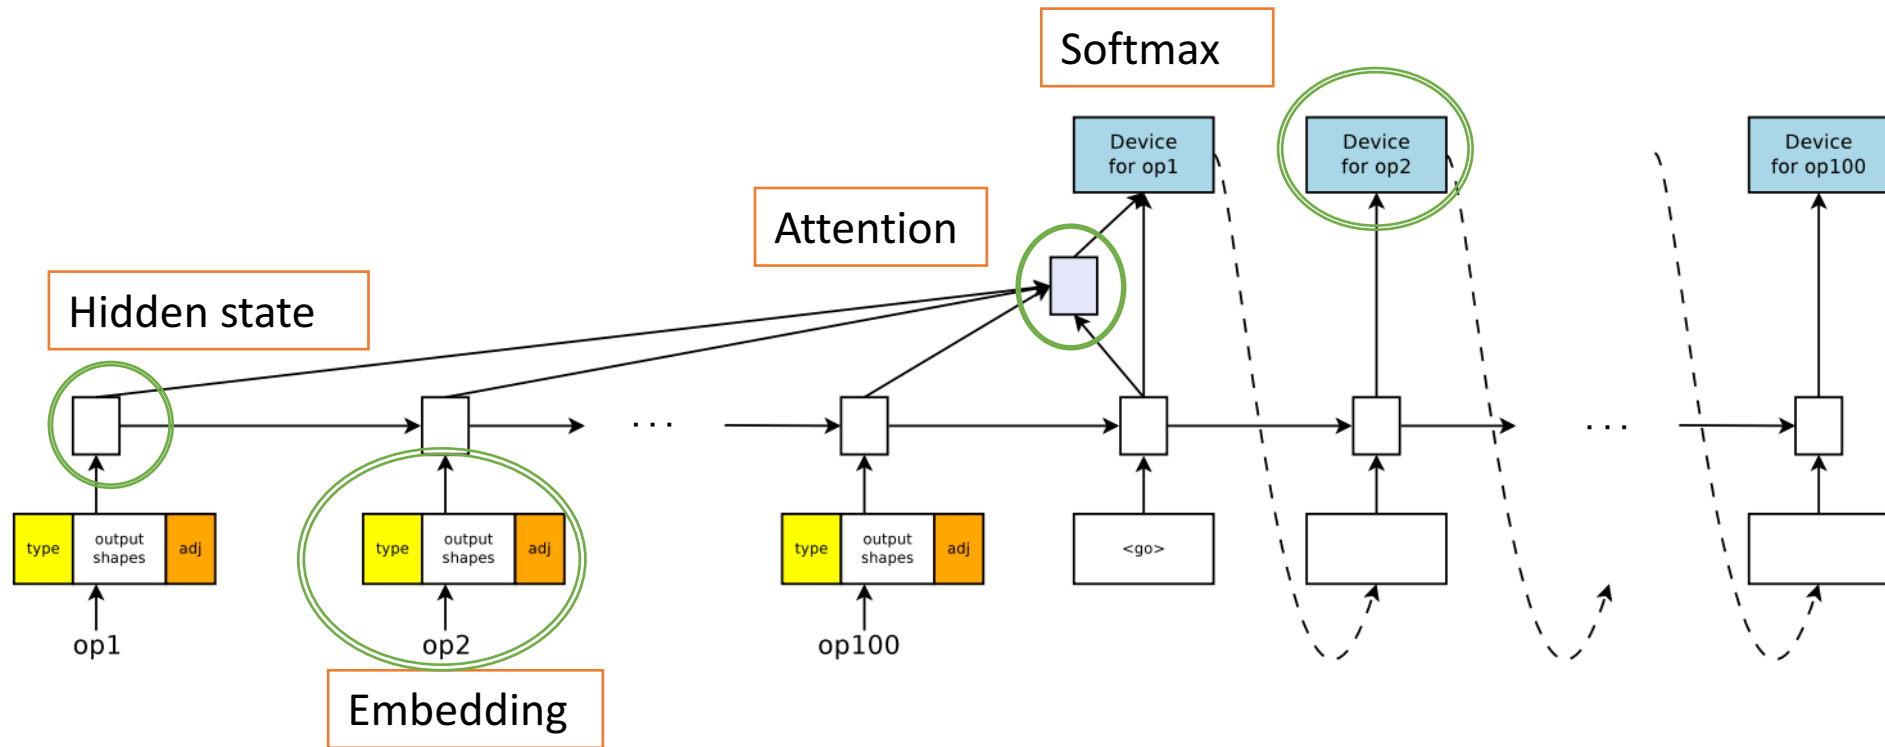


Content-based attention mechanism

- Allow the decoder more direct access to the input
- Allow the decoder to peek into the input at every decoding step



Architecture Details



Architecture Details

- Input:
- The sequence of operations of the input graph
- Encoder:
- Collect the types of its operations
- Concatenate the size of each operation's list into a fixed-sized, zero-padded list, i.e., output shape
- Take the one-hot encoding vector that represents the adjacency information

Architecture Details

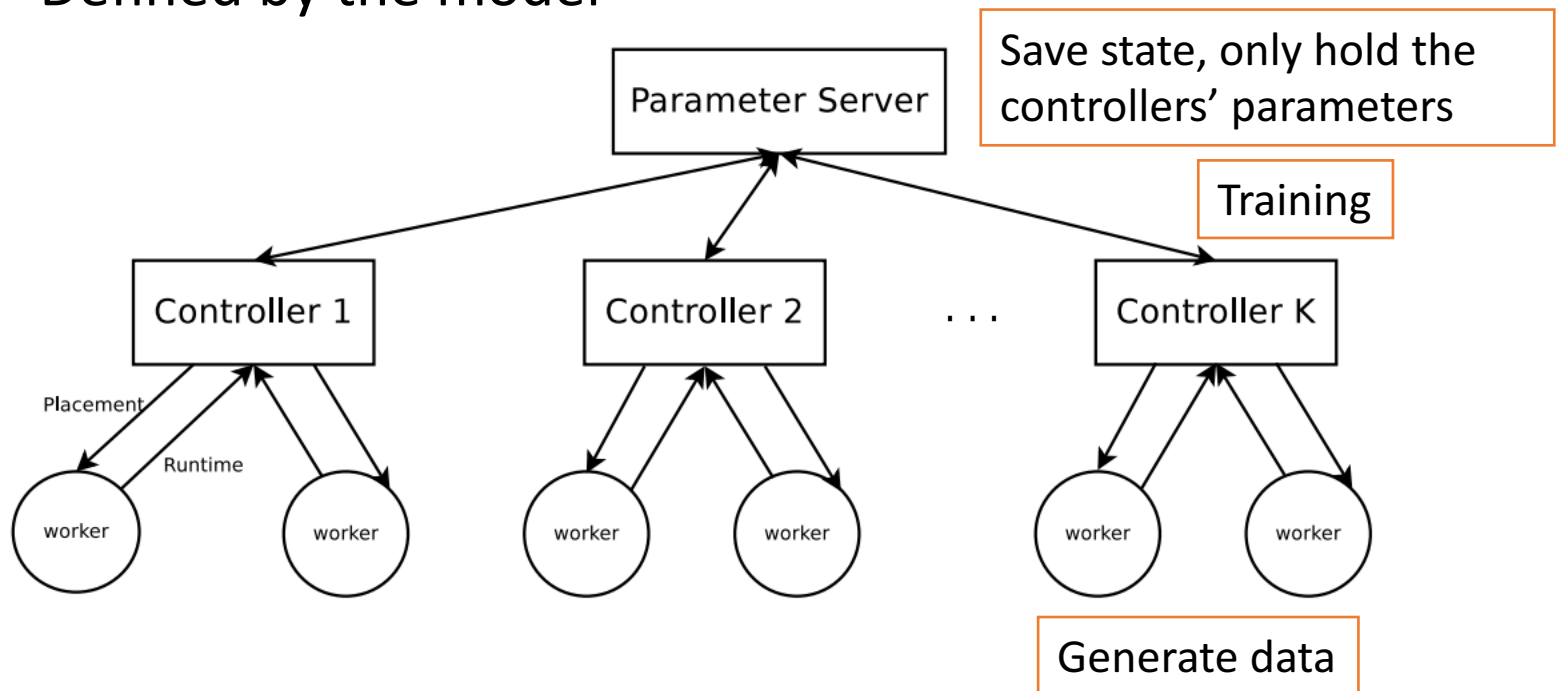
- Decoder:
- An attentional LSTM
- A fixed number of steps
 - Equal to the number of operations
- Output the device of the operator at the same encoder time step

Co-locating Operations

- TensorFlow computational graphs have thousands of operations
- Reduce the number of objects to place on different devices
 - Manually forcing several operations to be located on the same device
- E.g., the output of an operation X is consumed only by another operation Y \rightarrow operations X and Y are co-located
- E.g., initialization operations

Distributed Training

- Asynchronous distributed training
- Each controller execute the current policy
 - Defined by the model



First Phase

- Each worker receives a signal
 - Indicate that it should wait for placements from its controller
- Each controller receives a signal
 - Indicate it should sample K placements
- Each controller then independently sends the placements to their workers

Second Phase

- Each worker executes the placement and measures the running time
- Each controller waits for all of its workers
 - Finish executing their assigned placements
 - Return their running times
- The controller scales the gradients to asynchronously update the controller parameters

Benchmarks

- Recurrent Neural Network Language Model
 - Grid structure
 - Parallel executions
- Neural Machine Translation with attention mechanism
 - large number of hidden states due to the source and target sentences
 - Require model parallelism
- Inception-V3
 - Image recognition and visual feature extraction
 - Sequential processing

Baselines


- Single-CPU
- Single-GPU
- Scotch
 - Estimate the computational costs, amount of flow data
- MinCut
 - Run on GPU if no CPU implementation
- Expert-designed
 - Put each layer on a device
 - Partition the model into contiguous parts with roughly the same number of layers

Single-Step Runtime Efficiency

- Only information
 - Running times of the placement
 - The number of available devices
- Learn subtle tradeoffs
 - Performance gain by parallelism
 - The costs induced by inter-device communications


| Tasks | Single-CPU | Single-GPU | #GPUs | Scotch | MinCut | Expert | RL-based | Speedup |
|----------------------------|------------|-------------|-------|--------|--------|--------|-------------|---------|
| RNNLM (batch 64) | 6.89 | 1.57 | 2 | 13.43 | 11.94 | 3.81 | 1.57 | 0.0% |
| | | | 4 | 11.52 | 10.44 | 4.46 | 1.57 | 0.0% |
| NMT (batch 64) | 10.72 | OOM | 2 | 14.19 | 11.54 | 4.99 | 4.04 | 23.5% |
| | | | 4 | 11.23 | 11.78 | 4.73 | 3.92 | 20.6% |
| Inception-V3 (batch 32) | 26.21 | 4.60 | 2 | 25.24 | 22.88 | 11.22 | 4.60 | 0.0% |
| | | | 4 | 23.41 | 24.52 | 10.65 | 3.85 | 19.0% |

End-to-End Runtime Efficiency




Analysis of Found Placements

- Per-device computational load
 - Load-balancing



Analysis of Found Placements

- Memory copy time



Summary

- Weakness
 - No data parallelism
 - Overhead (hundreds of machines, one day)
 - Not general
- Use reinforcement learning to optimize placement in the machine learning system
 - Combine with theory work (bandit)
- Model parallelism
 - Complementary to current work