


Introduction to Reinforcement Learning

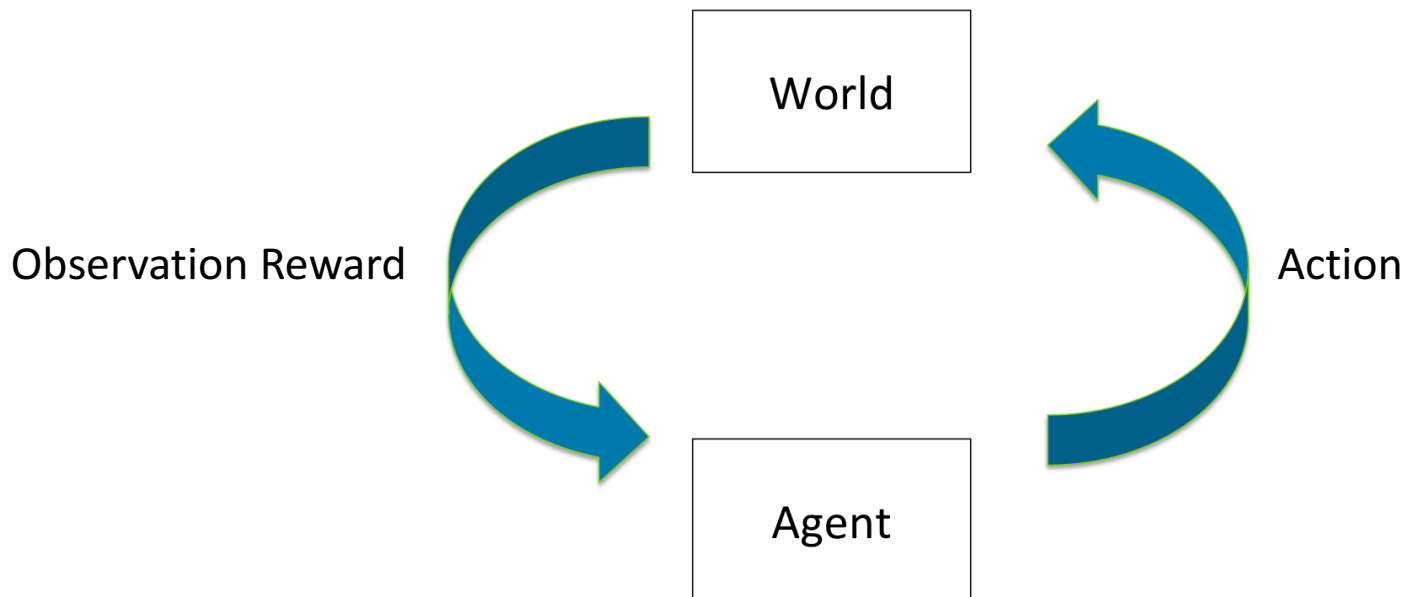
Overview

- Learn to make good **sequences** of decisions
- Don't know in advance how world works
- Repeated **interactions** with environment
- Reward for sequence of decisions



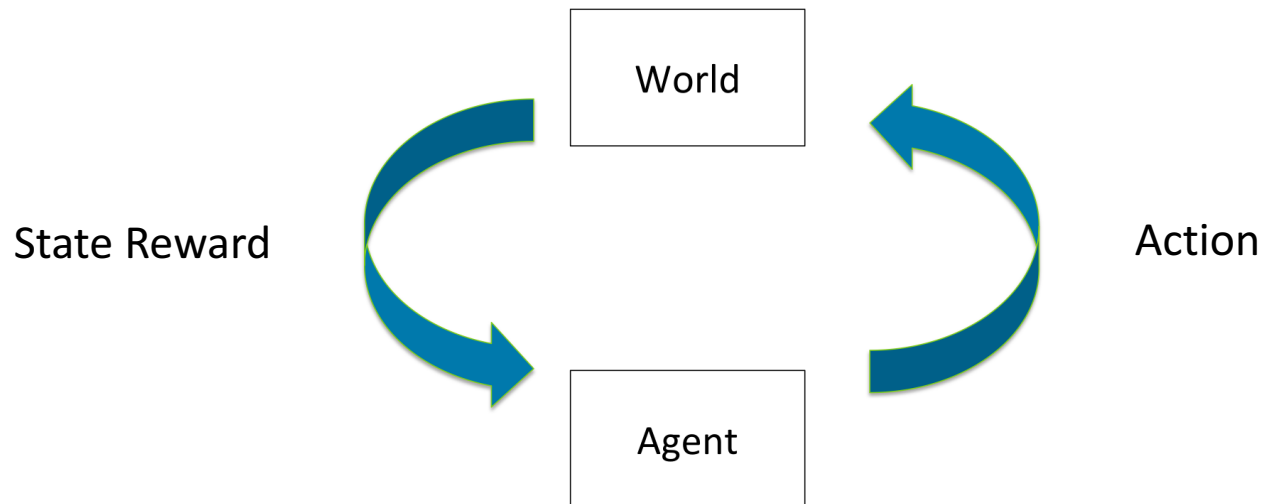
Decision Making Under Uncertainty

- Multi-armed bandit
- A set of arms to select
 - Obtain reward which follows **unknown distribution**
- Actions **do not change** the state of the world



Markov Decision Process

- Actions **change** the state of the world
- State = Observation
- Sufficient statistic that captures how world behaves
- Policy: mapping from state to action



Markov Decision Process: $\langle S, A, R, T, \gamma \rangle$

- S : set of states
- A : set of actions
- R : immediate reward $R(s) / R(s, a) / R(s, a, s')$
- T : dynamics model $p(s_{t+1} | s_t, a_t)$
- γ : discount factor
 - the difference in importance between future rewards and present rewards
- Memoryless
 - The outcome of an action depends only on the **current** state (vs entire history)
- Policy $\pi: S \rightarrow A$
 - Specifies what action to take in each state

MDP Policy Value

- For a given state s
- Value of policy $V^\pi(s)$: Expected discounted sum of rewards obtain if the agent follows policy π starting in state s
 - $V^\pi(s) = \mathbb{E}[\sum_{t=0}^{\infty} \gamma^t r(s_t, \pi(s_t)) | s_0 = s]$
- Optimal policy: $\operatorname{argmax}_{\pi} V^\pi(s)$
- **Immediate** reward + Discounted sum of **future** rewards
 - $V^\pi(s) = r(s, \pi(s)) + \gamma \sum_{s' \in \mathcal{S}} p(s' | \pi(s), s) V^\pi(s')$

Q: state-policy value

- Expected **immediate reward** for taking action a and expected **future reward** get after taking that action from that state and following π
- $Q^\pi(s, a) = r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|a, s) V^\pi(s')$

Optimal Value, Q & Policy

- Optimal V
 - Highest possible value for each s (under any possible policy)
 - Satisfies the Bellman Equation
 - $V^*(s) = \max_a [r(s, a) + \gamma \sum_{s' \in S} p(s'|a, s) V^*(s')]$
- Optimal Q function
 - $Q^*(s, a) = r(s, a) + \gamma \sum_{s' \in S} p(s'|a, s) V^*(s')$
- Optimal policy
 - $\pi^*(s) = \operatorname{argmax}_a Q^*(s, a)$

MDP Planning

- How to compute π^* ?
- Know full MDP
- Given the dynamics and reward model
 - Reward and state transition probability
- Computational challenge
 - Not learning

Value Iteration

- Bellman equation inspires an update rule
 - $V^*(s) = \max_a [r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|a, s) V^*(s')]$
- First compute value for each state as if only get to take 1 action
- Then what if take 2 actions...
- Estimate the optimal value
- Output
 - The solution: derive the optimal policy π^* from the optimal value
 - The discounted sum of the rewards to be earned (on average) by following that solution from state s

Value Iteration

1. Initialize $V_0(s) = 0$ for all states s ,
2. Set $k = 1$
3. Loop until [finite horizon, convergence]
 - For each state s
 - $V_{k+1}(s) = \max_a [r(s, a) + \gamma \sum_{s' \in S} p(s'|a, s) V_k(s')]$
4. Extract Policy $\pi(s)$

Policy Iteration

- Search **directly** for the optimal policy π^*
- Compute infinite horizon value of a policy
- Use to select another (better) policy
- Closely related to a very popular method in RL
 - policy gradient

Policy Iteration

1. Initialize $\pi_0(s)$ randomly for all states s
2. Loop until [finite horizon, convergence]
 - Policy evaluation: Compute V^{π_i}
 - Policy improvement:
 - Compute Q value of different 1st action and then following π_i
 - $Q^{\pi_i}(s, a) = r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|a, s) V^{\pi_i}(s')$
 - Use to extract a new policy
 - $\pi_{i+1}(s) = \operatorname{argmax}_a Q^{\pi_i}(s, a)$

Convergence

- Converge to a unique solution
 - for discrete state and action space
 - when $\gamma < 1$
 - all state-action pairs are visited infinitely often

Model-based Passive Reinforcement Learning

- Estimate MDP model parameters from data
 - Reward
 - State transition probability
- If finite set of states and actions
 - count & average
- Use estimated MDP to do policy evaluation of π

Model-free Passive Reinforcement Learning

- Only maintain estimate of Q value
- Temporal Difference learning
 - Approximate expectation with samples
 - Approximate future reward with estimate
- Maintain estimate of $V^\pi(s)$ for all states
 - Update $V^\pi(s)$ each time after each transition (s, a, s', r)
 - Likely outcomes s' will contribute updates more often
 - Approximate expectation over next state with samples

Q-learning

- Update $Q(s, a)$ every time experience (s, a, s', r)
- Create new sample estimate
 - $Q_{samp}(s, a) = r(s, a) + \gamma \max_{a'} Q(s', a')$
- Update estimate of $Q(s, a)$
 - $Q(s, a) = (1 - \alpha)Q(s, a) + \alpha Q_{samp}(s, a)$
- If acting randomly, Q-learning converges Q^*
- Optimal Q values
- Finds optimal policy

Simple Approach: ϵ -greedy

- With probability $1 - \epsilon$
 - Choose $\operatorname{argmax}_a Q(s, a)$
- With probability ϵ
 - Select random action
- Even after millions of steps still won't always be following argmax of $Q(s, a)$

Example

- State: the amount of occupied resource in the cloud
- Action: whether to accept newly arrived job
- Reward: revenue earned by the infrastructure provider
- Deterministic state transition
 - Consider the average time the system stays at state s

Scaling Up



- Want to be able to tackle problems with enormous or infinite state spaces
- Tabular representation is insufficient
- Don't want to have to explicitly store
 - dynamics or reward model
 - value
 - state-action value
 - policy
- for every single state

Generalization

- Smoothness assumption
- If two states are close, then (at least one of)
 - Dynamics are similar
 - Reward is similar
 - Q functions are similar
 - optimal policy is similar
- More generally, dimensionality reduction or compression
 - Unnecessary to individually represent each state
 - Compact representations possible

Function Approximation

- Key idea: replace lookup table with a function
- Replace table with **general parameterized** form



- Examples:
 - Linear combinations of features
 - Neural networks

Linear Value Function Approximation


- Represent state by a feature vector $x(s) = \begin{pmatrix} x_1(s) \\ x_2(s) \end{pmatrix}$
- For example
 - Distance of robot from landmarks
 - Trends in the stock market
 - Piece and pawn configurations in chess
- Represent value function by a **linear combination** of features
 - $\hat{v}(s, w) = x(s)^T w$

Linear Value Function Approximation

- Objective function is quadratic in parameters w
 - $J(w) = \mathbb{E}_{\pi}[(v_{\pi}(s) - x(s)^T w)^2]$
- Update = step-size \times prediction error \times feature value
 - $\Delta w = \alpha(v_{\pi}(s) - \hat{v}(s, w))x(s)$
 - Use historical average for $v_{\pi}(s)$

Deep Neural Networks

- Input activation
 - $a(x) = b + w^T x$
- Output activation
 - $h(x) = g(a(x)) = g(b + w^T x)$
- x : features
- w : weights (parameters)
- b : bias term
- $g(\cdot)$: activation function
 - Sigmoid function
 - Rectified linear function



Single Hidden Layer Neuro Net

- Hidden layer pre-activation


- $a(x) = b^{(1)} + w^{(1)T}x$

- Hidden layer activation

- $h(x) = g(a(x))$

- Output layer activation

- $f(x) = o(b^{(2)} + w^{(2)T}h^{(1)}(x))$



How to train neural nets

- To train a neural net, we need:
- Loss function
 - $l(f(x^{(t)}, \theta), y^{(t)})$
- A procedure to compute gradients
 - $\nabla_{\theta} l(f(x^{(t)}, \theta), y^{(t)})$
- Regularizer and its gradient
 - $\Omega(\theta), \nabla_{\theta} \Omega(\theta)$
 - Prevent overfitting

Stochastic Gradient Descent

- Perform updates after seeing each example
- For each training epoch
 - For each training example $(x^{(t)}, y^{(t)})$
 - $\Delta = -\nabla_{\theta} l(f(x^{(t)}, \theta), y^{(t)}) - \lambda \nabla_{\theta} \Omega(\theta)$
 - $\theta = \theta + \alpha \Delta$
- Backpropagation with gradient descent
 - Calculate the error contribution of each neuron after a batch of data is processed
 - From upper layer to lower layer

Example

- Feature: available resource, job resource profile
- Action: whether to schedule job in one time slot
- Reward: job completion time

Summary

- Standard Value iteration / Q learning is not very useful
 - State space is large
 - Convergence
 - Infinitely visiting each state-action pair
- Other work provide regret bound on modified Q learning
- DNN is useful
 - But there is no theoretical support behind

Thank you! ^^