

Model-Agnostic Meta-Learning for Fast Adaptation of Deep Networks

Chelsea Finn, Pieter Abbeel, Sergey Levine
University of California, Berkeley and OpenAI

ICML 2017

Versatility of intelligence

- The capability of doing many different things
- Current AI systems excel at mastering a **single** skill
- But struggle in doing a **variety** of simple problems
 - A champion can not hold a conversation
 - A expert helicopter controller for aerobatics can not navigate in new, simple situations

Meta-learning: Learning to learn

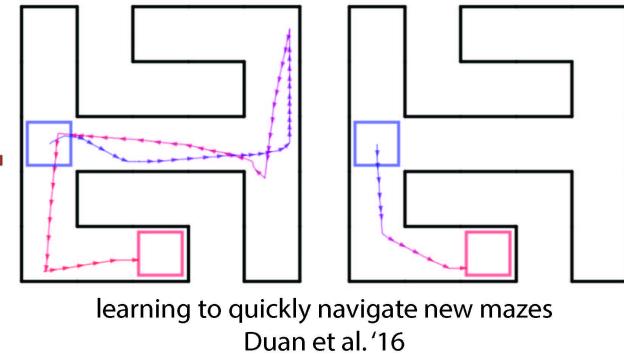
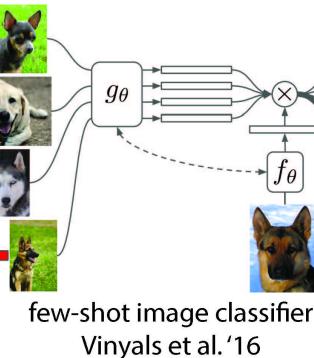
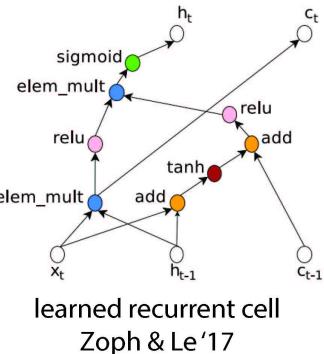
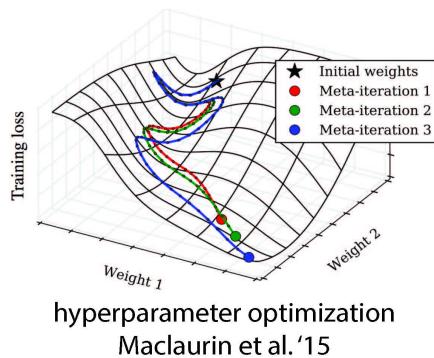
- Current AI systems
 - Master a complex skill from scratch
 - Using an understandably large amount of time and experience
- Acquire many skills and adapt to many environments
 - Cannot afford to train each skill in each setting from scratch.
- Learn how to **learn new tasks** faster by **reusing** previous experience
 - Rather than considering each new task in isolation

Why learning to learn

- Effectively reuse data on other tasks
- Replace manual engineering of architecture, hyperparameters, etc.
- Learn to quickly adapt to unexpected scenarios
 - Inevitable failures
 - Long tail
- Learn how to learn with weak supervision

The use of learning to learn

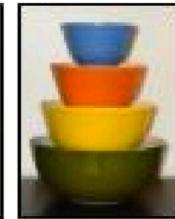
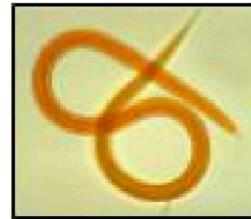
- Hyperparameter and neural network optimization
- Finding good network architectures
- Few-shot image recognition
- Learning to efficiently navigate a new maze with only one traversal through the maze



Typical Objective of Few-Shot Learning

Image recognition

Given 1 example of 5 classes:



Classify new examples



Human Concept Learning

Given 1 *positive* example:



Classify new examples:



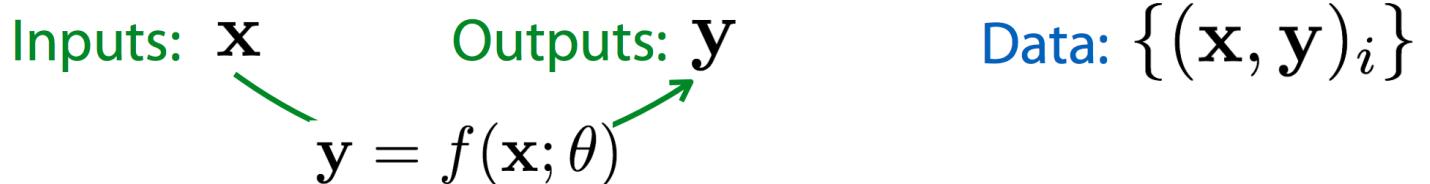
Beyond how humans learn, this setting is also more interesting.

Recent Meta-learning Approaches

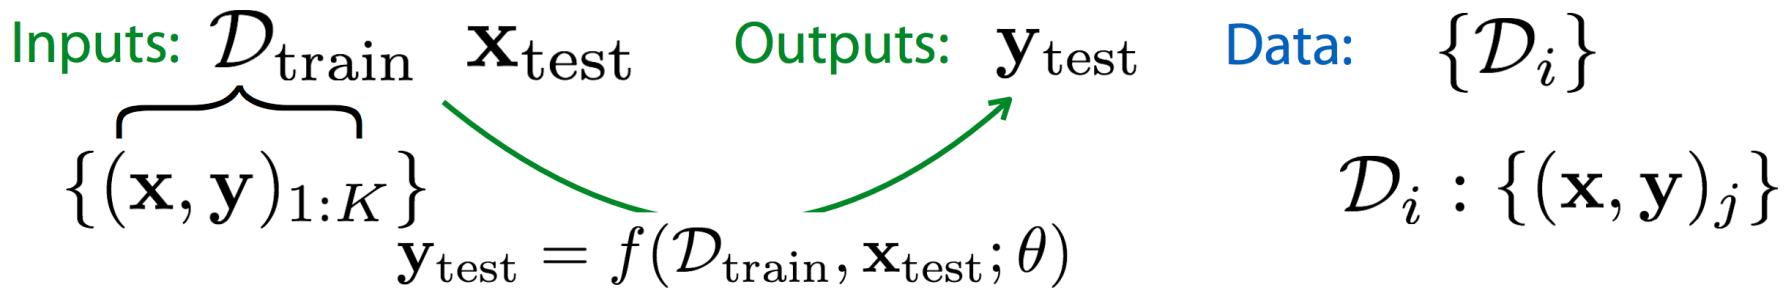
- Standard ML:
 - Training on **a single task** and testing on held-out examples from that task
- Meta-learning:
 - Training with **a large number of tasks**
 - Testing in their ability to learn **new tasks**
- Two optimizations at the same time
 - The learner: learn new tasks
 - The meta-learner: train the learner

The Meta-Learning Problem

Supervised Learning:

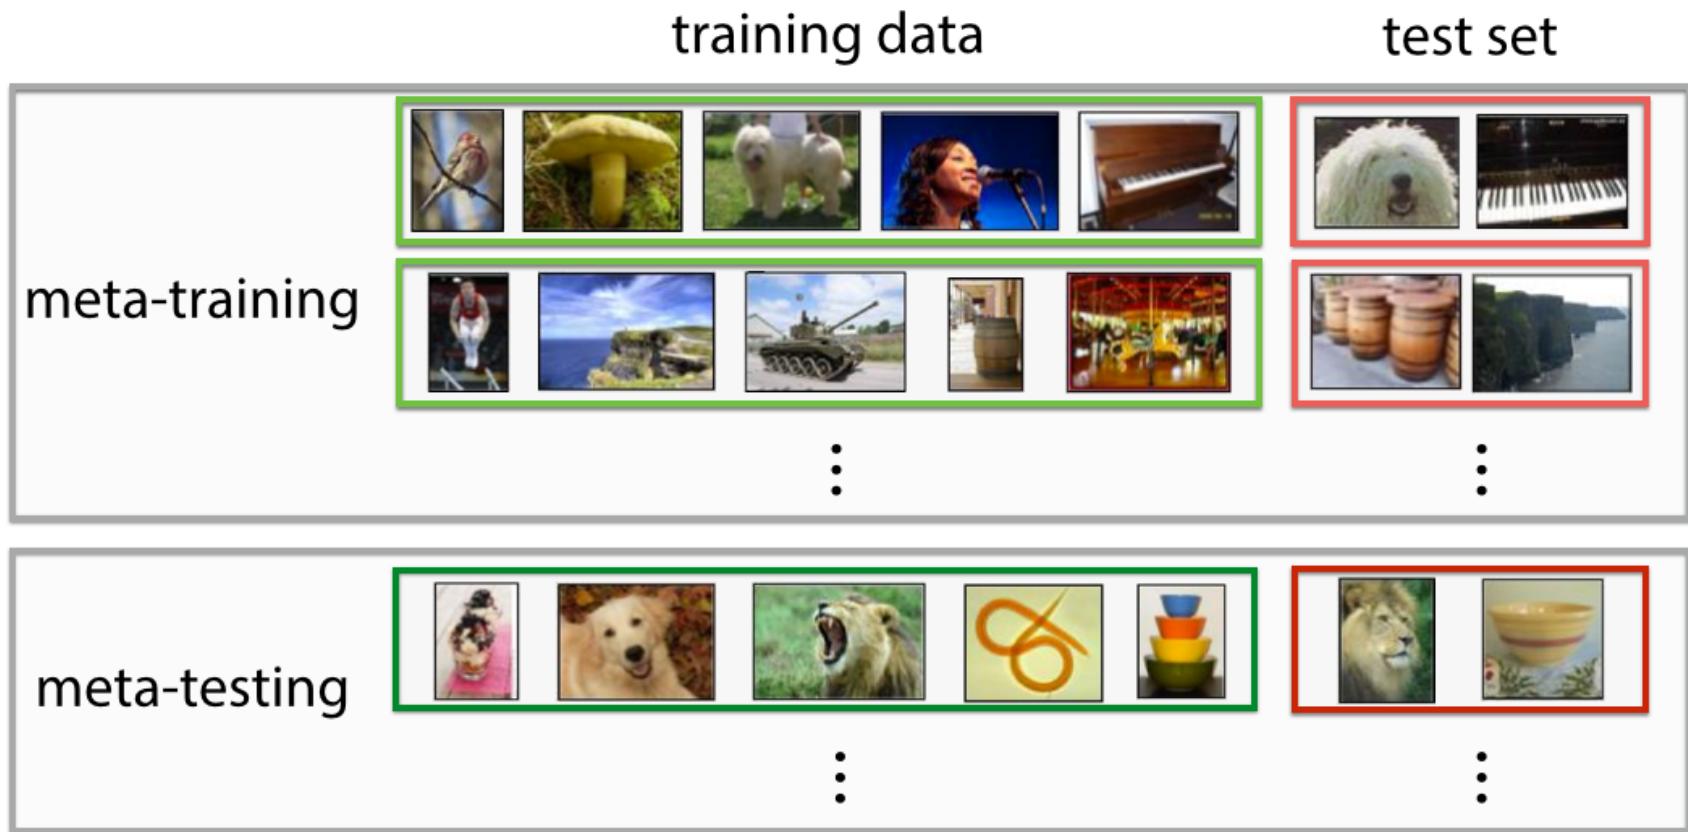


Meta-Supervised Learning:



Meta-learning Example

- Classifying a new image within 5 possible classes

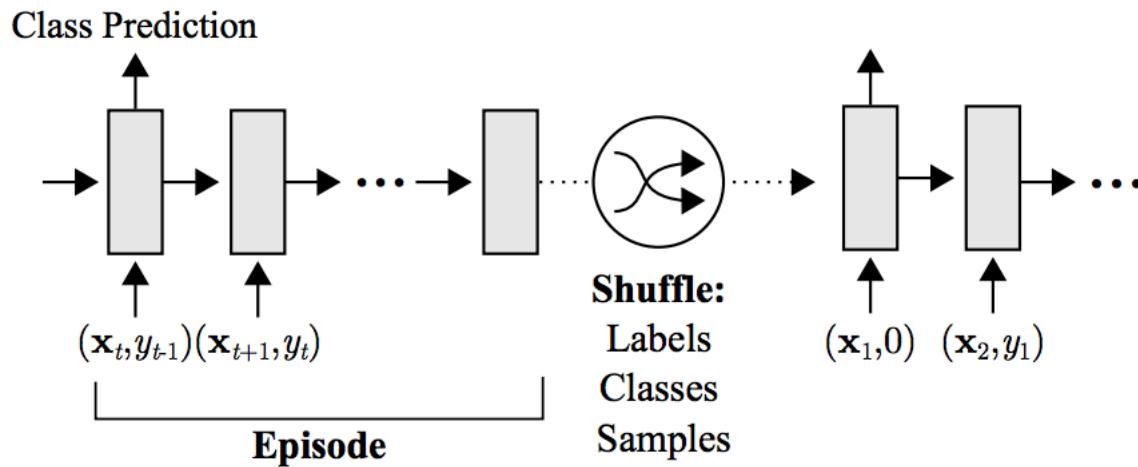


Recurrent Models

- Train a recurrent model
 - an LSTM
- Take in the dataset **sequentially** and then process new inputs from the task
- The recurrent learner is trained to adapt to new tasks as it is rolled out

Recurrent Models

- Image classification
 - Passing in the set of (image, label) pairs of a dataset sequentially
 - Followed by new examples which must be classified



Metric Learning

- Learn a **metric** in input space
 - A way of measuring the distance (or similarity) between objects
- Specialized to one/few-shot classification
- Can not be used in other problem
 - RL
- Meta-learner: gradient descent
- Learner: a comparison scheme
 - Nearest neighbors in the meta-learned metric space
 - Euclidean distance

Learning Optimizers

- Better neural network optimization
- Learn **parameter update** when given gradients
 - Search space includes SGD, RMSProp, Adam, etc.
- Meta-learner: a recurrent network
 - remember how it previously updated the learner model
 - Trained with reinforcement learning or supervised learning
- Few-shot image classification

Model-Agnostic Meta Learning

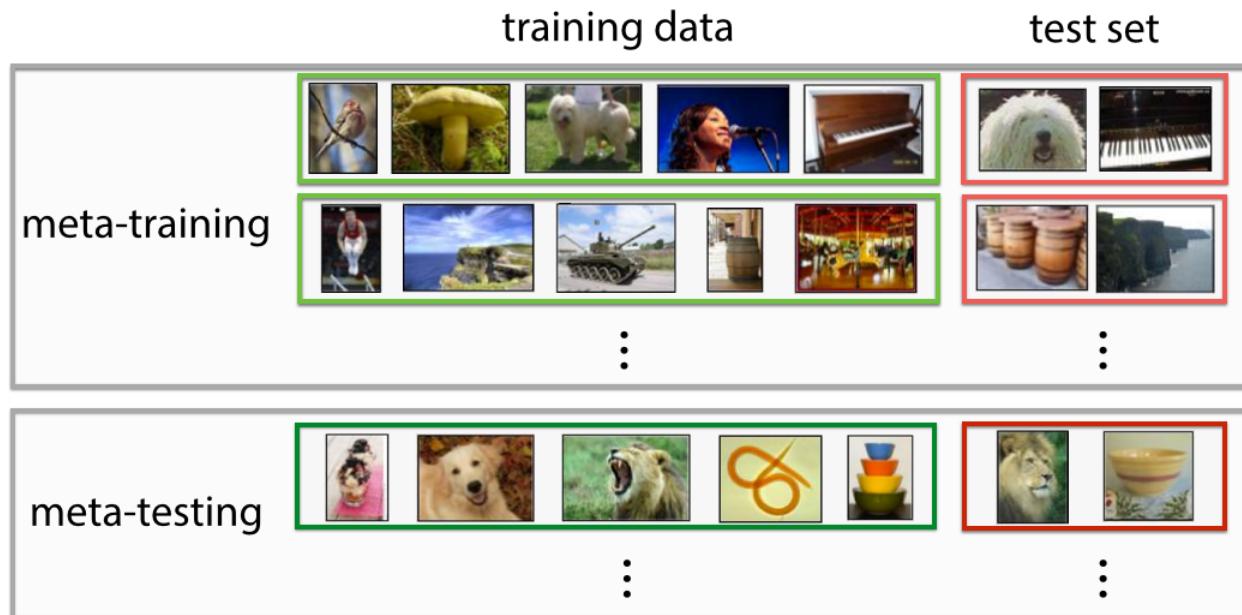
- Trained on a set of tasks
 - Adapt to **new tasks** quickly using only a small number of examples or trials
- Model f : map observations x to outputs a
- Learning task $\mathcal{T} = \{\mathcal{L}(x_1, a_1, \dots, x_H, a_H), q(x_1), q(x_{t+1}|x_t, a_t), H\}$
 - A loss function \mathcal{L} : a mis-classification loss or a cost function in MDP
 - A distribution over initial observations $q(x_1)$
 - A transition distribution $q(x_{t+1}|x_t, a_t)$
 - An episode length H

MAML Training Phase

- Consider a distribution over tasks $p(\mathcal{T})$
- Loop for:
 - Learn a **new task** \mathcal{T}_i drawn from $p(\mathcal{T})$ from only K samples and feedback
 - Test on new samples from \mathcal{T}_i
 - Consider how test error on new data changes w.r.t the parameters

Model-Agnostic Meta Learning

- Training error of meta learning process
 - Test error
- Meta-performance
 - The model's performance after learning from K samples

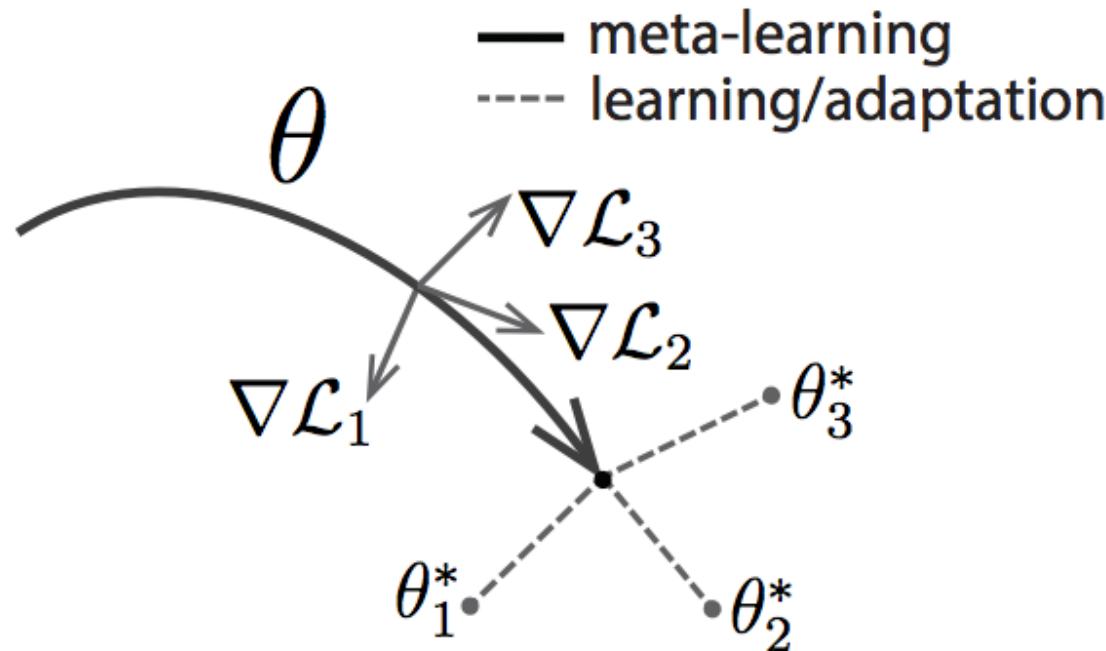


Algorithm

- Some internal representations are more transferable than others
- Learn internal features that are broadly applicable to all tasks in $p(\mathcal{T})$
 - Rather than a single individual task

Algorithm

- Find model parameters that are **sensitive** to changes in the task
 - Make rapid progress on new tasks
 - Small changes in the parameters will produce **large improvements** on the loss function of any task



fine-tuning: $\theta' \leftarrow \theta - \alpha \nabla_{\theta} \mathcal{L}(\theta)$

[test-time]

**Model-Agnostic
Meta-Learning:
(MAML)**

$$\min_{\theta} \sum_{\text{tasks}} \mathcal{L}_{\text{v}}(\theta - \alpha \nabla_{\theta} \mathcal{L}_{\text{tr}}(\theta))$$

Key idea: Train over many tasks, to learn parameter vector θ that transfers

In-distribution task: k-shot learning

Base case: learning from scratch

Related but out-of-distribution task: somewhere in between

Algorithm

Algorithm 1 Model-Agnostic Meta-Learning

Require: $p(\mathcal{T})$: distribution over tasks

Require: α, β : step size hyperparameters

- 1: randomly initialize θ
 - 2: **while** not done **do**
 - 3: Sample batch of tasks $\mathcal{T}_i \sim p(\mathcal{T})$ Sample different tasks
 - 4: **for all** \mathcal{T}_i **do**
 - 5: Evaluate $\nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$ with respect to K examples
 - 6: Compute adapted parameters with gradient descent: $\theta'_i = \theta - \alpha \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$ Update parameters of each task
 - 7: **end for**
 - 8: Update $\theta \leftarrow \theta - \beta \nabla_{\theta} \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i})$
 - 9: **end while**
-

Minimize the meta-objective: find parameter θ to minimize the loss of all tasks

Supervised Regression Classification

- Regression tasks: Use mean-squared error
- Discrete classification tasks: Use a cross entropy loss

Algorithm 1 Model-Agnostic Meta-Learning

Require: $p(\mathcal{T})$: distribution over tasks

Require: α, β : step size hyperparameters

- 1: randomly initialize θ
- 2: **while** not done **do**
- 3: Sample batch of tasks $\mathcal{T}_i \sim p(\mathcal{T})$
- 4: **for all** \mathcal{T}_i **do**
- 5: → Evaluate $\nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$ with respect to K examples
- 6: Compute adapted parameters with gradient descent: $\theta'_i = \theta - \alpha \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$
- 7: **end for**
- 8: Update $\theta \leftarrow \theta - \beta \nabla_{\theta} \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i})$
- 9: **end while**

Sample K data points using
 f_{θ} in \mathcal{T}_i

→ Evaluate $\nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$ with respect to K examples

Compute adapted parameters with gradient descent: $\theta'_i = \theta - \alpha \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$

Reinforcement Learning

- Query K sample trajectories for few-shot learning
 - Trajectory: a sequence of state-action pairs from starting state to ending state

Algorithm 1 Model-Agnostic Meta-Learning

Require: $p(\mathcal{T})$: distribution over tasks

Require: α, β : step size hyperparameters

- 1: randomly initialize θ
 - 2: **while** not done **do**
 - 3: Sample batch of tasks $\mathcal{T}_i \sim p(\mathcal{T})$
 - 4: **for all** \mathcal{T}_i **do**
 - 5: Evaluate $\nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$ with respect to K examples
 - 6: Compute adapted parameters with gradient descent: $\theta'_i = \theta - \alpha \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$
 - 7: **end for**
 - 8: Update $\theta \leftarrow \theta - \beta \nabla_{\theta} \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i})$
 - 9: **end while**
-

Sample K trajectories using
 f_{θ} in \mathcal{T}_i



Reinforcement Learning

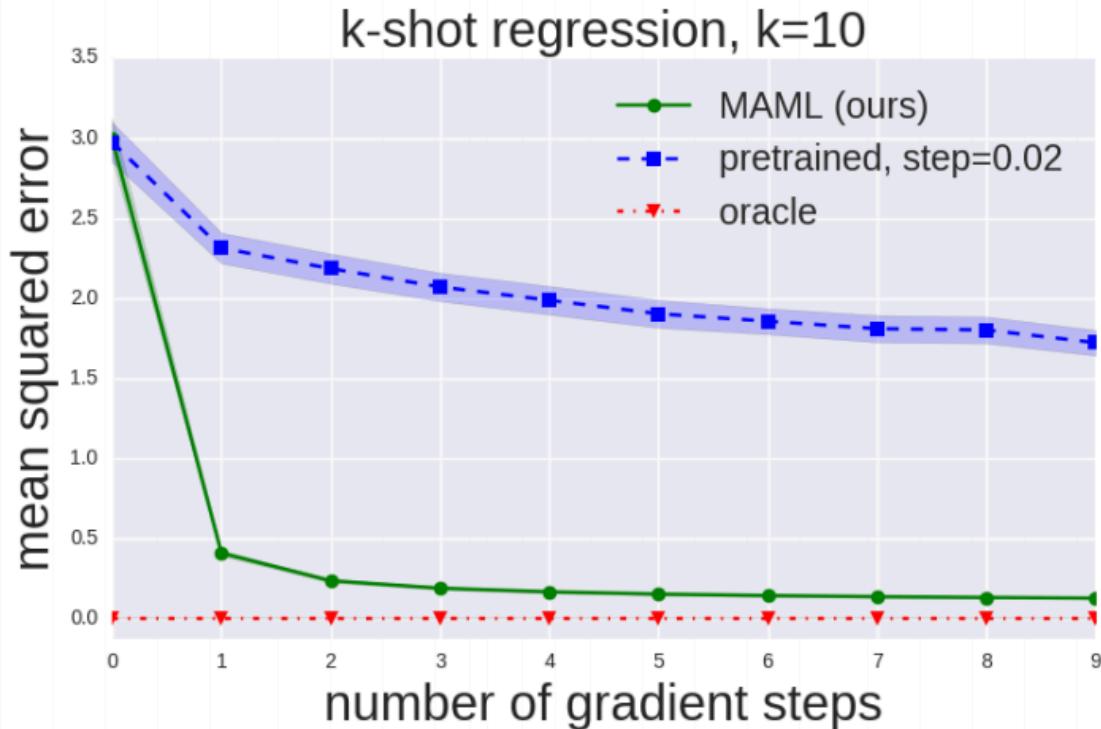
- Learn a policy
 - Map from states to a distribution over actions at each time step
- Loss function
 - Sum of negative reward

Evaluation: Regression

- Each regression task maps input to the output of a sine wave
 - The amplitude and phase of the sinusoid are varied
 - $f_i(x) = A \sin(x - \phi)$
- Comparisons:
 - Pre-training on all tasks and fine-tuning on new tasks
 - Use SGD on the K provided points from new tasks
 - An oracle which receives the true A and ϕ as input

Evaluation: Regression

- In-distribution: $A \in [0, 5]$, $\phi \in [0, \pi]$
- Out-of-distribution: $A \in [5, 10]$, $\phi \in [\pi, 2\pi]$

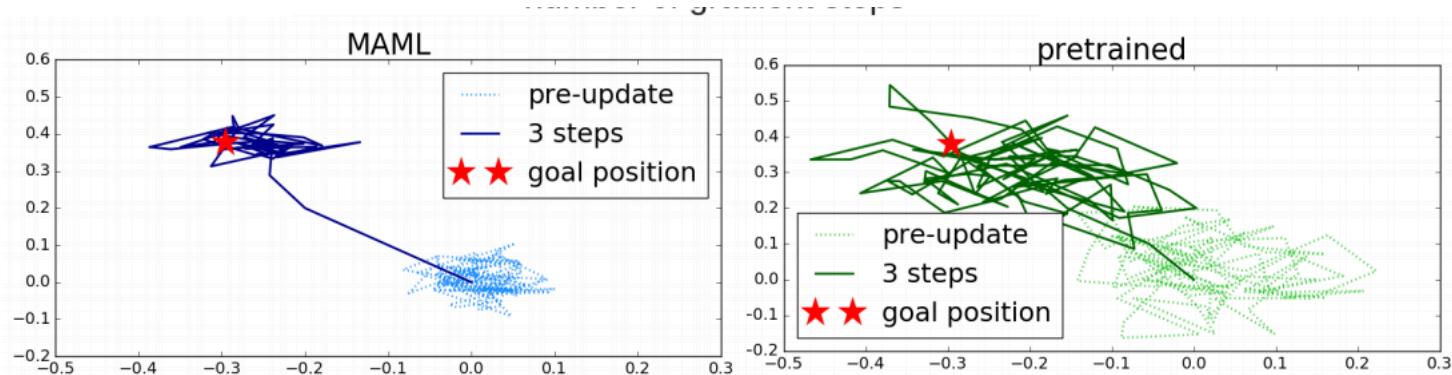


Evaluation: Reinforcement Learning

- A point agent must move to different goal positions in 2D square
 - Different tasks have different goal positions
 - Randomly chosen for each task within a unit square
- Observation
 - The current 2D position

Evaluation: Reinforcement Learning

- Actions
 - Velocity commands
 - With the range $[-0.1, 0.1]$
- Reward
 - The negative squared distance to the goal



Summary

- **Pros:**

- Reusing knowledge from past tasks
 - A crucial ingredient in making high-capacity scalable models
 - Lifelong few-shot learning
- Fast training with small datasets

- **Cons:**

- But needs to train on different types of tasks
- Multiple tasks are still related
 - The NN architecture needs to be the same
- The performance may not be as good as designing NN for a specific task